

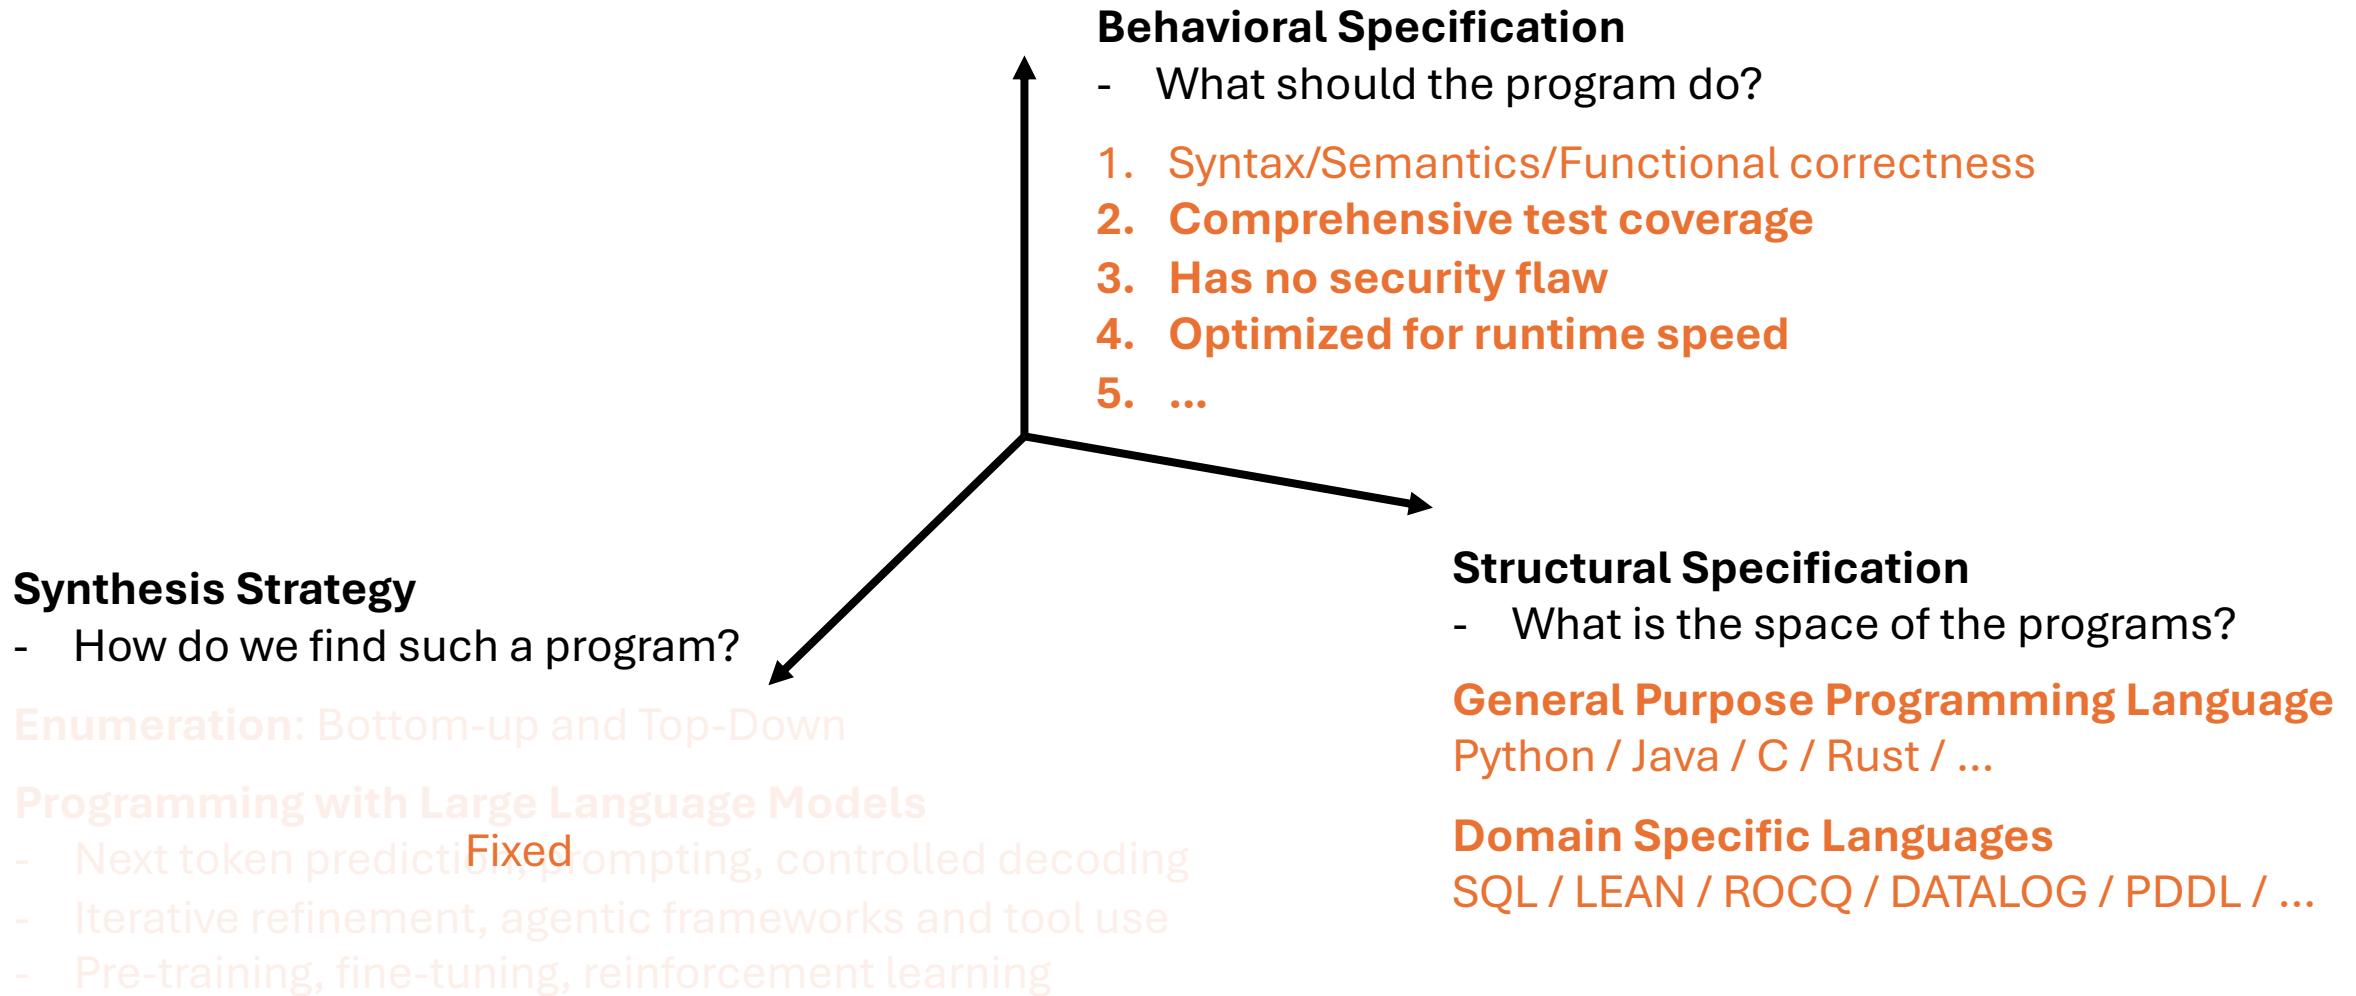
Machine Programming

Lecture 16 – LLM for Software Testing: Static Analysis

Logistics – Week 9

- Assignment 3: Coding LLM Agents
 - <https://github.com/machine-programming/assignment-3>
 - Fully functional web-app agent. Due: Oct 23 (Thu)
- Forming groups for your final projects!
 - Form a group of 2-3 before This Thursday (~~Oct 19~~ Oct 23)

Module 3: Overview



Software Analysis

Testing and Dynamic Analysis

Forms of Software Analysis

Static Analysis

Examine code without concrete execution

- Taint analysis
- Reachability analysis
- Abstract interpretation
- Symbolic execution
- ...

Dynamic Analysis

Observes program behavior while executing

- Unit testing
- Fuzzing
- Property-based testing
- Penetration testing
- ...

Dynamic Testing

Unit Testing

Human-written tests
in the form of concrete
input-output

```
assert f(2) == 4
```

Dynamic Testing

Unit Testing

Fuzzing

Human-written tests
in the form of concrete
input-output

Randomly
generating
test inputs

```
assert f(2) == 4          for i in 0..10^4:  
                           x = gen_input()  
                           f(x)
```

Dynamic Testing

Unit Testing

Human-written tests
in the form of concrete
input-output

```
assert f(2) == 4
```

Fuzzing

Randomly
generating
test inputs

```
for i in 0..10^4:  
    x = gen_input()  
    f(x)
```

Property-based testing

Randomly generating
test inputs, but with
properties checking outputs

```
for i in 0..10^4:  
    x = gen_input()  
    y = f(x)  
    assert property(x,y)
```

Dynamic Testing

Unit Testing

Human-written tests
in the form of concrete
input-output

```
assert f(2) == 4
```

Fuzzing

Randomly
generating
test inputs

```
for i in 0..10^4:  
    x = gen_input()  
    f(x)
```

Property-based testing

Randomly generating
test inputs, but with
properties checking outputs

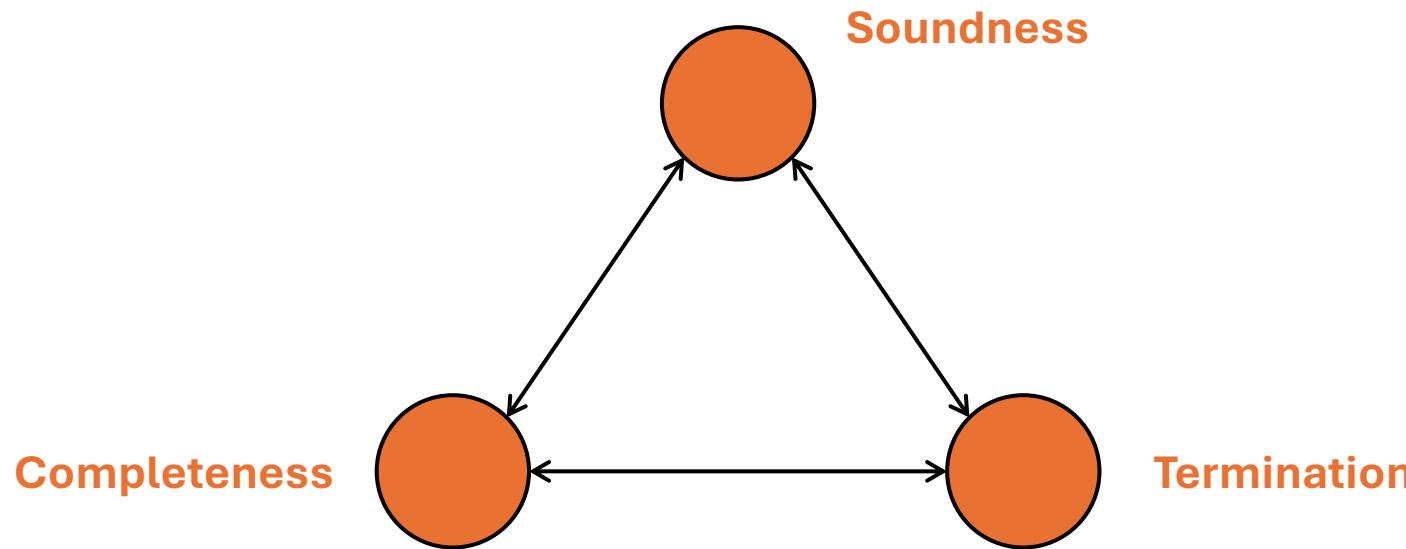
```
for i in 0..10^4:  
    x = gen_input()  
    y = f(x)  
    assert property(x,y)
```

Penetration testing

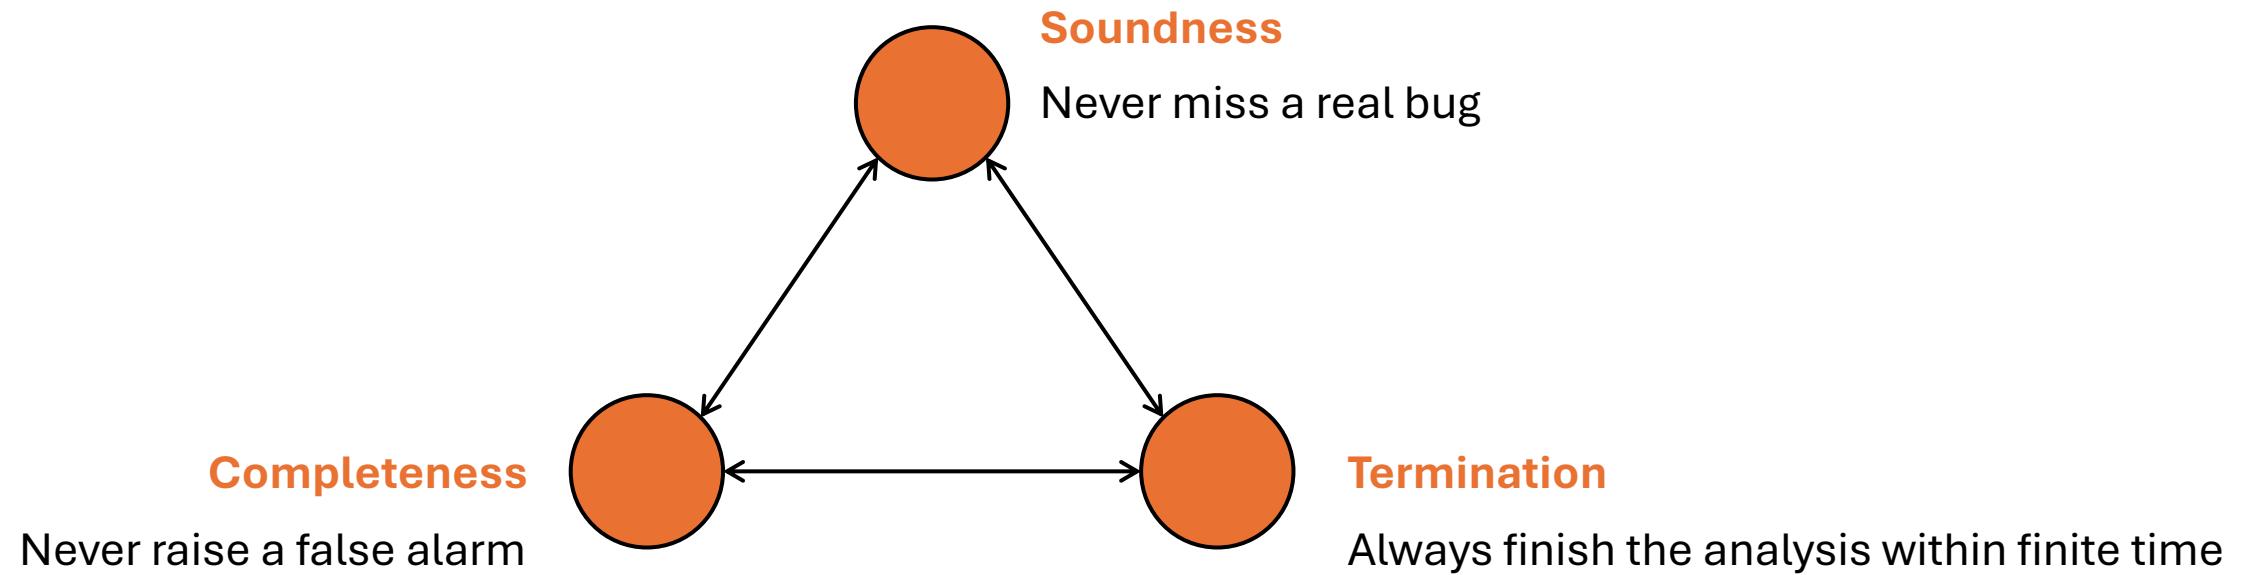
Crafting adversarial inputs
to trigger vulnerabilities rather
than getting outputs

```
try:  
    f(very_bad_input)  
    assert false  
except:  
    assert no_bad_behavior
```

Software Testing: Impossible Triangle

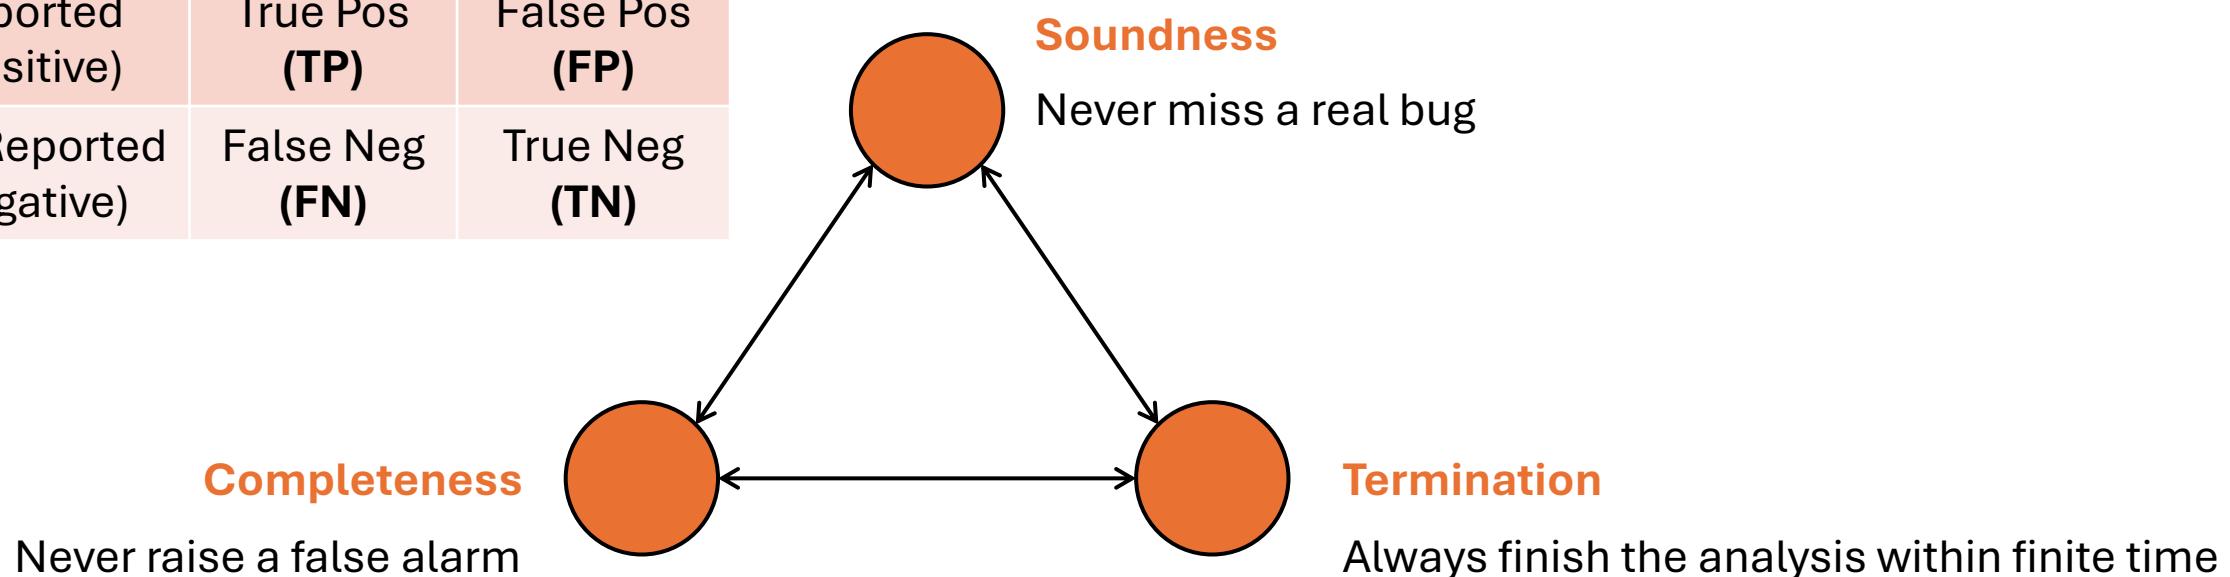


Software Testing: Impossible Triangle



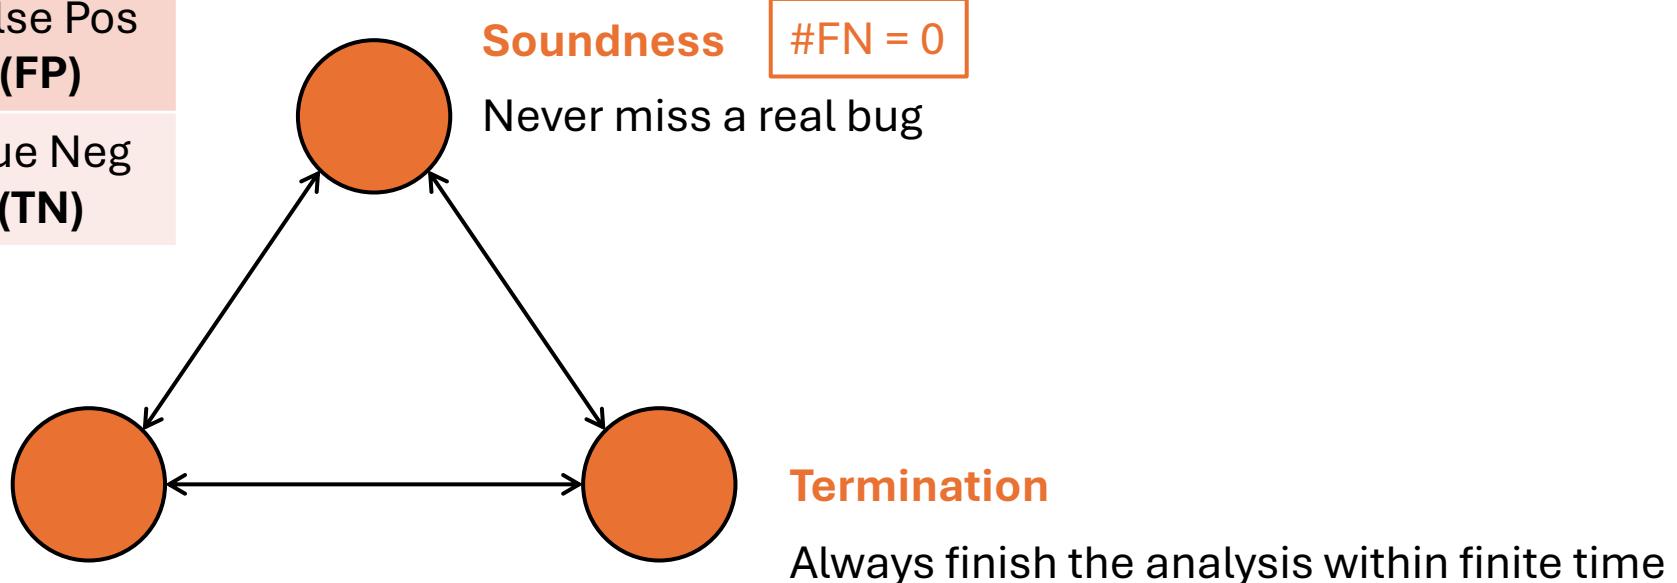
Software Testing: Impossible Triangle

	Real Bug	False Bug
Reported (positive)	True Pos (TP)	False Pos (FP)
Not Reported (negative)	False Neg (FN)	True Neg (TN)



Software Testing: Impossible Triangle

	Real Bug	False Bug
Reported (positive)	True Pos (TP)	False Pos (FP)
Not Reported (negative)	False Neg (FN)	True Neg (TN)



Software Testing: Impossible Triangle

Dynamic Analysis:

Unit Testing

- (assuming tests are correct)

Fuzzing

- (assuming harness is real)

Completeness

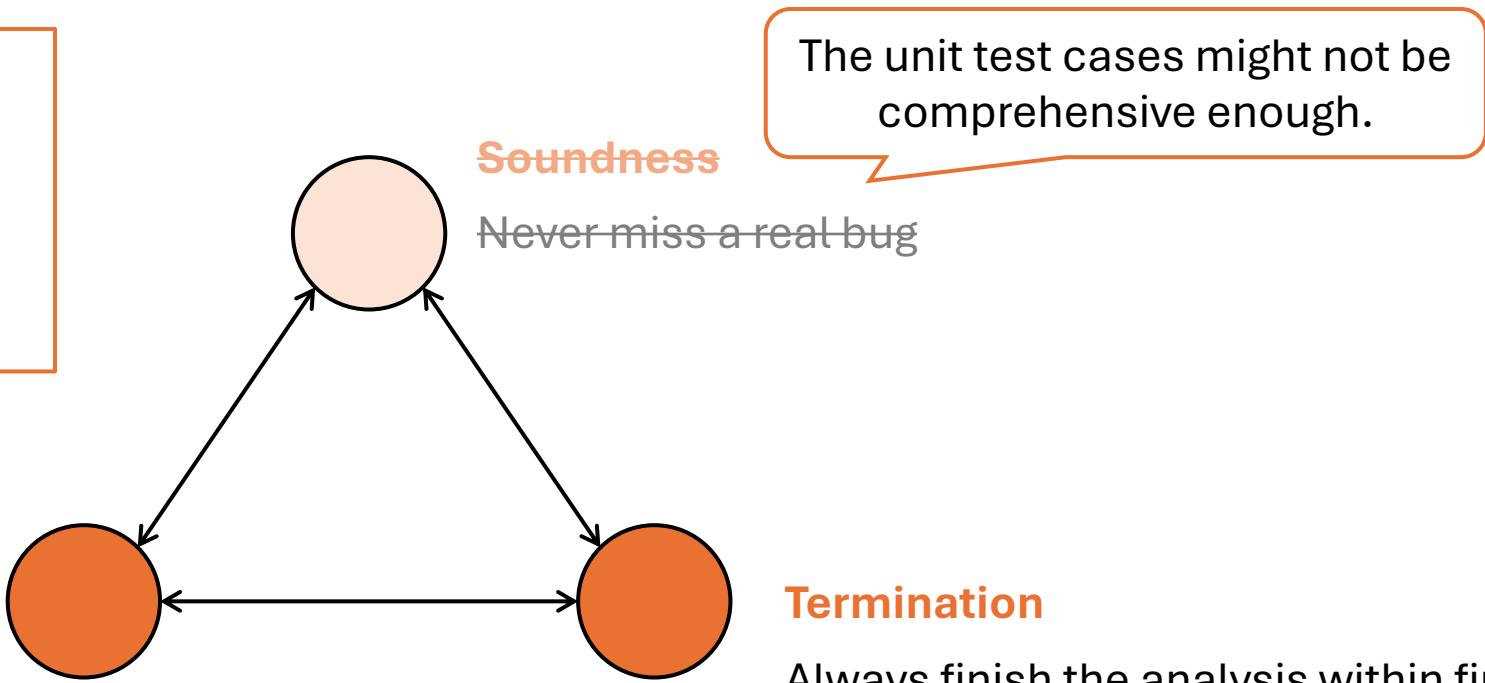
Never raise a false alarm

Soundness

Never miss a real bug

Termination

Always finish the analysis within finite time



Program – Line coverage: 5/5 (100%)

```
def calculate_interest(balance: float, months: int) -> float:  
    rate = 0.02 # 2% per month  
    if months > 12:  
        rate = 0.03 # higher rate for long-term deposits  
    return balance * rate / months
```

Test 1

```
def test_short_term_interest():  
    assert calculate_interest(1000, 6) == pytest.approx(3.33)
```

Test 2

```
def test_long_term_interest():  
    assert calculate_interest(1000, 24) == pytest.approx(1000 * 0.03 / 24)
```

Test 3 (Assume that this test case is not there)

```
def test_zero_months_divide_by_zero(): # covers implicit error path  
    assert calculate_interest(1000, 0) == 0
```

Software Testing: Impossible Triangle

Static Analysis:

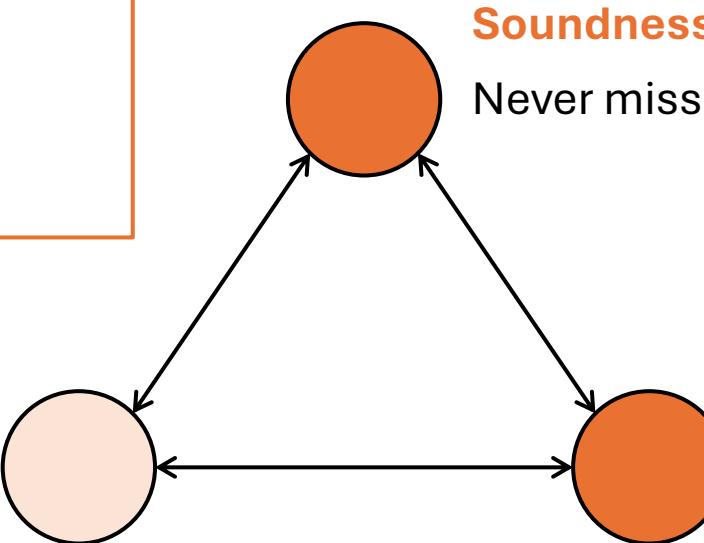
Taint Analysis

- (assuming correct specs)

Abstract Interpretation

- (assuming correct domains)

Completeness
Never raise a false alarm



Termination
Always finish the analysis within finite time

May report an alarm for developer confirmation, and that is rejected

Software Testing: Impossible Triangle

Formal Verification:

Halting Problem

- Is the program terminating?

Functional Correctness Verification

- (assuming specs are correct)

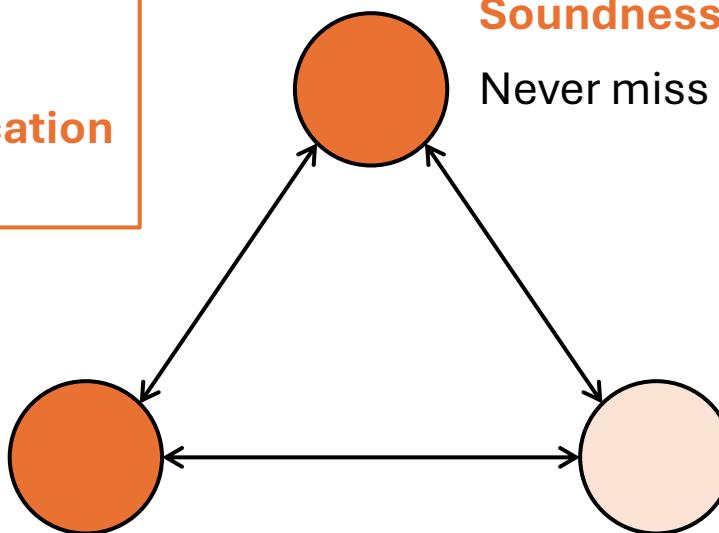
Completeness
Never raise a false alarm

Soundness

Never miss a real bug

Termination
~~Always finish the analysis within finite time~~

It may be impossible for a verifier to terminate



Software Analysis

Static Analysis

Software Testing: Impossible Triangle

Static Analysis:

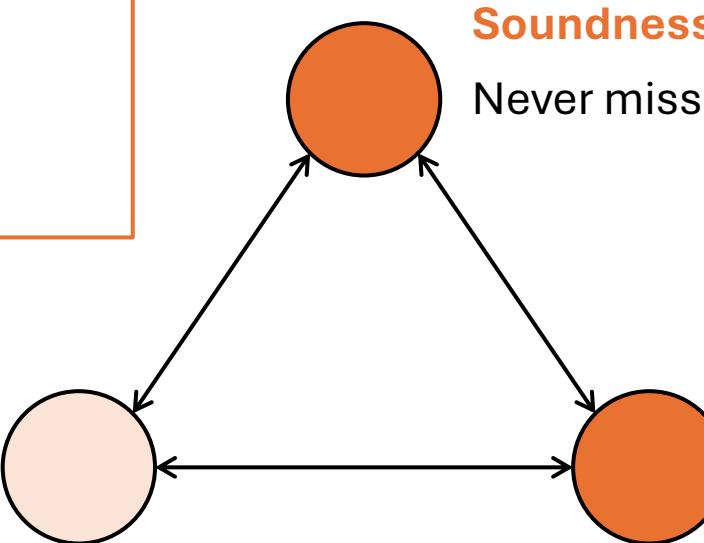
Taint Analysis

- (assuming correct specs)

Abstract Interpretation

- (assuming correct domains)

Completeness
Never raise a false alarm



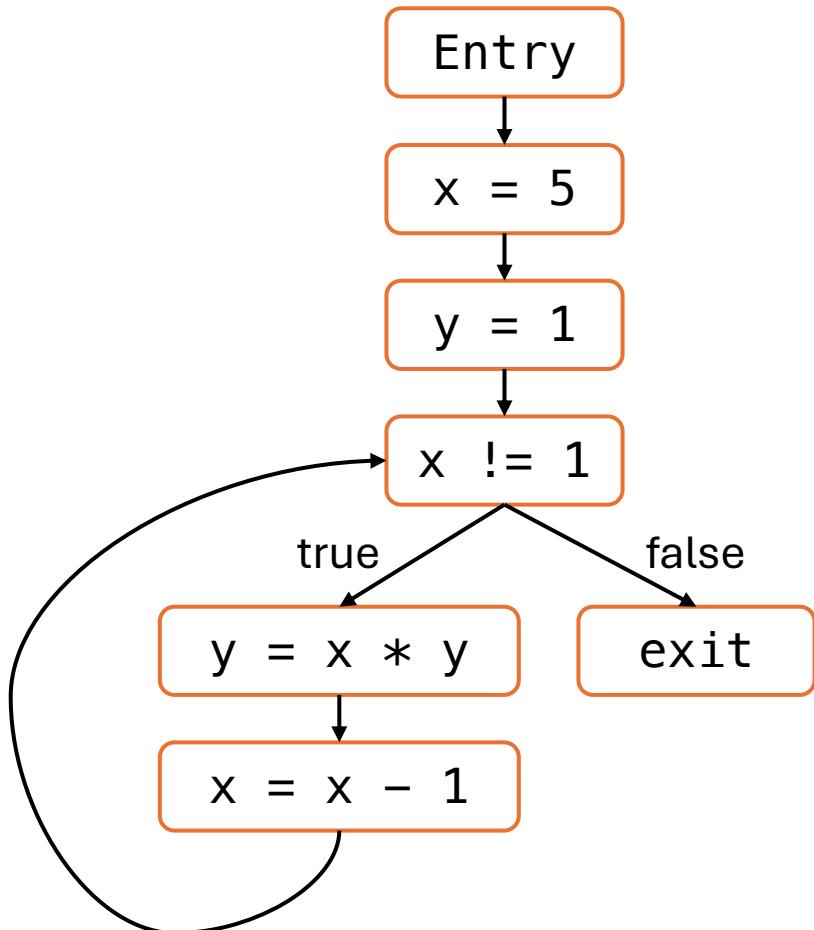
Termination
Always finish the analysis within finite time

May report an alarm for developer confirmation, and that is rejected

Control-Flow Graphs

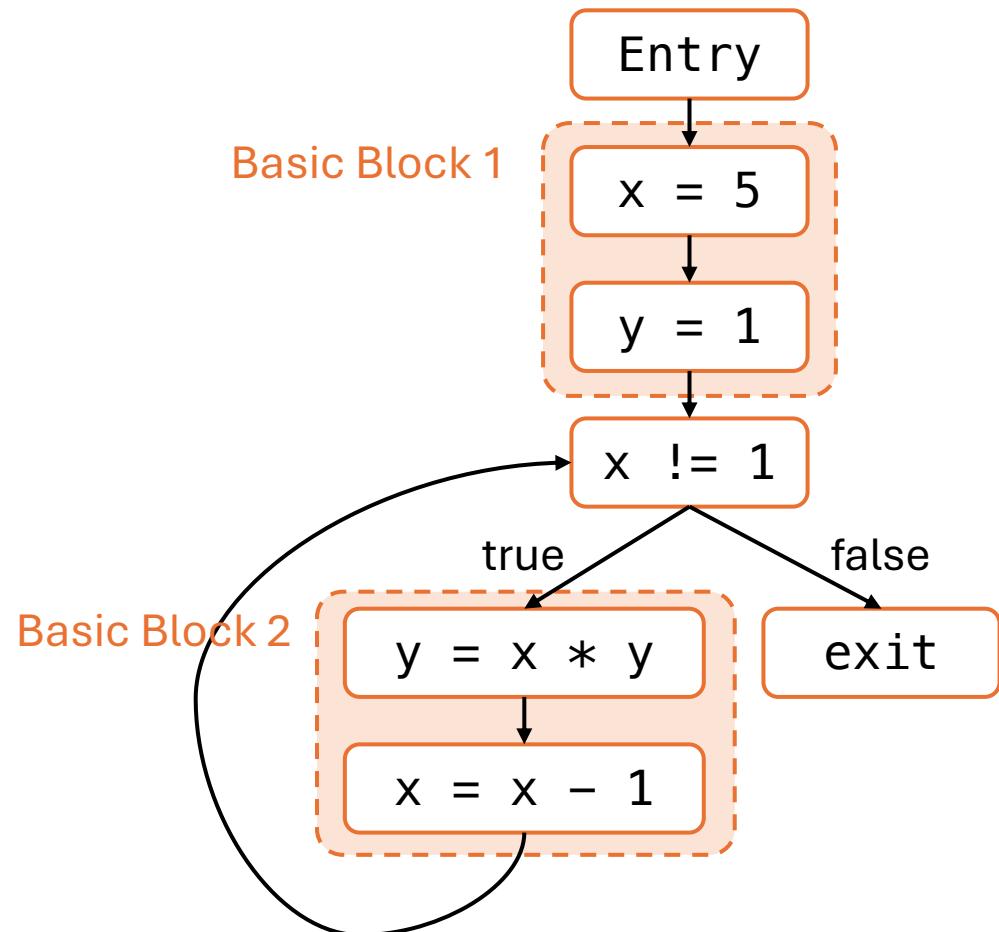
- Premise:
 - An abstract representation of imperative programs
 - Illustrating how “control” (execution) flows through the code
- A graph representation
 - Summarizes the flow of control in all possible runs of the program
- Goal:
 - Help all sorts of static analysis by providing the birds eye view

Control-Flow Graphs



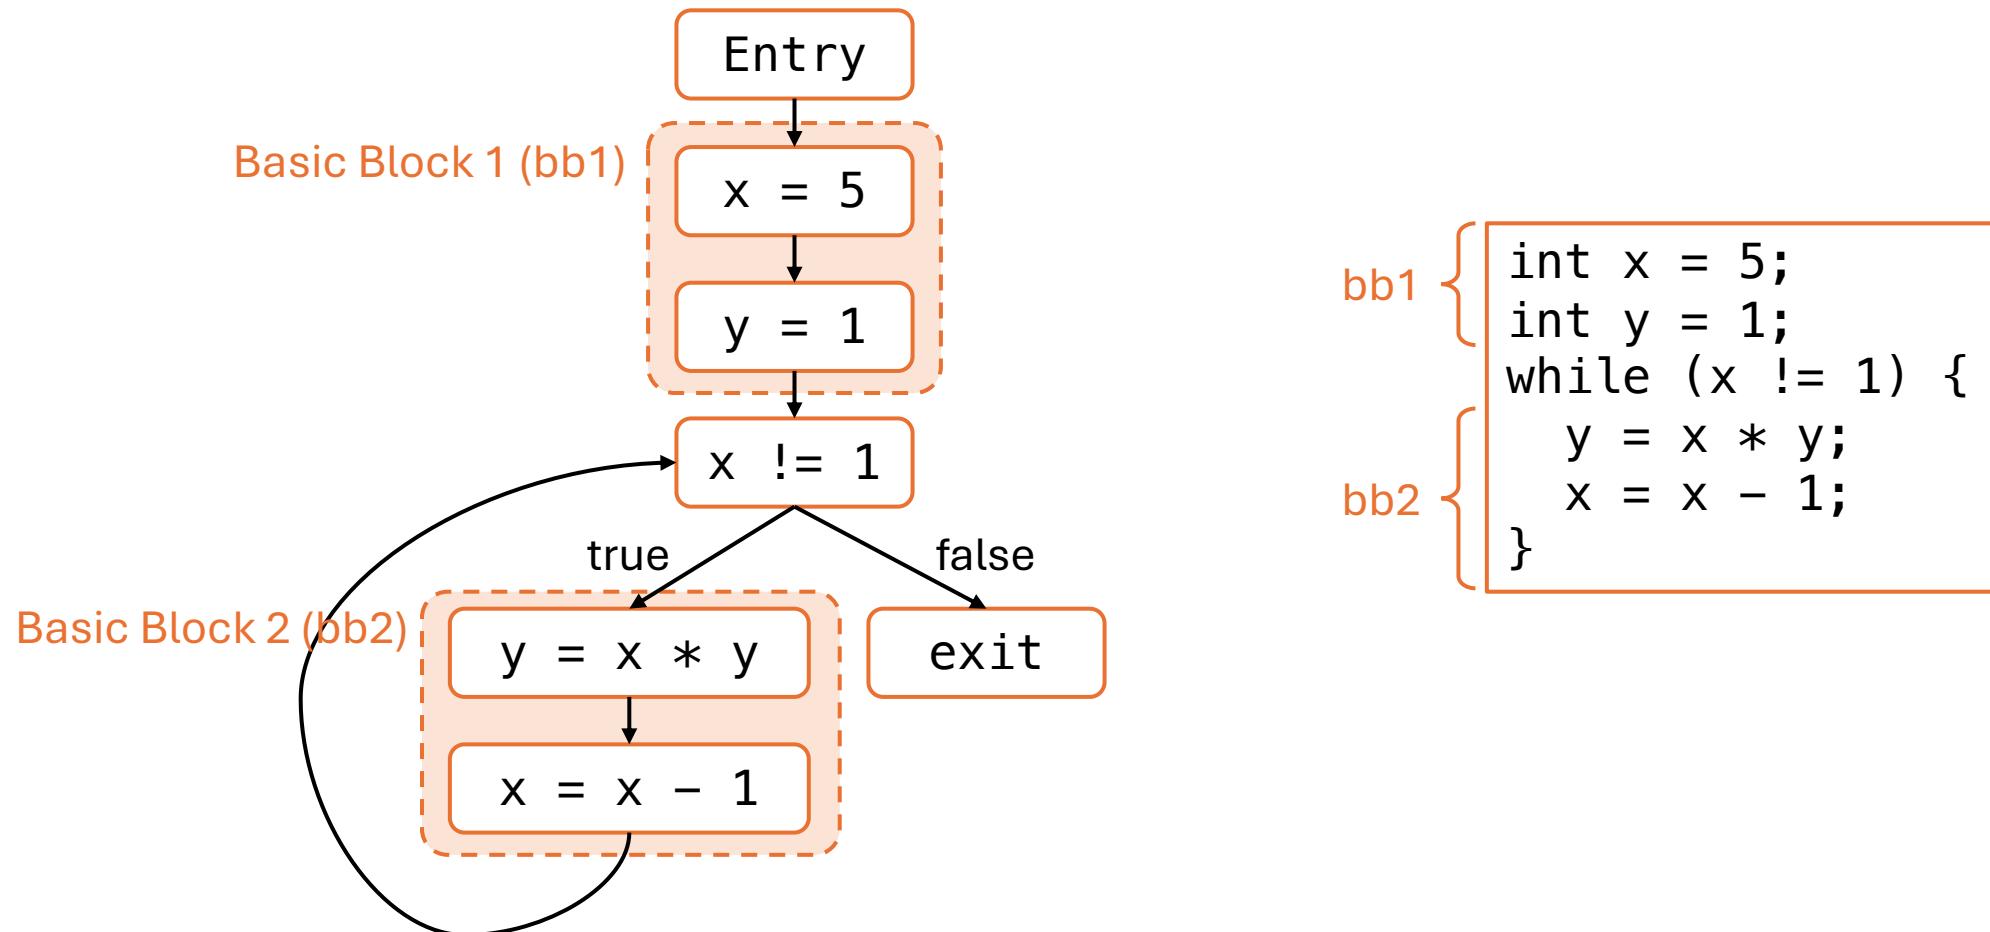
```
int x = 5;
int y = 1;
while (x != 1) {
    y = x * y;
    x = x - 1;
}
```

Control-Flow Graphs

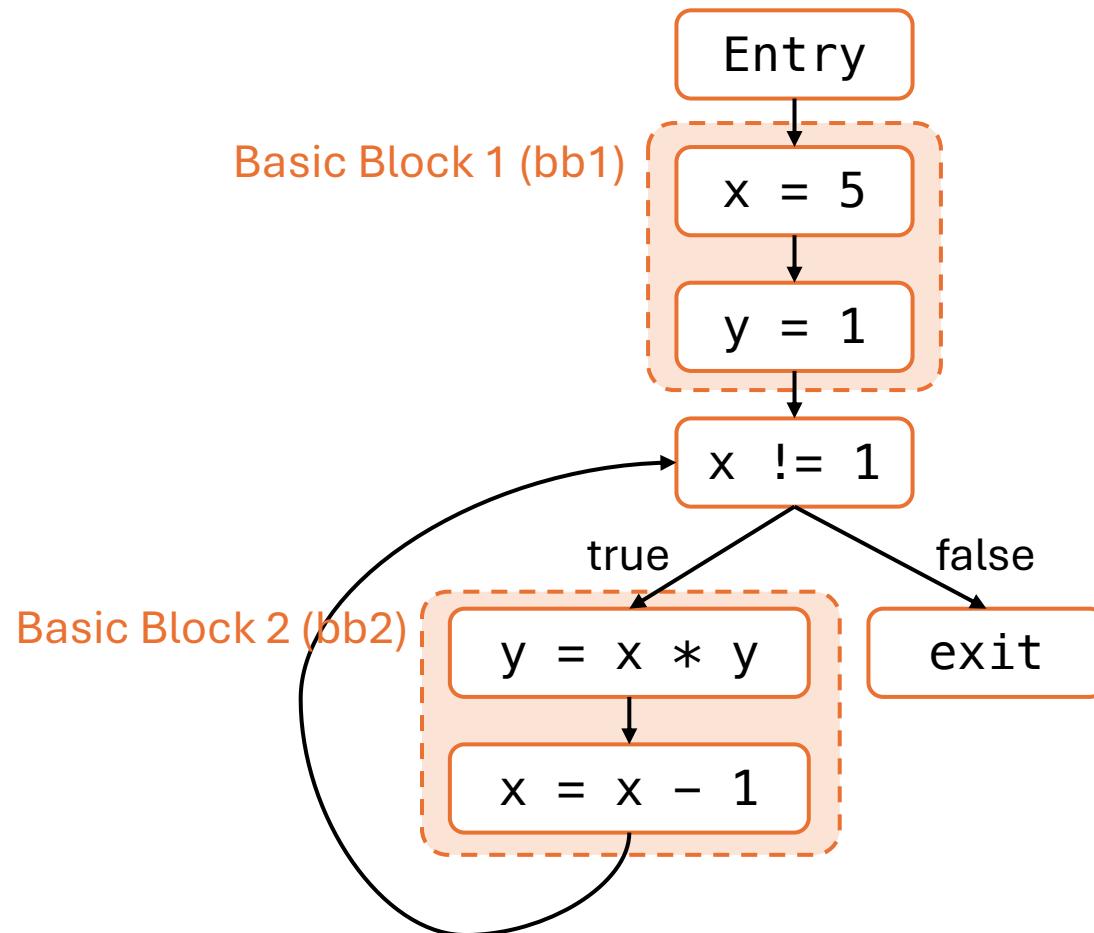


```
int x = 5;
int y = 1;
while (x != 1) {
    y = x * y;
    x = x - 1;
}
```

Control-Flow Graphs



Control-Flow Graphs: Traces



int x = 5;
int y = 1;
while (x != 1) {
 y = x * y;
 x = x - 1;
}

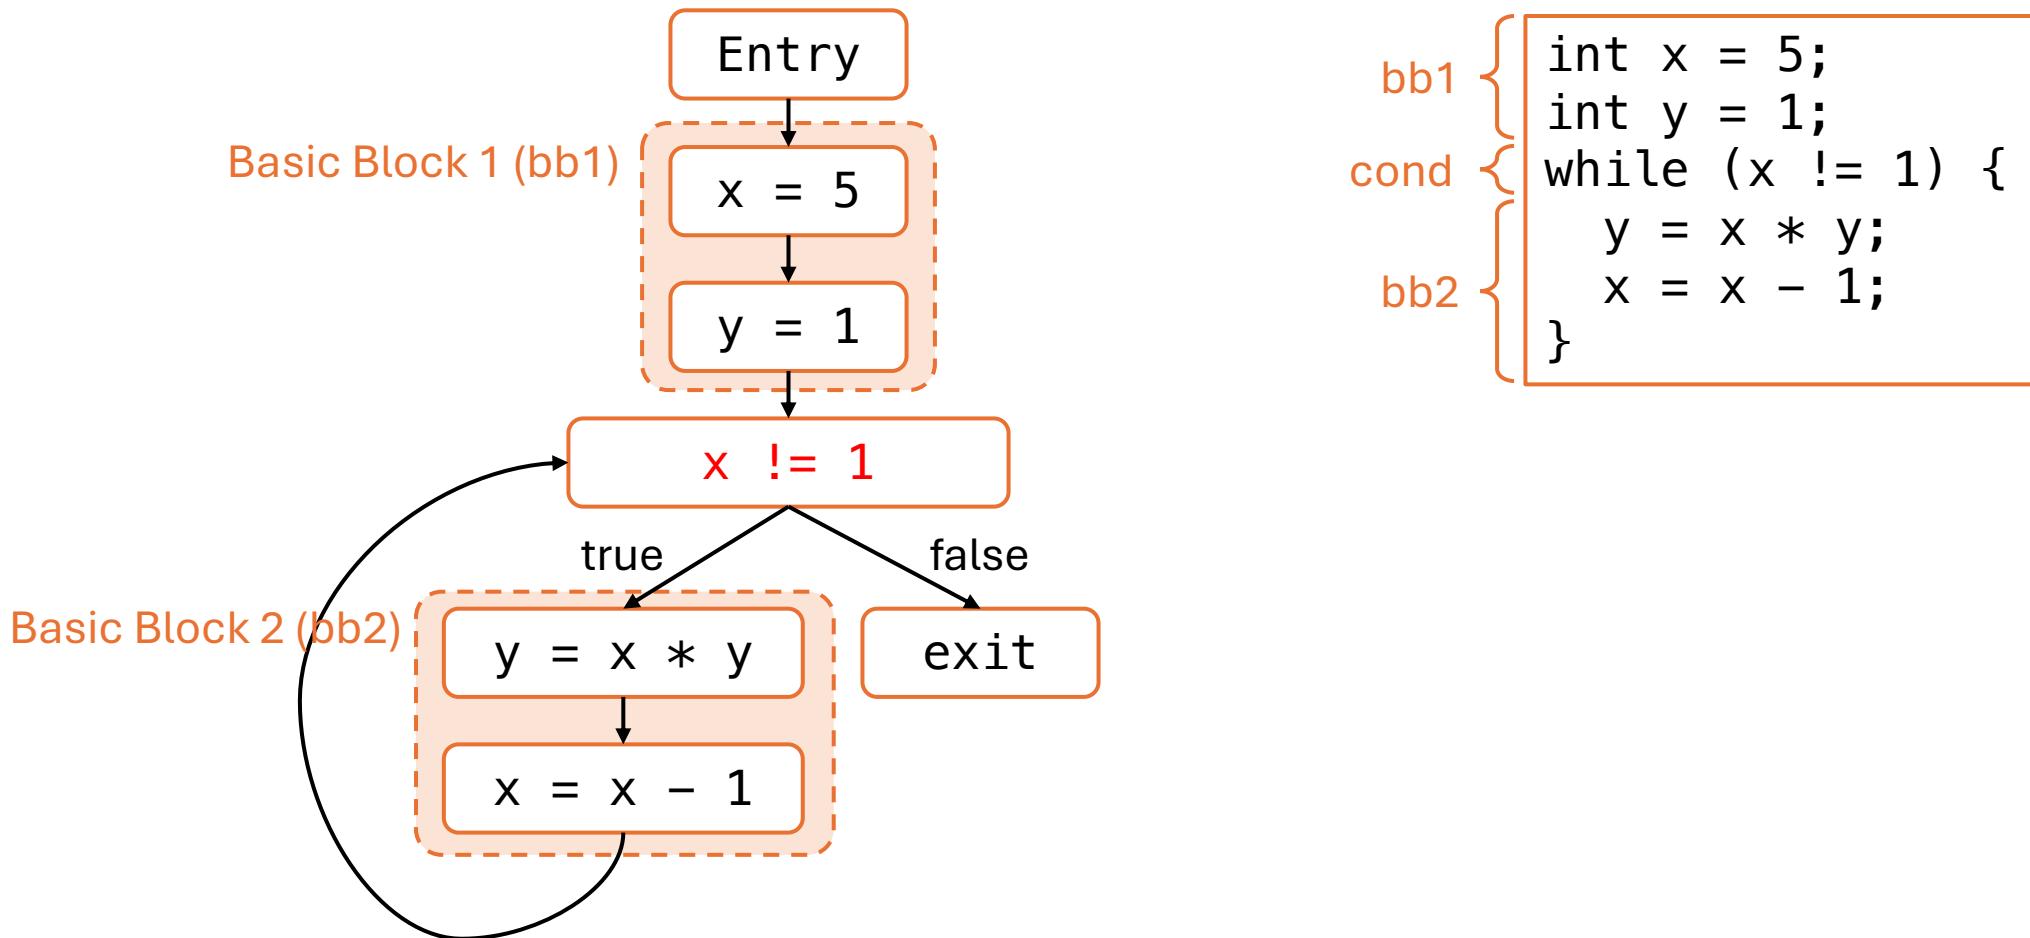
Execution Trace

Entry → bb1 → cond
→ bb2 → cond // (x == 5)
→ bb2 → cond // (x == 4)
→ bb2 → cond // (x == 3)
→ bb2 → cond // (x == 2)
→ bb2 → cond // (x == 1)
→ exit

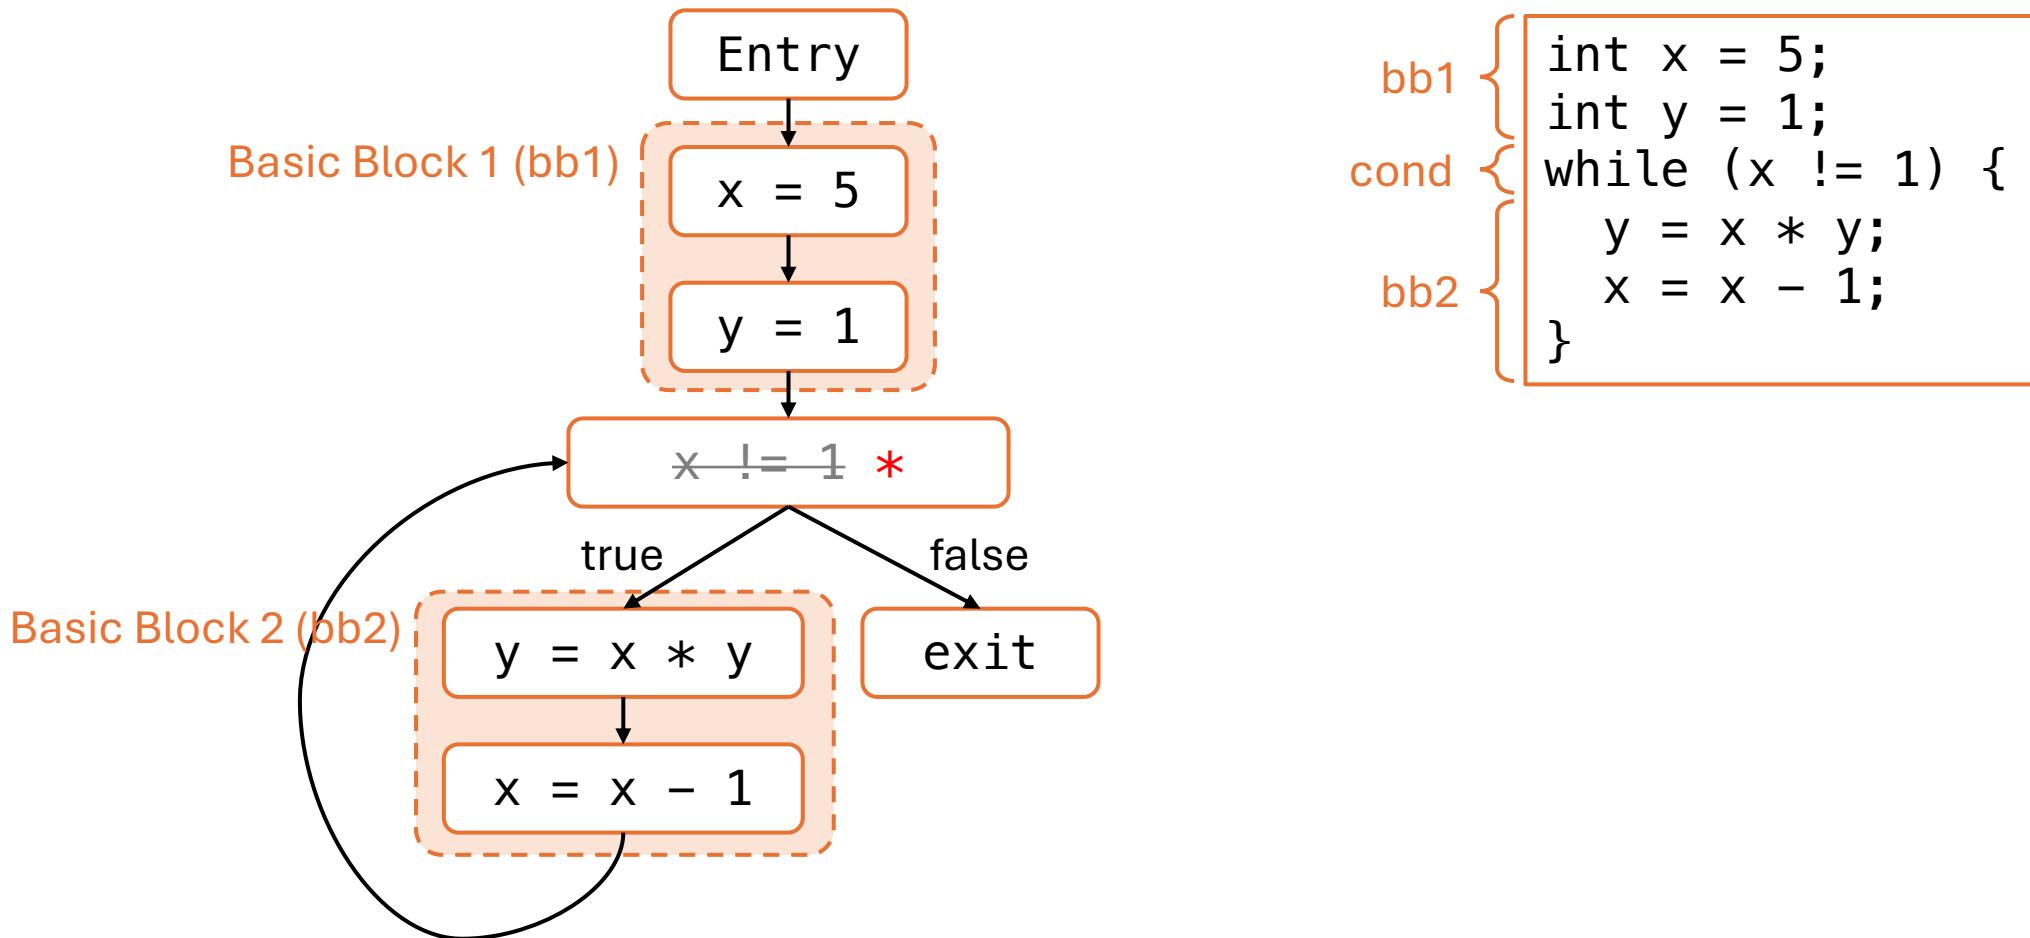
Dataflow Analysis

- Building on control-flow analysis
 - Abstracts away control-flow conditions
- Considers all paths possible in actual runs (**sound**)
- Including paths that are never realizable (**incomplete**)

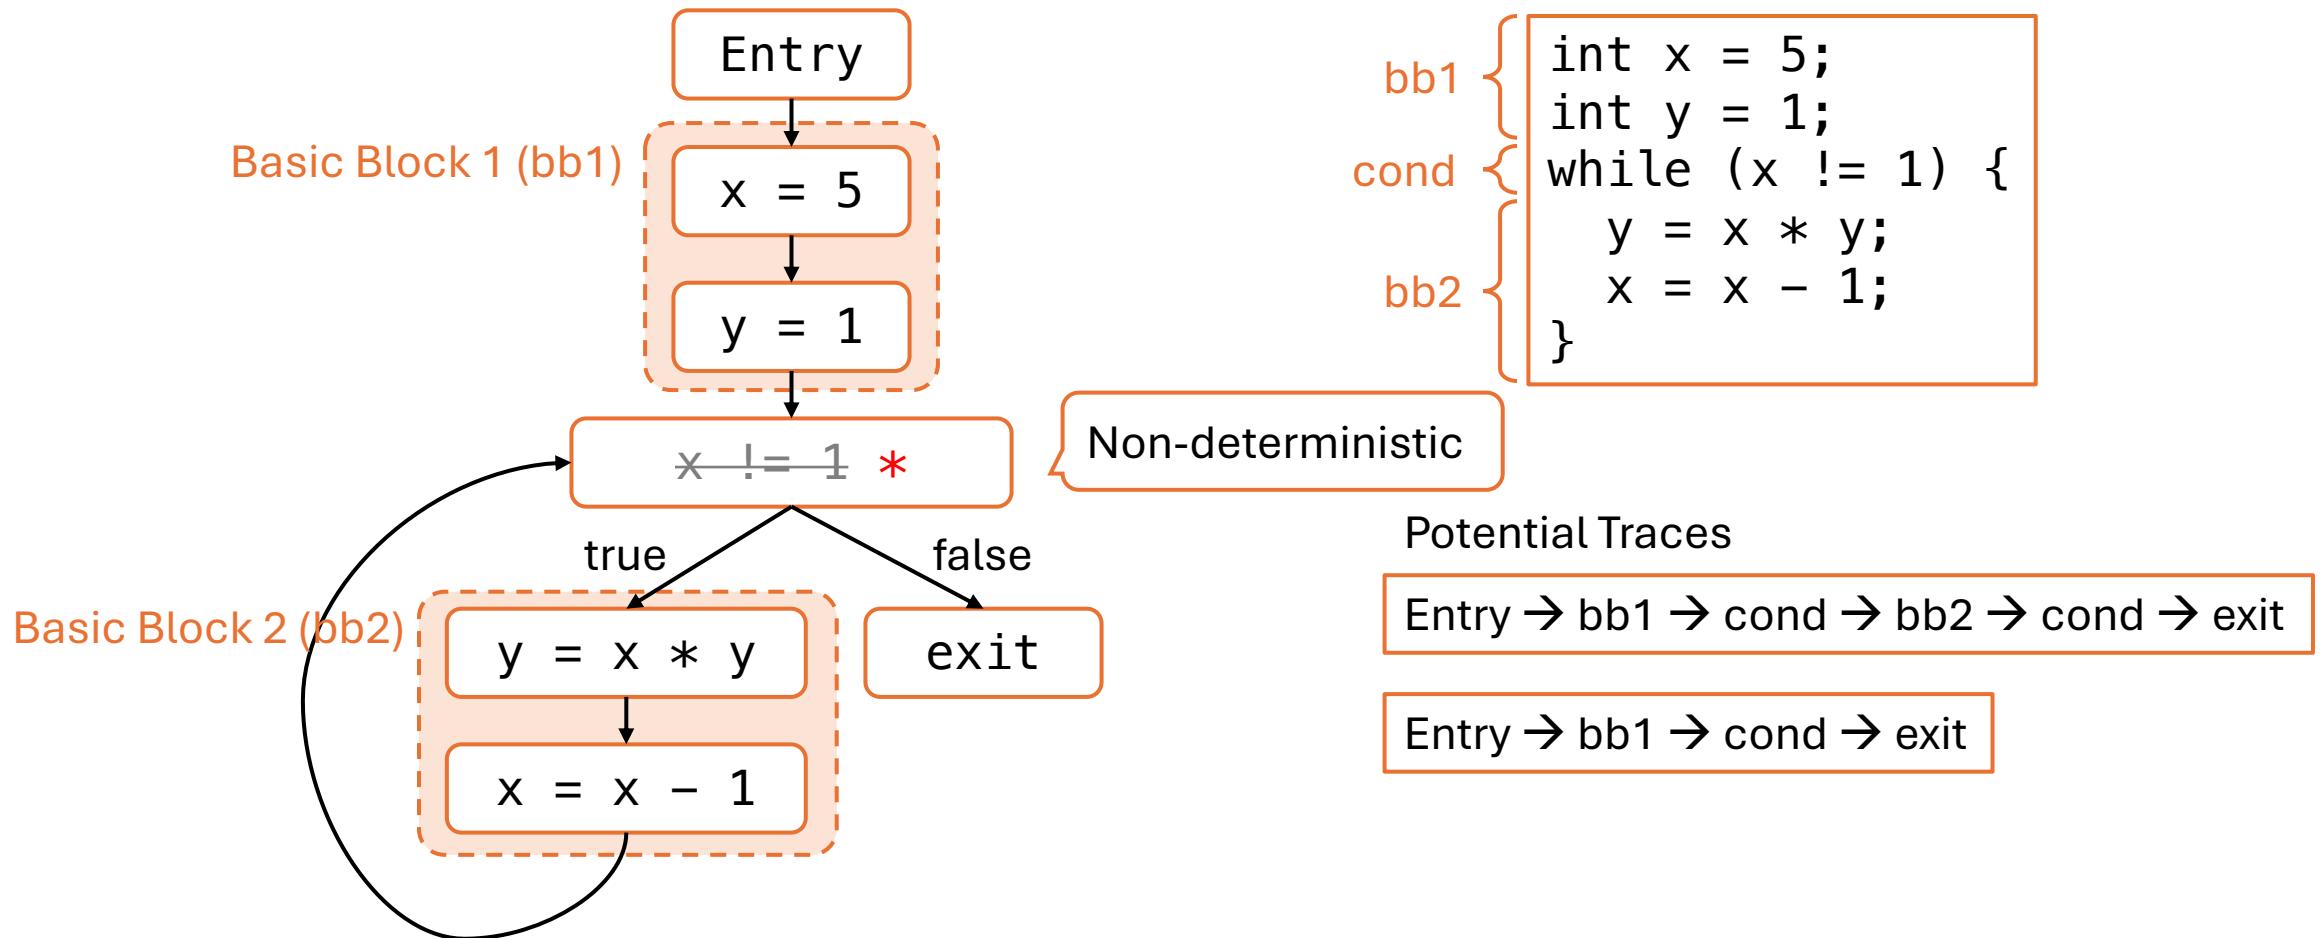
Abstracting Control-Flow Graphs



Abstracting Control-Flow Graphs



Abstracting Control-Flow Graphs



Modern Dataflow Analyses

Interval Analysis

- Check memory safety
(integer overflows, buffer overruns, ...)

Taint Analysis

- Check information flow
(Sensitive data leak, code injection, ...)

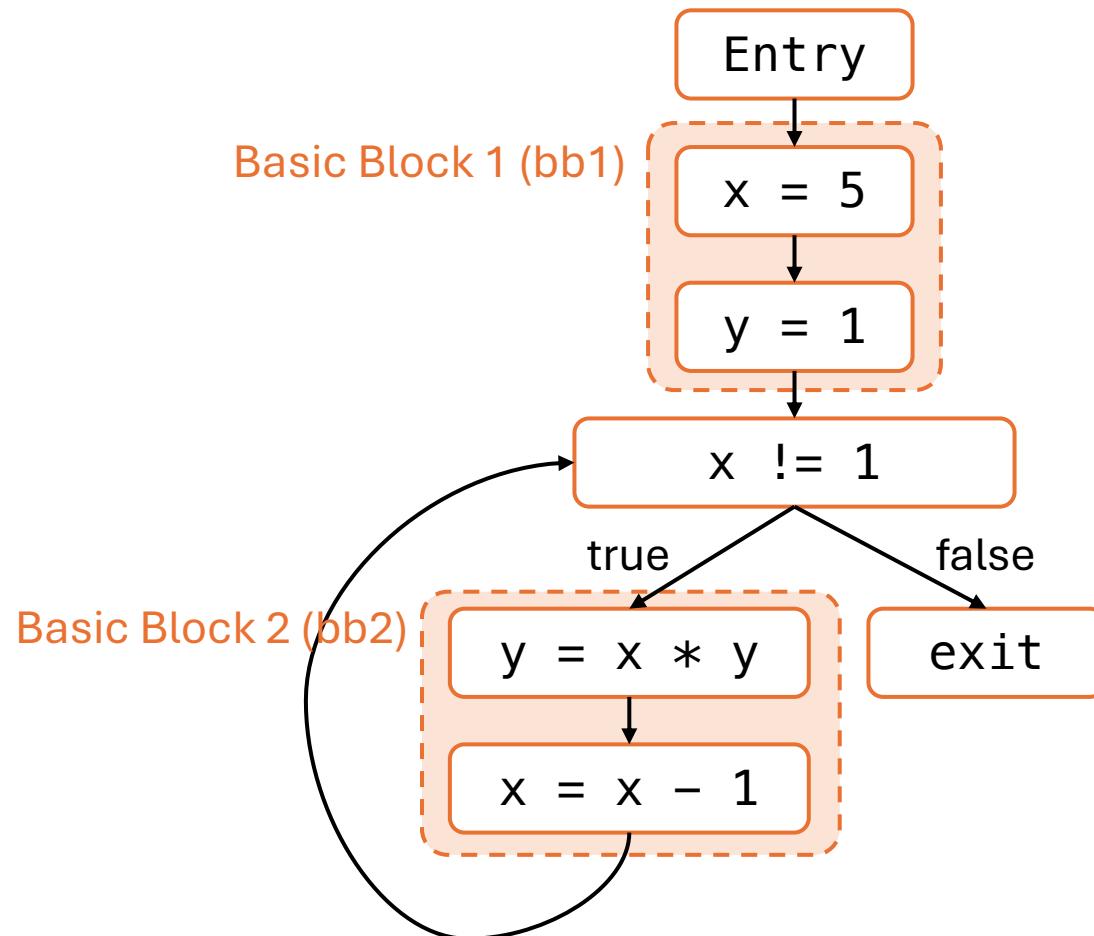
Type-State Analysis

- Temporal safety properties
(APIs of protocols, libraries, ...)

Concurrency Analysis

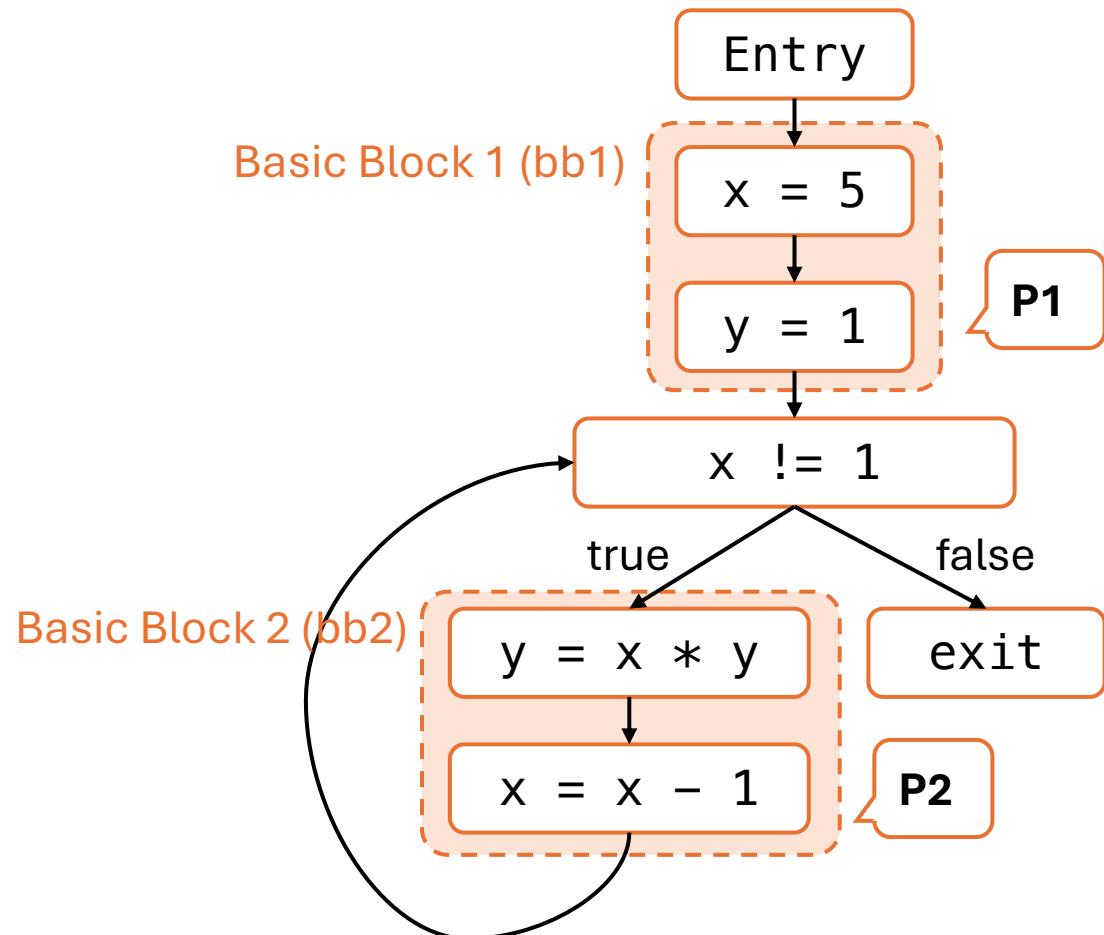
- Concurrency safety properties
(dataraces, deadlocks, ...)

Reaching Definitions Analysis



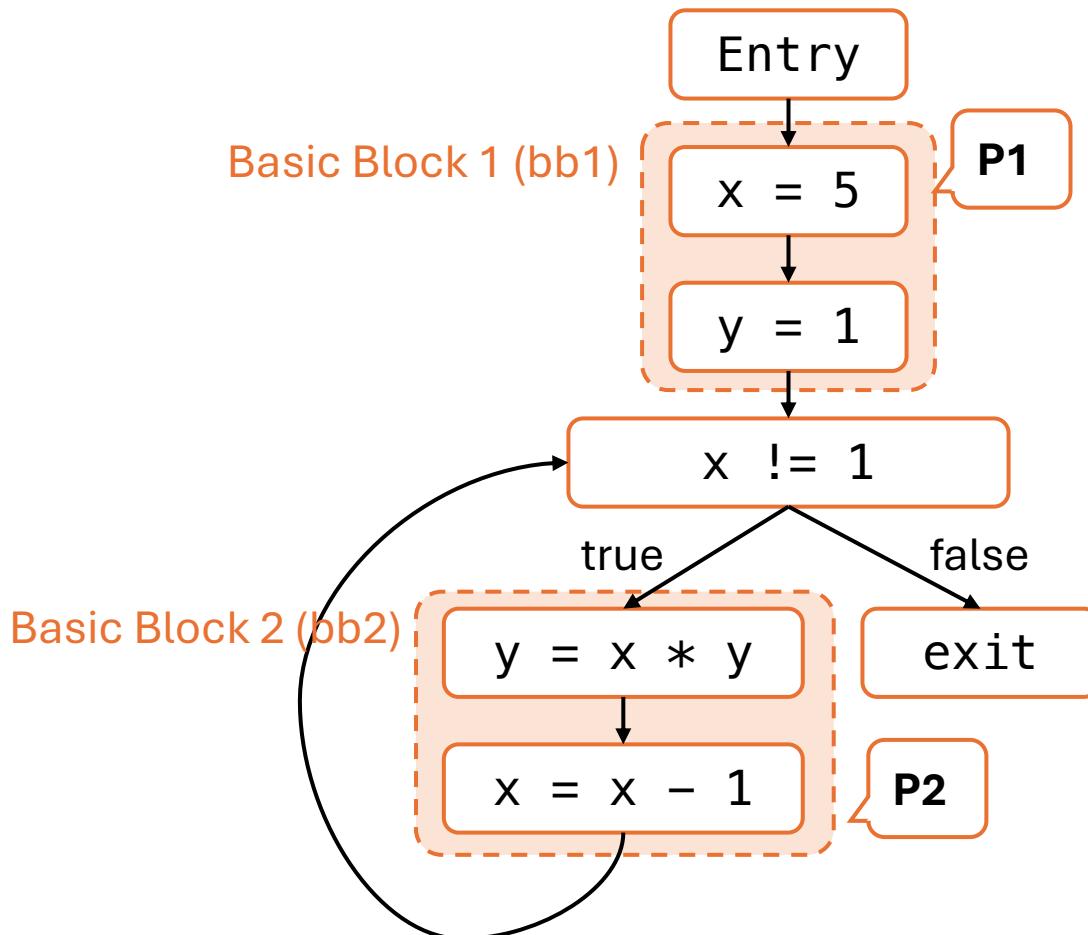
- Goal:
 - Determine whether a program point containing a piece of data (definition of a variable) can flow to another program point
- Reduced to...
 - Path finding problem in a graph

Reaching Definitions Analysis



- Goal:
 - Determine whether a program point containing a piece of data (definition of a variable) can flow to another program point
- Reduced to...
 - Path finding problem in a graph

Reaching Definitions Analysis



- Goal:

- Determine whether a program point containing a piece of data (definition of a variable) can flow to another program point

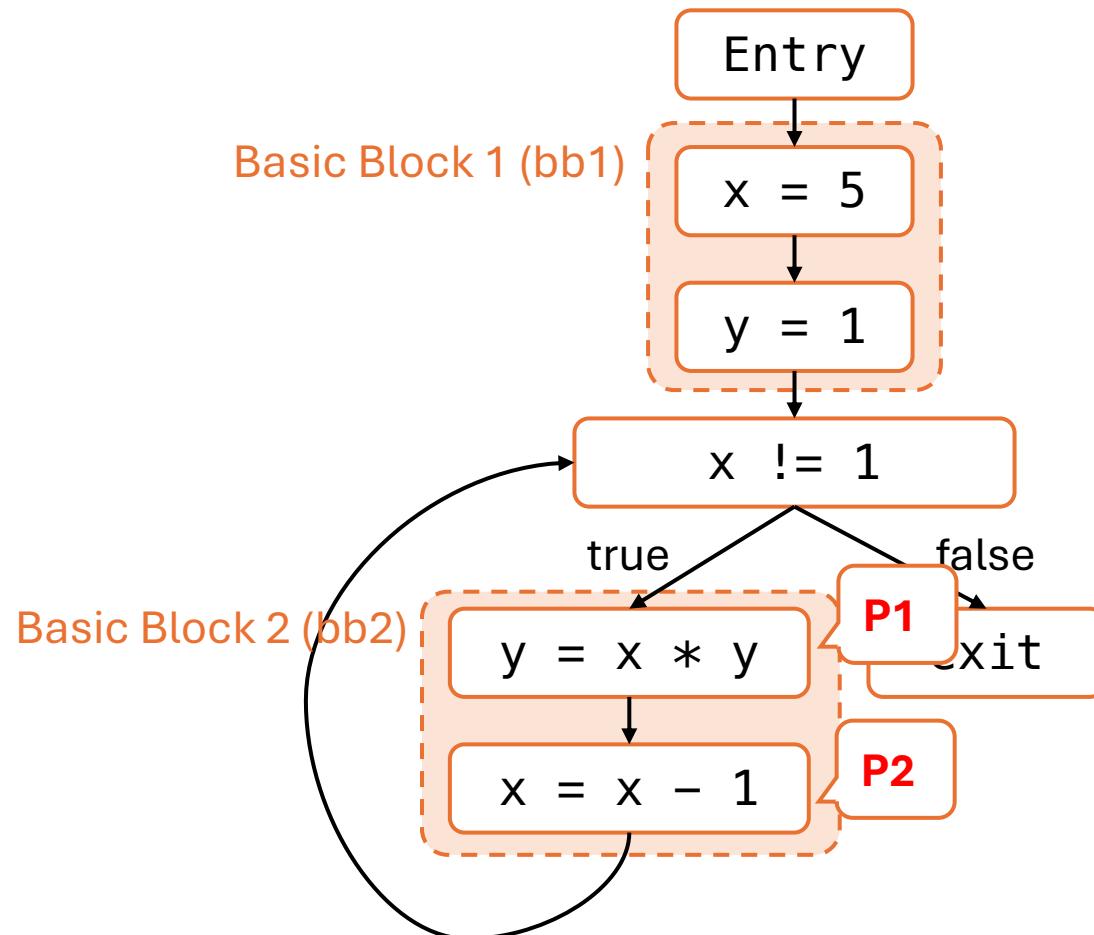
- Reduced to...

- Path finding problem in a graph

- Example:

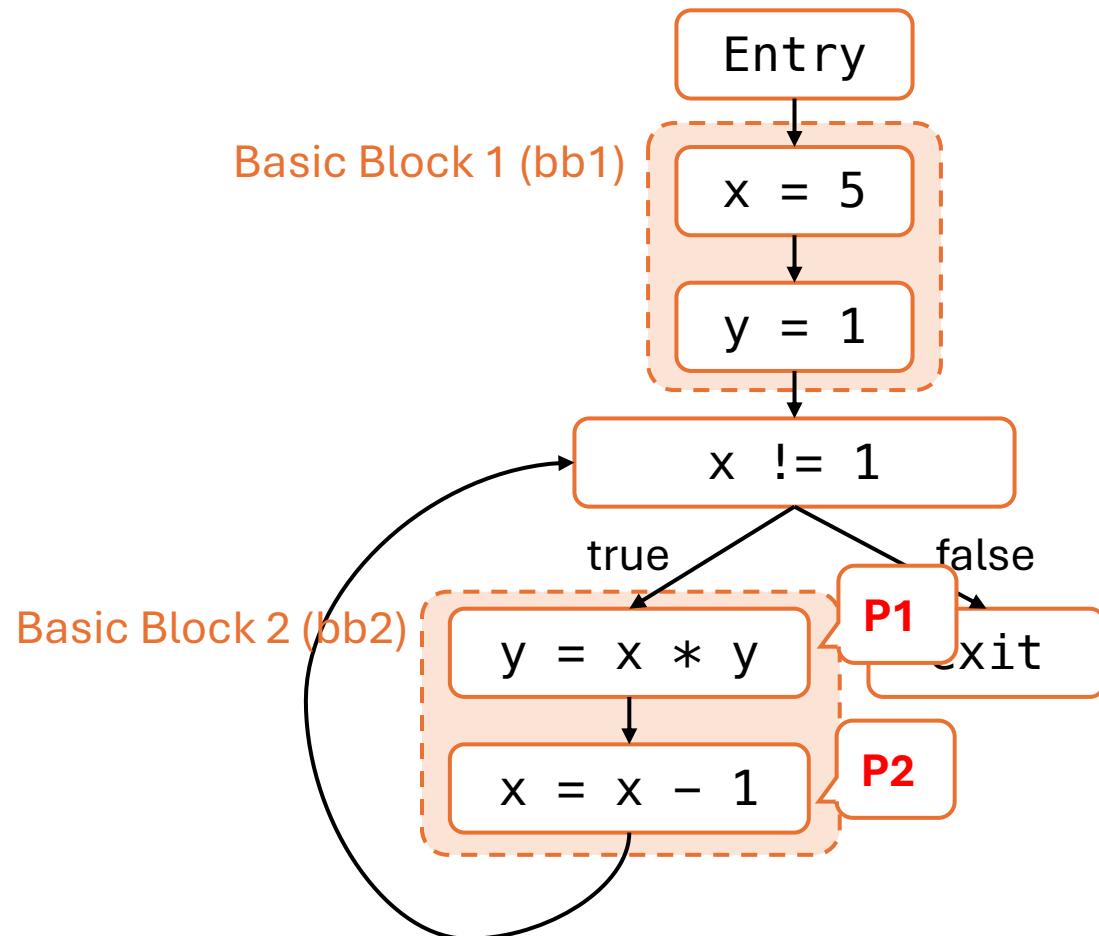
- Does information defined in **P1** flow to **P2**?

Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

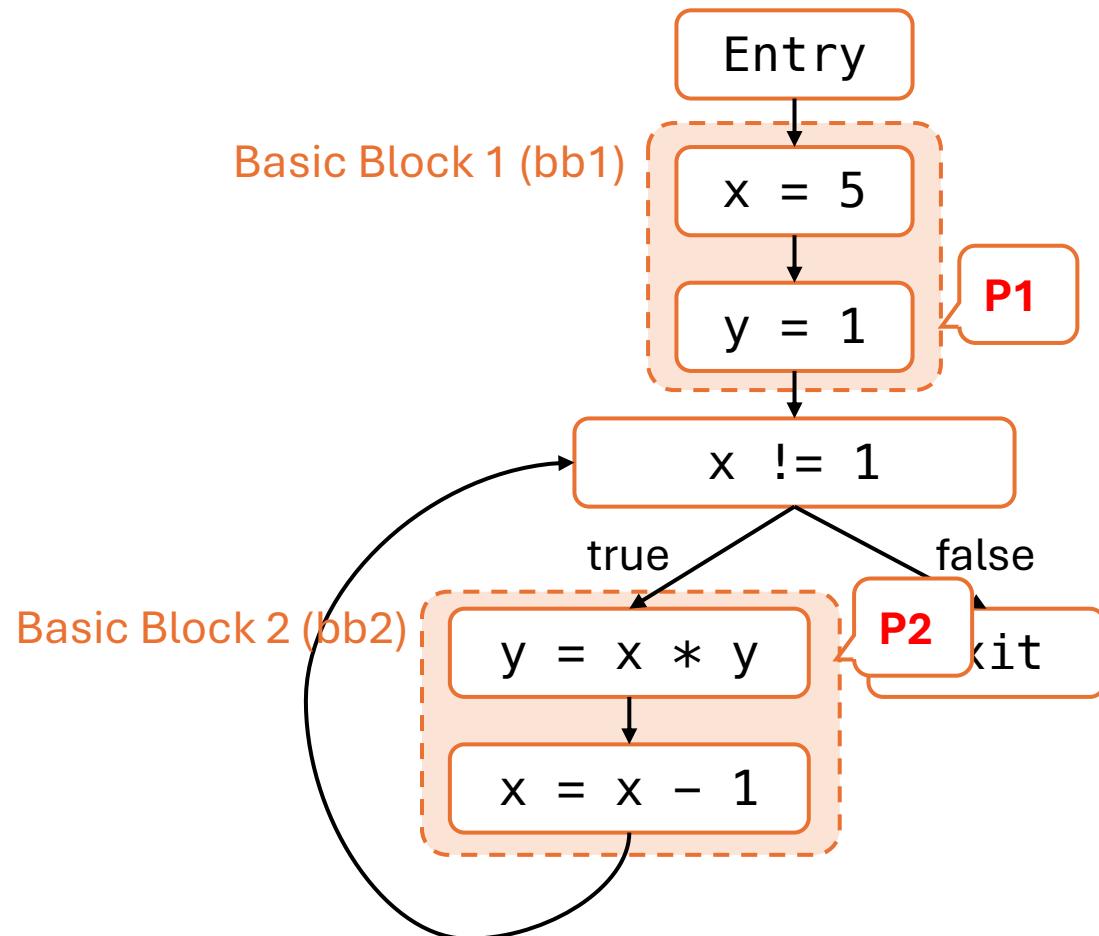
Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

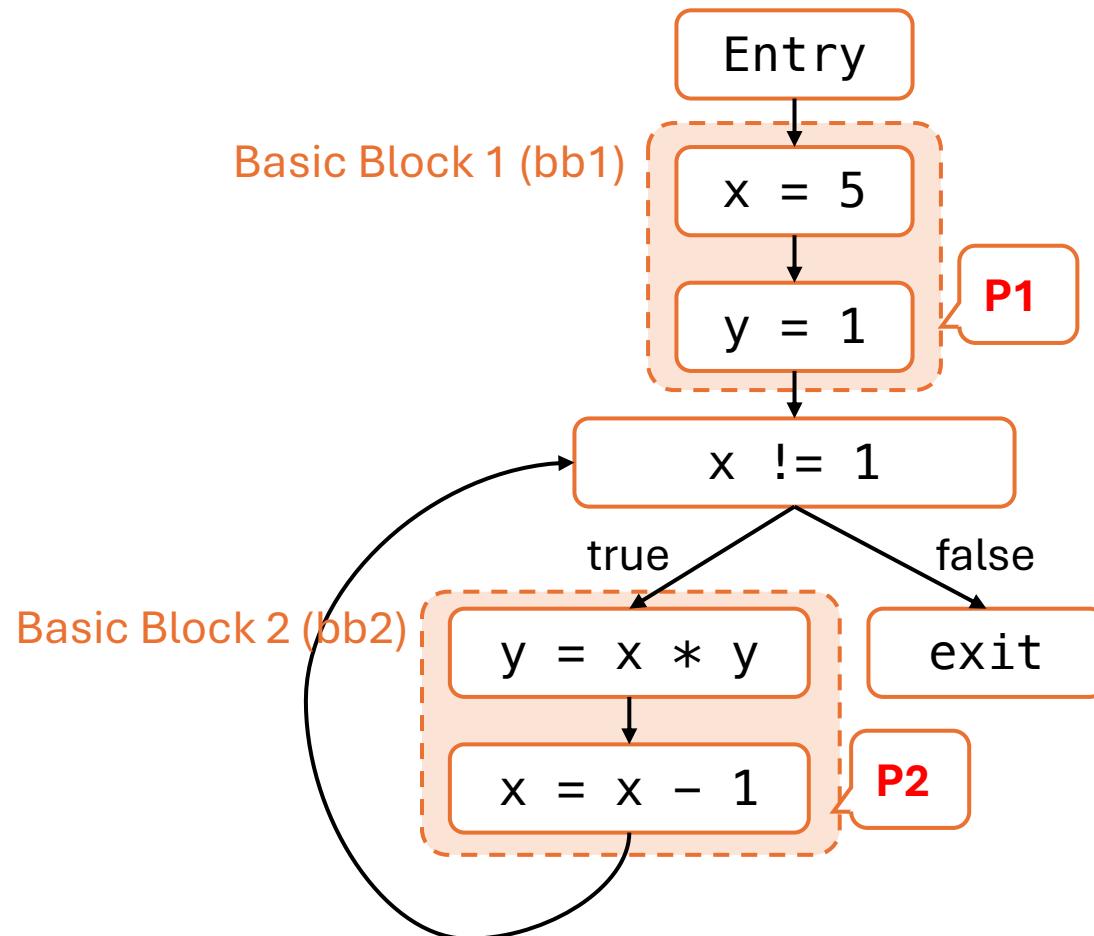
YES

Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

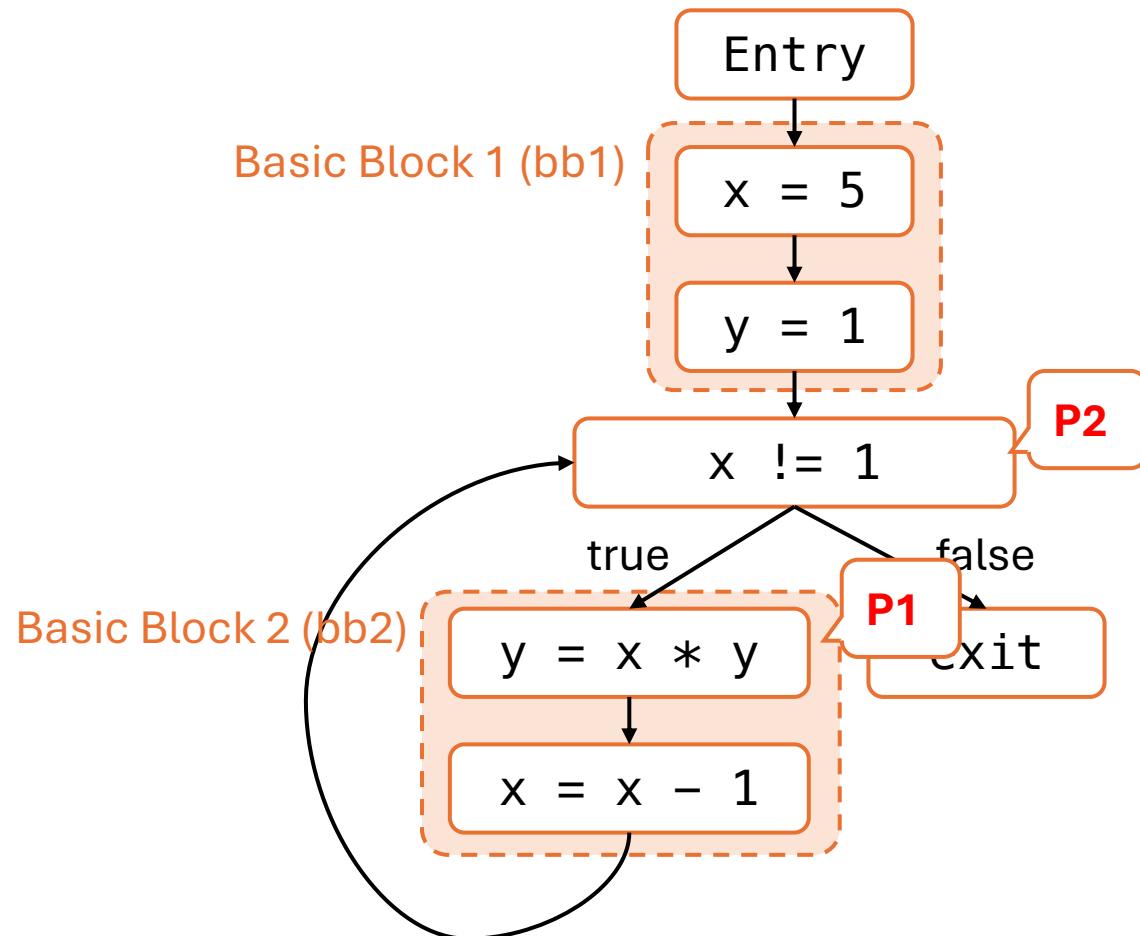
Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

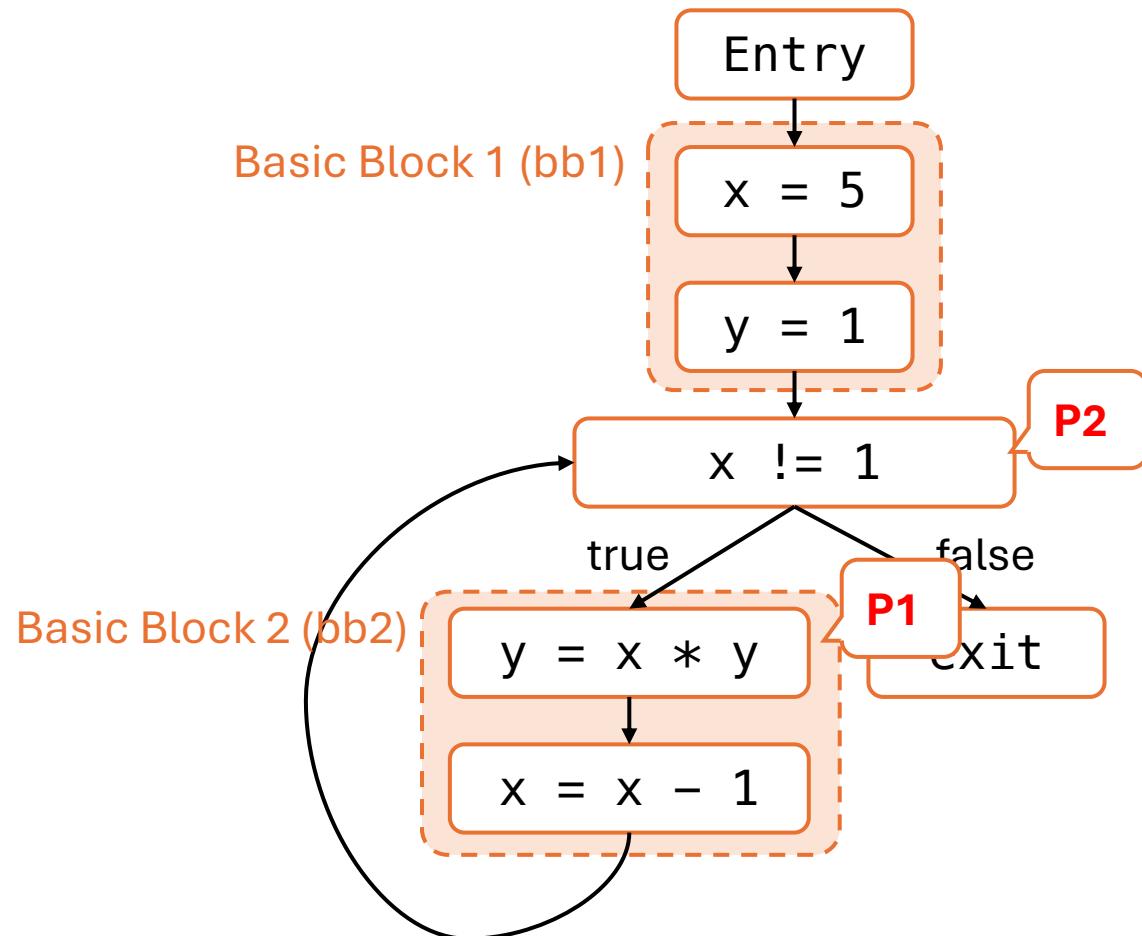
NO

Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

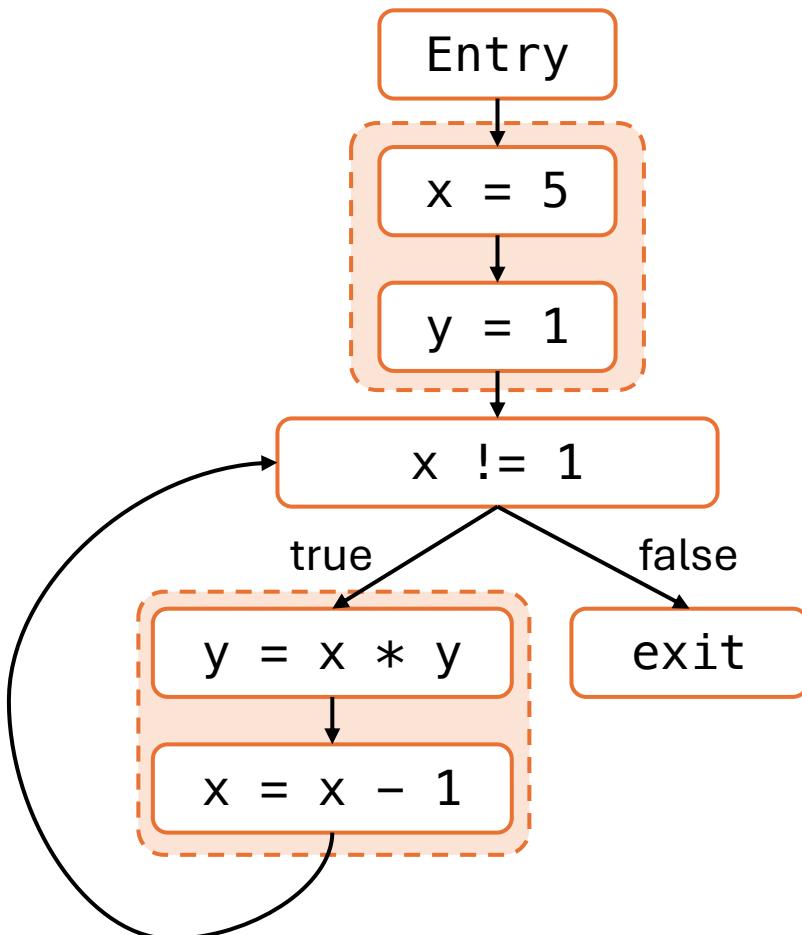
Reaching Definitions Analysis: Quiz



- Does information defined in **P1** flow to **P2**?

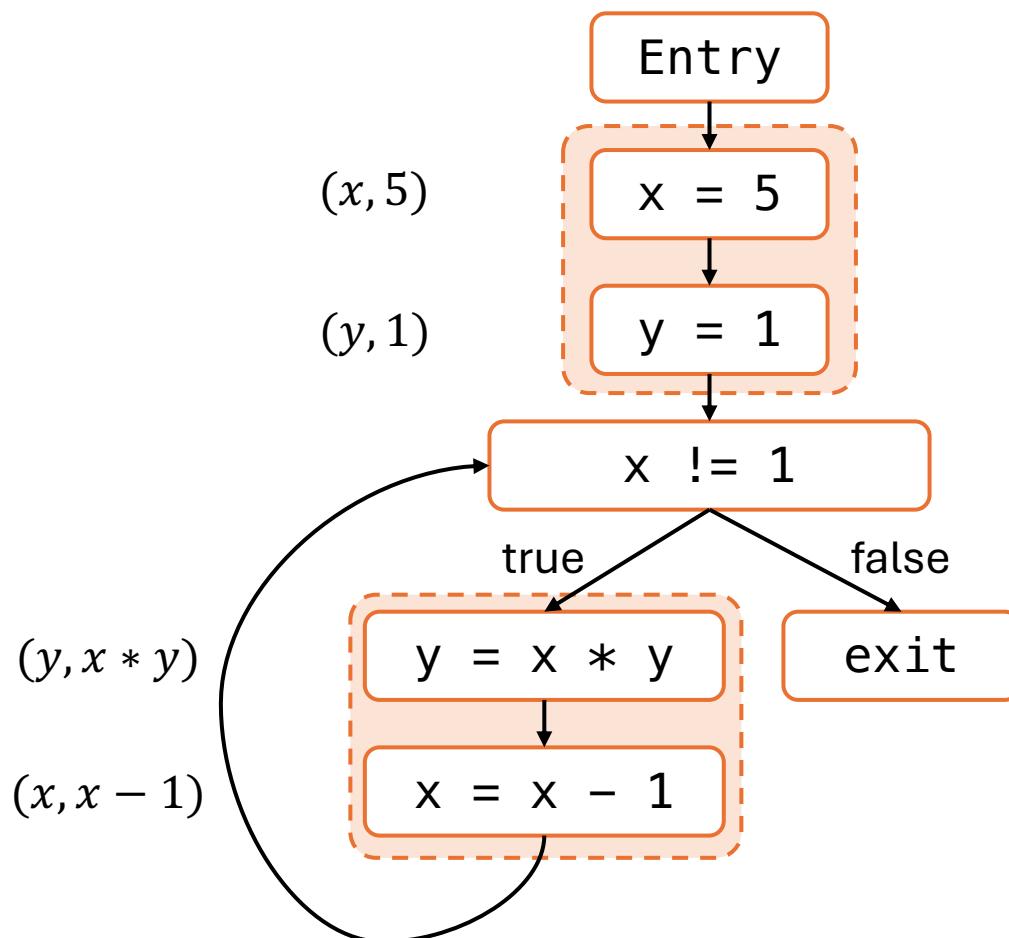
YES

How to Compute Reaching Definitions?



- For each node n , we denote
 - $\text{IN}(n)$ = nodes which can reach n
 - $\text{OUT}(n)$ = nodes which can go beyond n
- Dataflow analysis computes $\text{IN}(n)$ and $\text{OUT}(n)$ for each node
- Chaotic Iteration:
 - Starting from them being empty
 - Repeat operations to populate them
 - Until the two sets stop changing

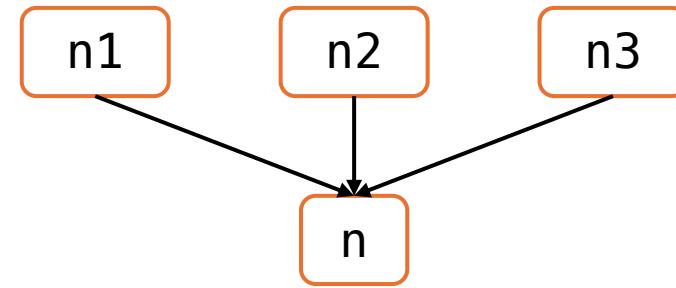
How to Compute Reaching Definitions?



- For each node n , we denote
 - $\text{IN}(n)$ = nodes which can reach n
 - $\text{OUT}(n)$ = nodes which can go beyond n
- Dataflow analysis computes $\text{IN}(n)$ and $\text{OUT}(n)$ for each node
- Chaotic Iteration:
 - Starting from them being empty
 - Repeat operations to populate them
 - Until the two sets stop changing

How to Compute? Flow-in

$$\text{IN}(n) = \bigcup_{m \in \text{predecessors}(n)} \text{OUT}(m)$$



$$\text{IN}(n) = \text{OUT}(n_1) \cup \text{OUT}(n_2) \cup \text{OUT}(n_3)$$

How to Compute? Flow-out

$$\text{OUT}(n) = (\text{IN}(n) - \text{KILL}(n)) \cup \text{GEN}(n)$$

n:

y = x * y

$$\text{GEN}(n) = \{(y, n)\}$$

$$\text{KILL}(n) = \{(y, 5)\}$$

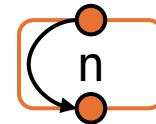
n:

x != 1

$$\text{GEN}(n) = \{\}$$

$$\text{KILL}(n) = \{\}$$

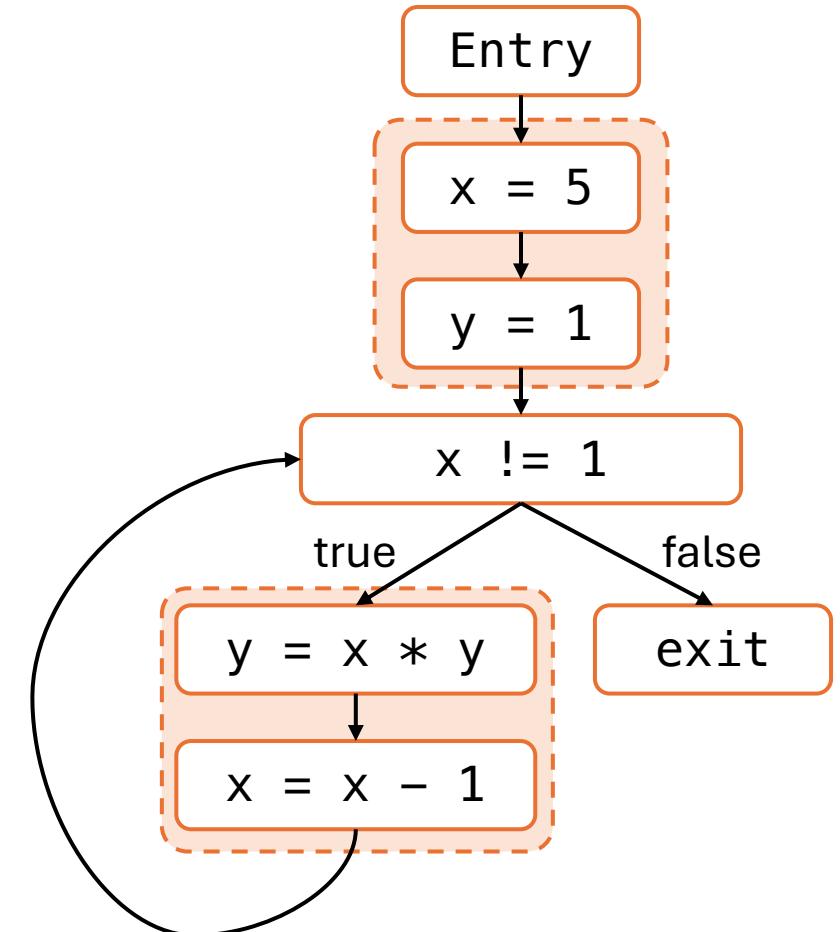
IN(n)



OUT(n)

Chaotic Iteration for Dataflow Analysis

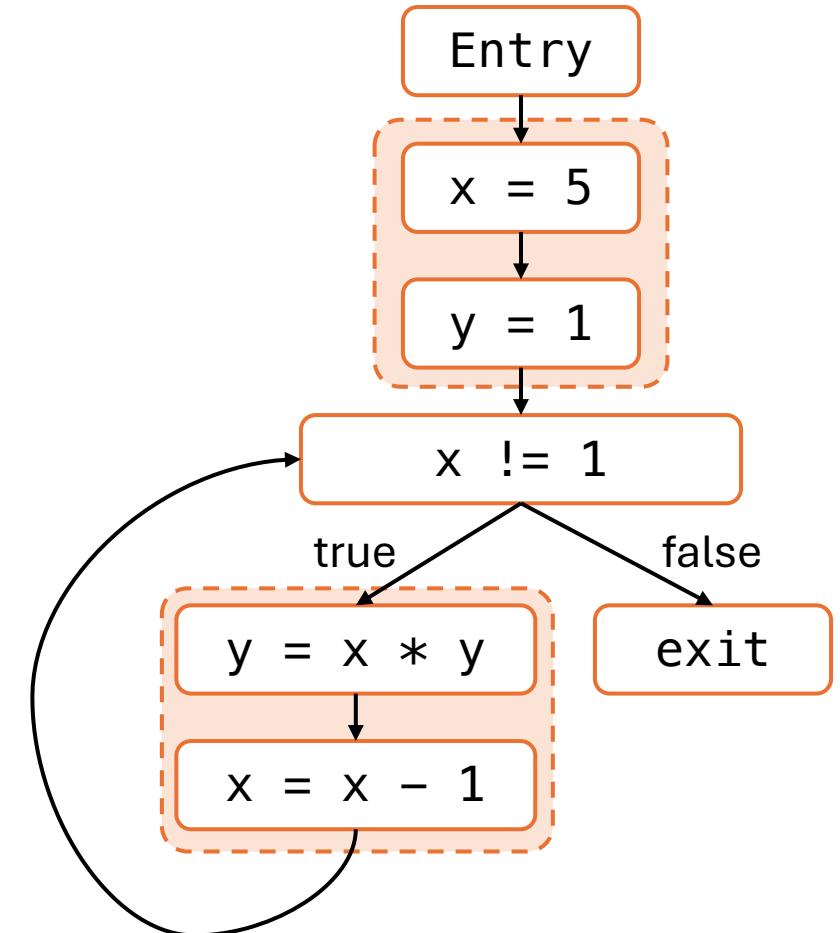
- for (each node n in control-flow graph):
 - $\text{IN}(n) = \{\}$
 - $\text{OUT}(n) = \{\}$
 - $\text{GEN}(n) = \{(v, n)\}$ if n is an assignment else $\{\}$
- repeat:
 - for (each node n):
 - $\text{IN}(n) = \text{flow-in}(n)$
 - $\text{OUT}(n) = \text{flow-out}(n)$
- until all the $\text{IN}(n)$ and $\text{OUT}(n)$ stops changing



Chaotic Iteration for Dataflow Analysis

- for (each node n in control-flow graph):
 - $\text{IN}(n) = \{\}$
 - $\text{OUT}(n) = \{\}$
 - $\text{GEN}(n) = \{(v, n)\}$ if n is an assignment else $\{\}$
- repeat:
 - for (each node n):
 - $\text{IN}(n) = \text{flow-in}(n)$
 - $\text{OUT}(n) = \text{flow-out}(n)$
- until all the $\text{IN}(n)$ and $\text{OUT}(n)$ stops changing

Question: Does it always terminate?

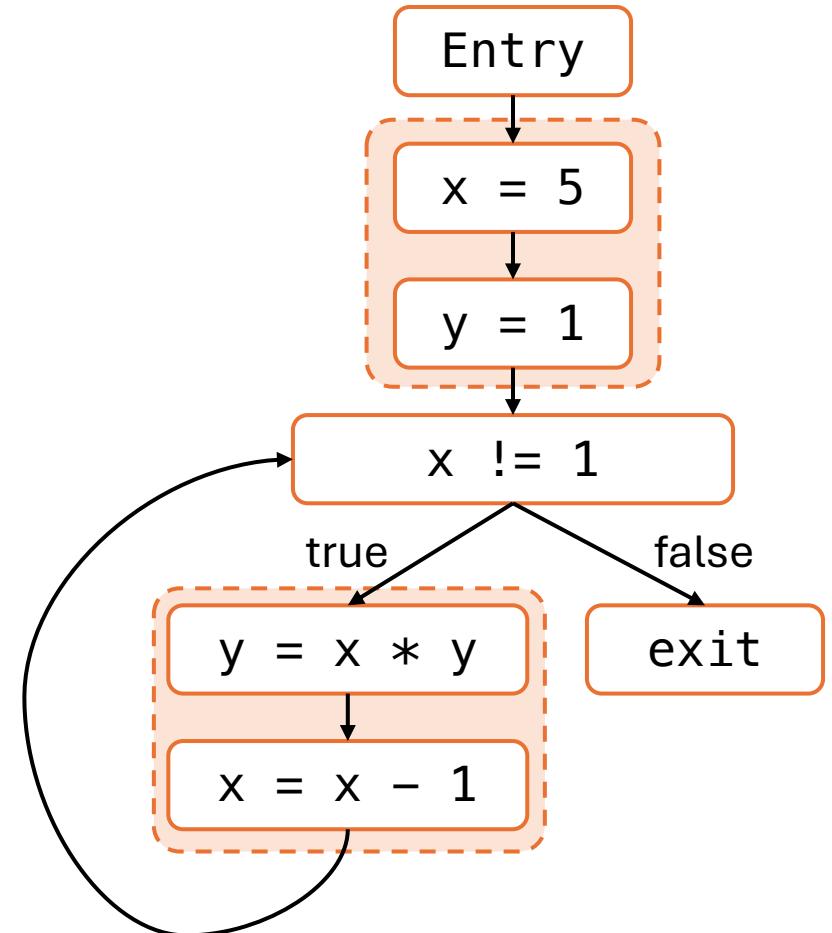


Chaotic Iteration for Dataflow Analysis

- for (each node n in control-flow graph):
 - $\text{IN}(n) = \{\}$
 - $\text{OUT}(n) = \{\}$
 - $\text{GEN}(n) = \{(v, n)\}$ if n is an assignment else $\{\}$
- repeat:
 - for (each node n):
 - $\text{IN}(n) = \text{flow-in}(n)$
 - $\text{OUT}(n) = \text{flow-out}(n)$
- until all the $\text{IN}(n)$ and $\text{OUT}(n)$ stops changing

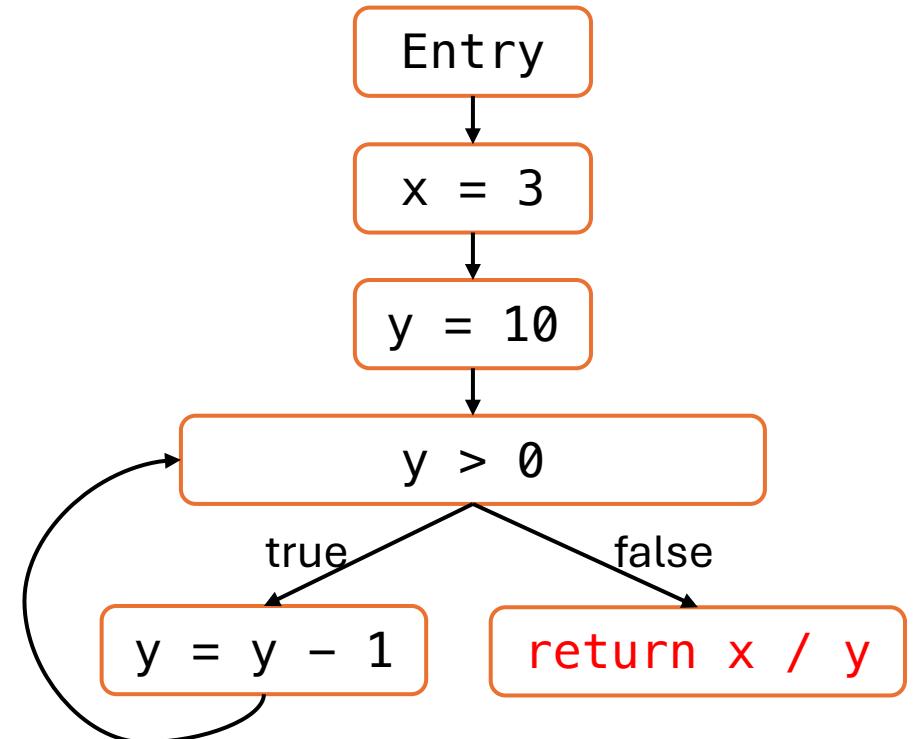
Question: Does it always terminate?

Answer: Yes!



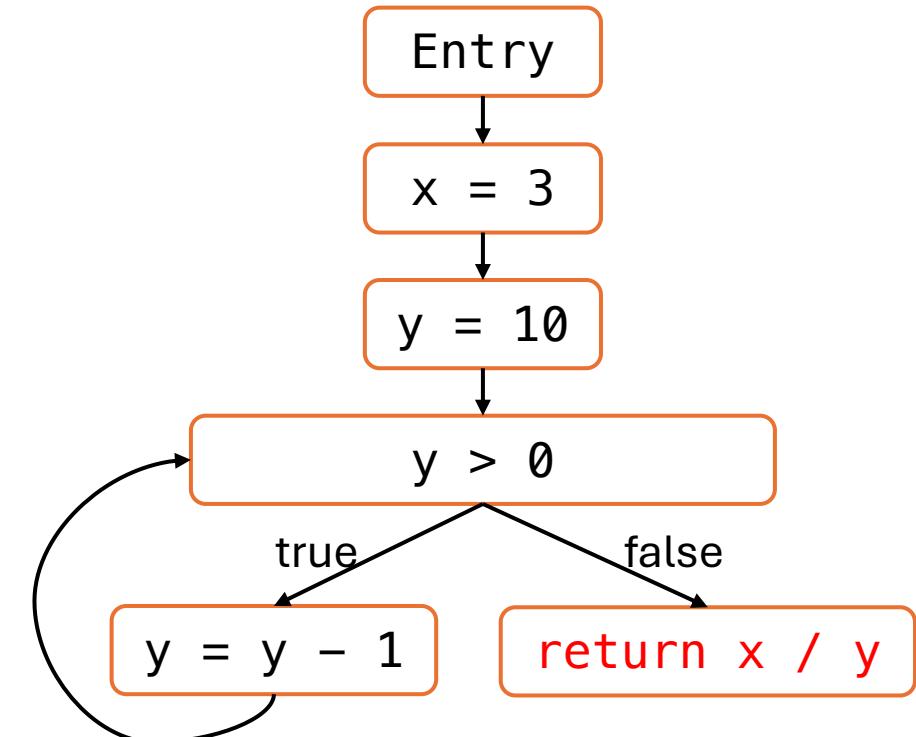
More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?



More Complicated Dataflow Analysis

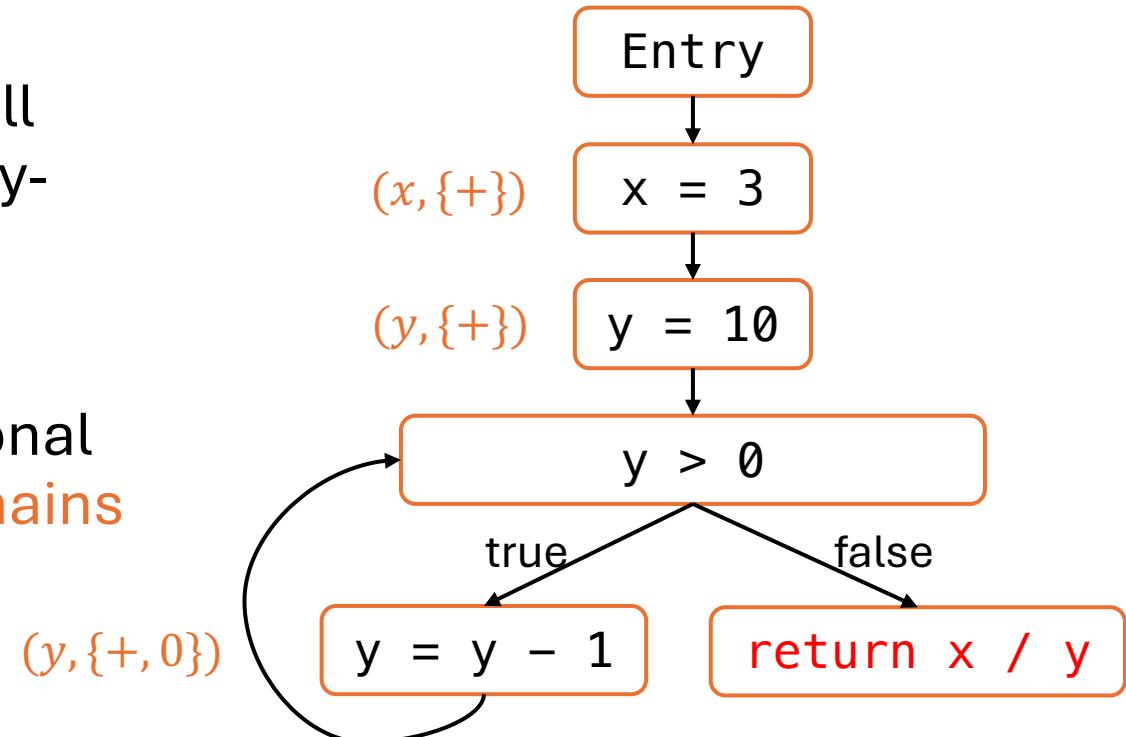
- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**



More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

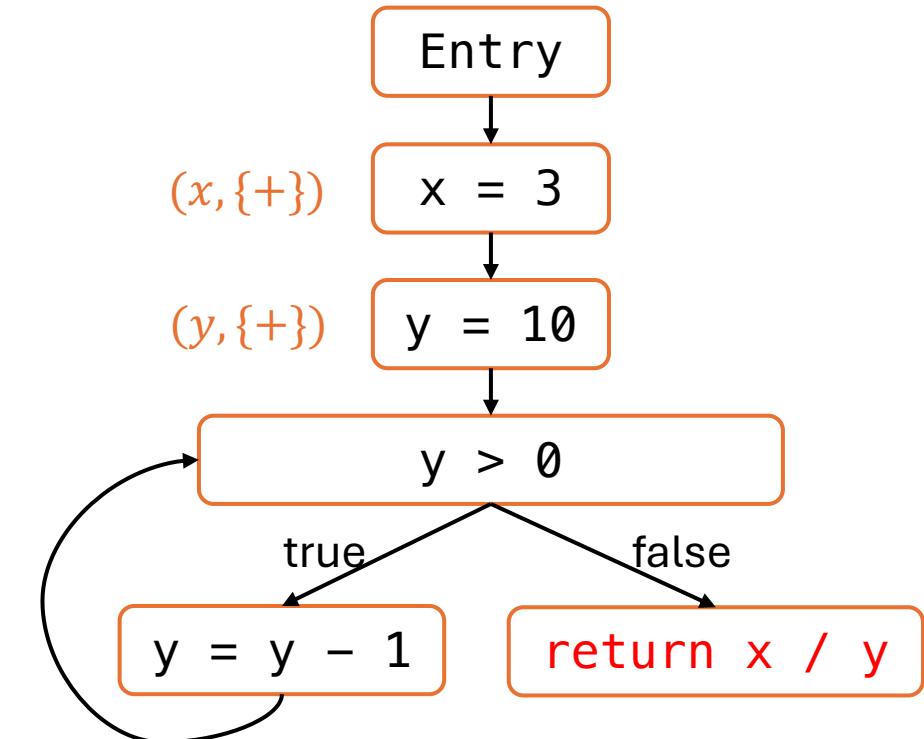
Abstract Domain: $\{+, 0, -\}$



More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

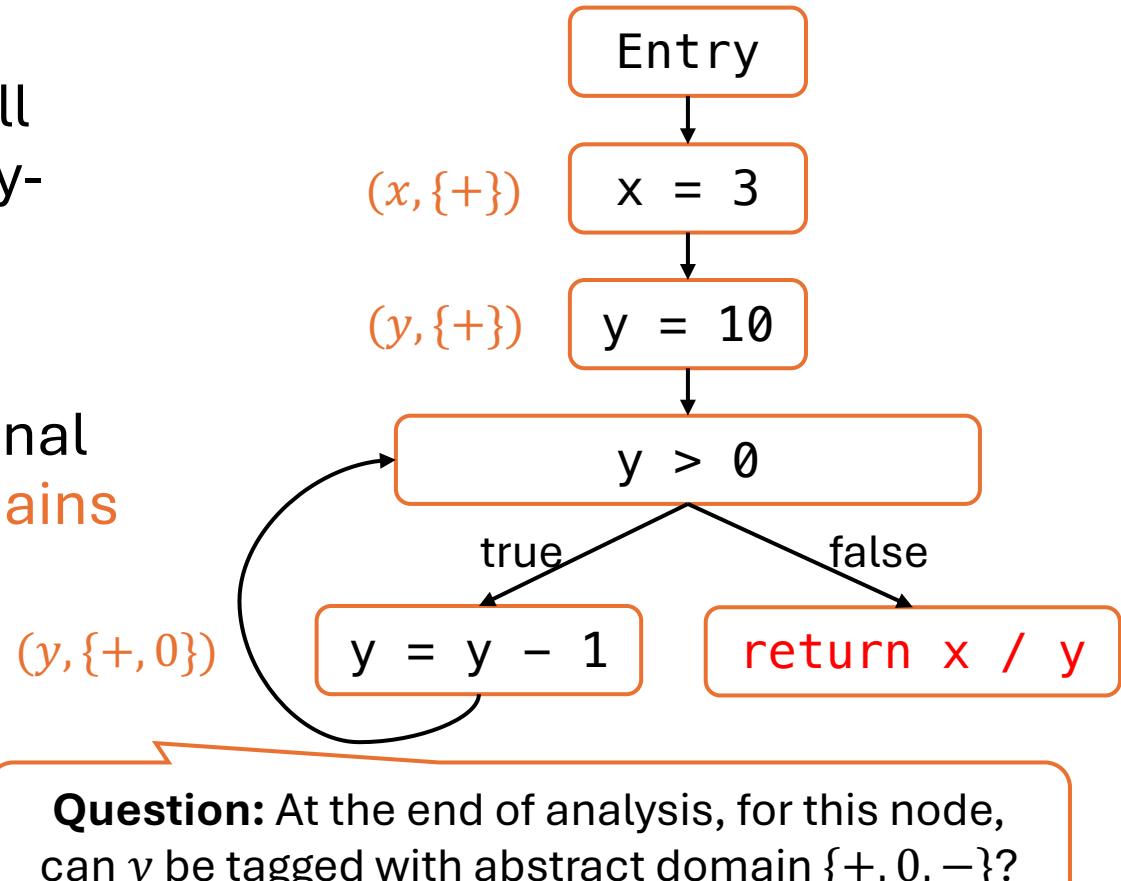
Abstract Domain: $\{+, 0, -\}$



More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

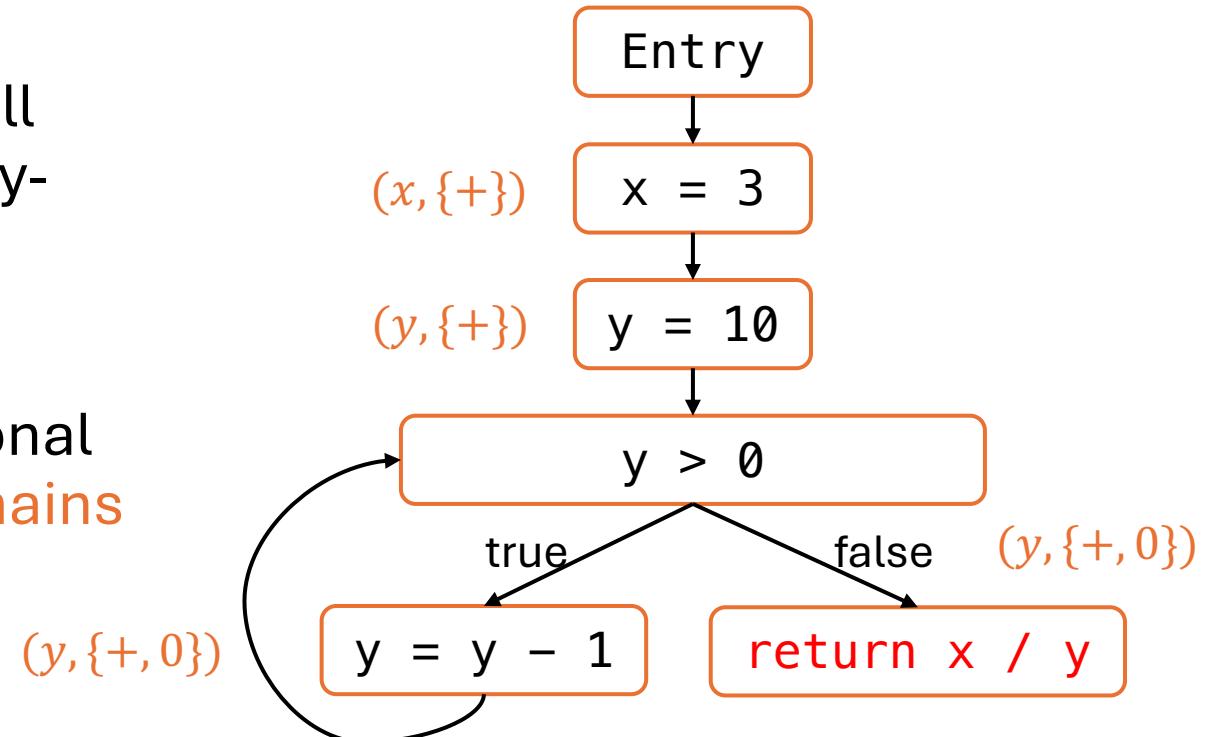
Abstract Domain: $\{+, 0, -\}$



More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

Abstract Domain: $\{+, 0, -\}$



Arithmetic of Abstract Domains

$$\{+\} + \{+\} \rightarrow \{+\}$$

$$\{+\} - \{+\} \rightarrow \{+, 0, -\}$$

Specialization: $\{+\} - 1 \rightarrow \{+, 0\}$

$$\{+\} * \{+\} \rightarrow \{+\}$$

$$\{+\} + \{0\} \rightarrow \{+\}$$

$$\{+\} - \{0\} \rightarrow \{+\}$$

$$\{+\} * \{0\} \rightarrow \{0\}$$

$$\{+\} + \{-\} \rightarrow \{+, 0, -\}$$

Specialization: $\{+\} + (-1) \rightarrow \{+, 0\}$

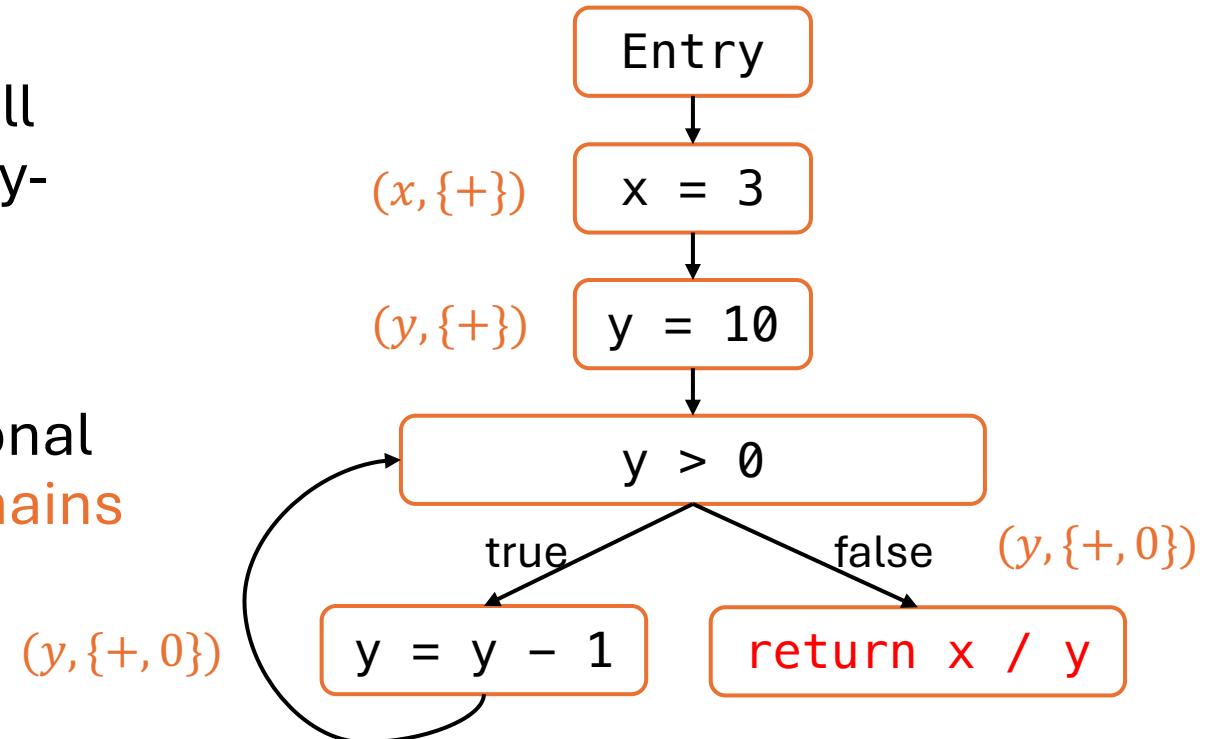
$$\{+\} - \{-\} \rightarrow \{+\}$$

$$\{+\} * \{-\} \rightarrow \{-\}$$

More Complicated Dataflow Analysis

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

Abstract Domain: $\{+, 0, -\}$

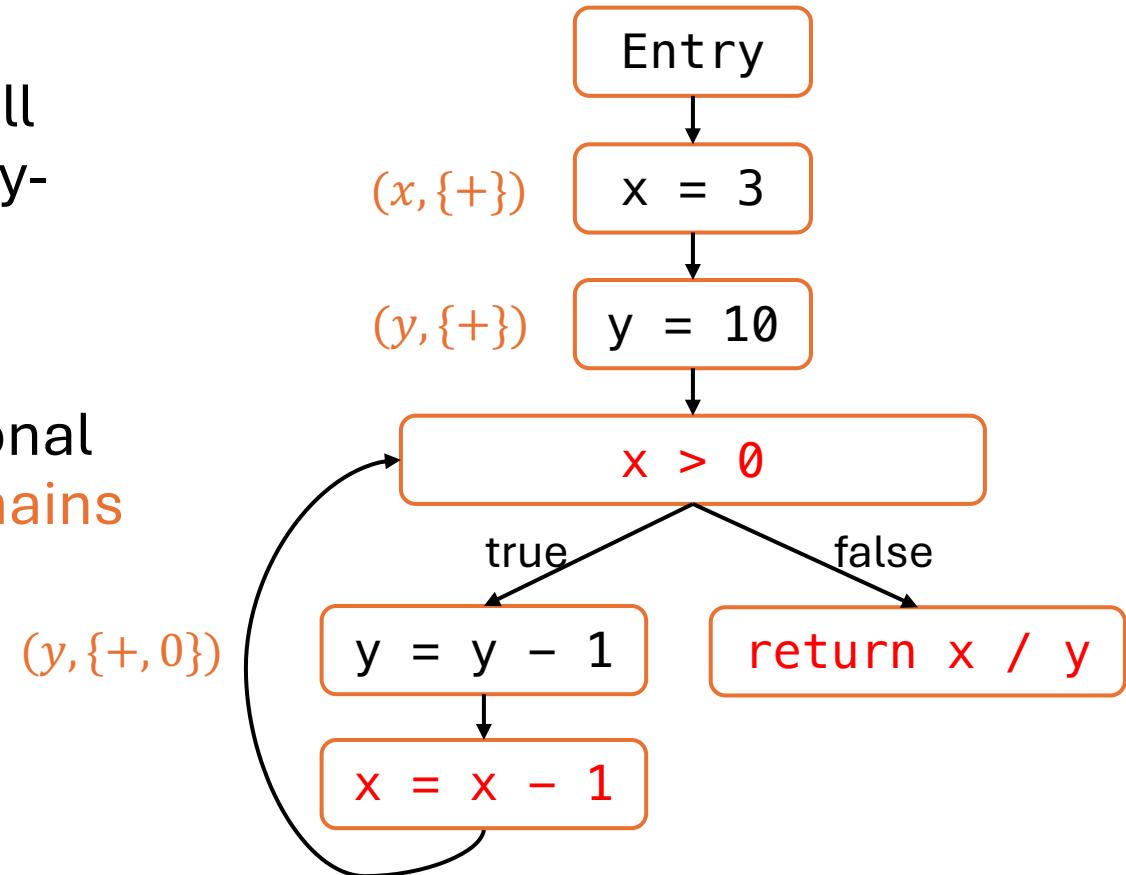


Since y is tagged with abstract domain 0, there **IS** a potential **divide-by-zero vulnerability!**

More Complicated Dataflow Analysis: Sensitivity

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

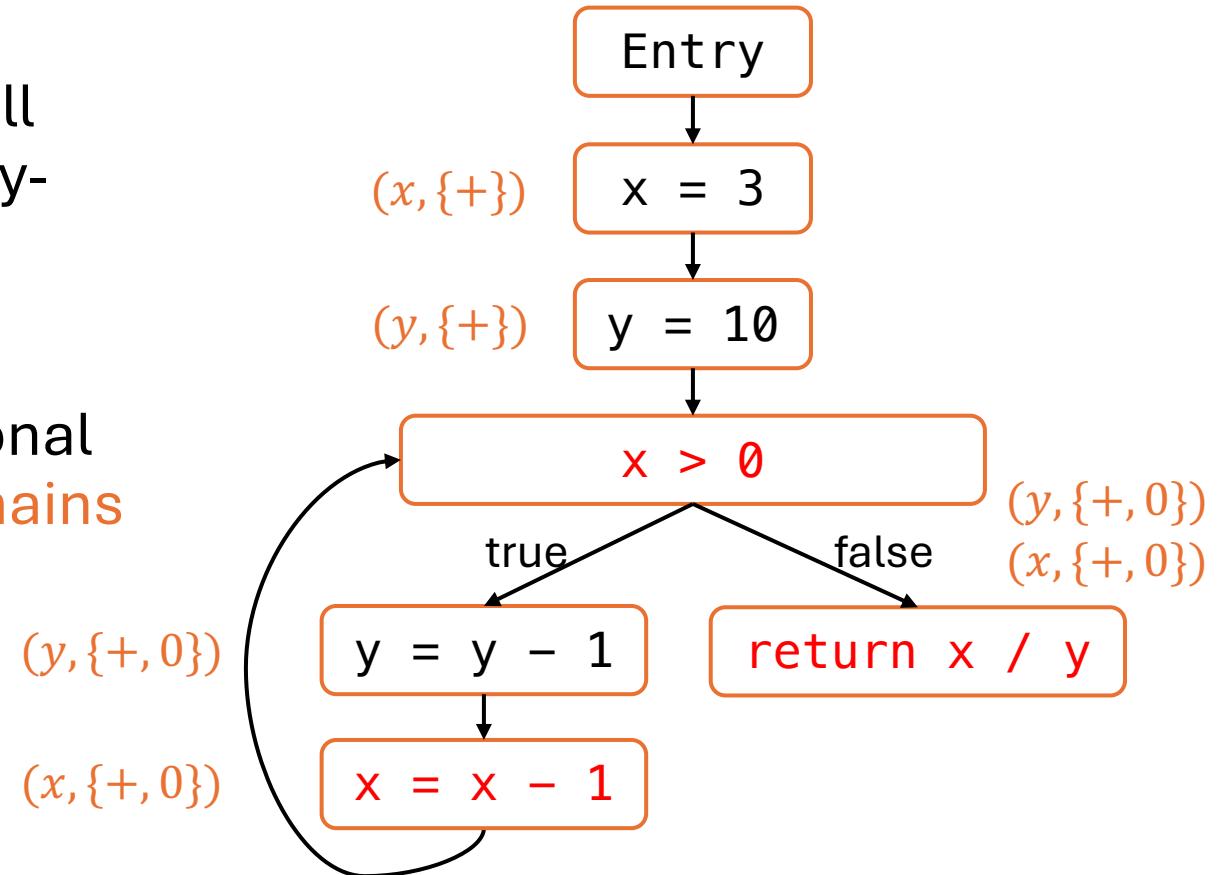
Abstract Domain: $\{+, 0, -\}$



More Complicated Dataflow Analysis: Sensitivity

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

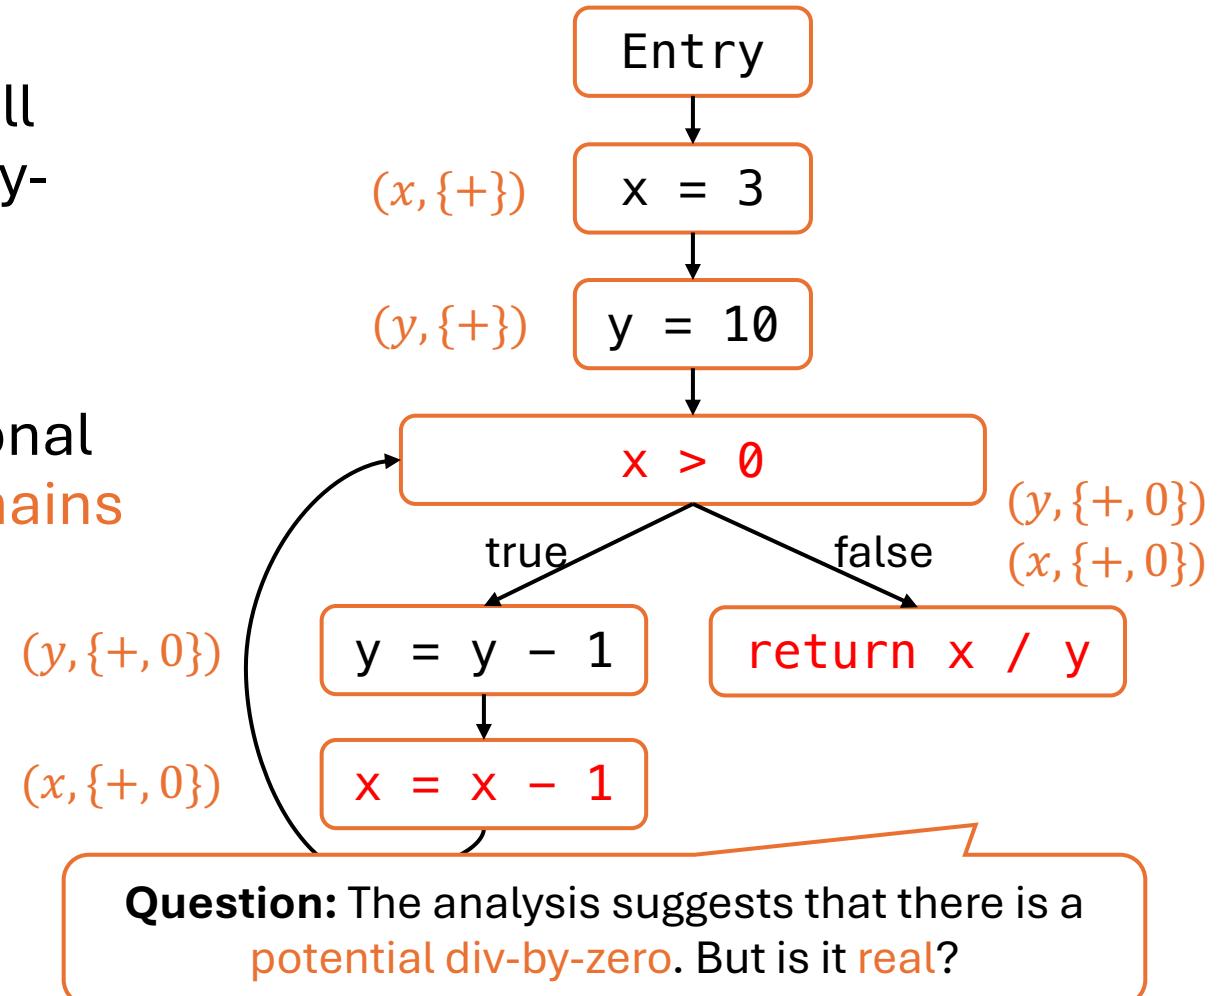
Abstract Domain: $\{+, 0, -\}$



More Complicated Dataflow Analysis: Sensitivity

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero vulnerability?
- Answer:
 - We can tag the nodes with additional information such as **abstract domains**

Abstract Domain: $\{+, 0, -\}$



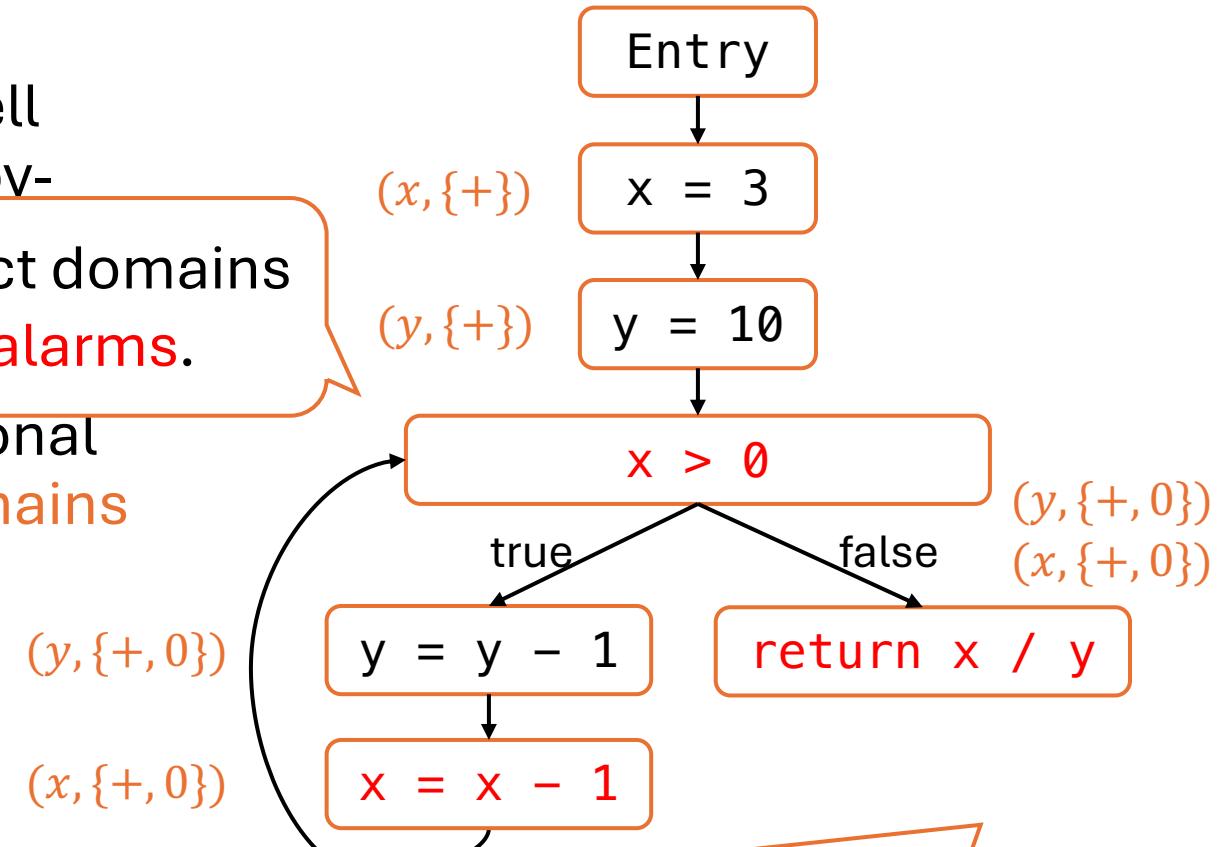
More Complicated Dataflow Analysis: Sensitivity

- Question:
 - With Dataflow analysis, can we tell whether there could be a divide-by-zero?

Dataflow Analysis with $\{+, 0, -\}$ abstract domains reported a **false positive div-by-zero alarms**.

- We can tag the nodes with additional information such as **abstract domains**

Abstract Domain: $\{+, 0, -\}$



Question: The analysis suggests that there is a potential div-by-zero. But is it **real? NO**

Software Testing: Impossible Triangle

Static Analysis:

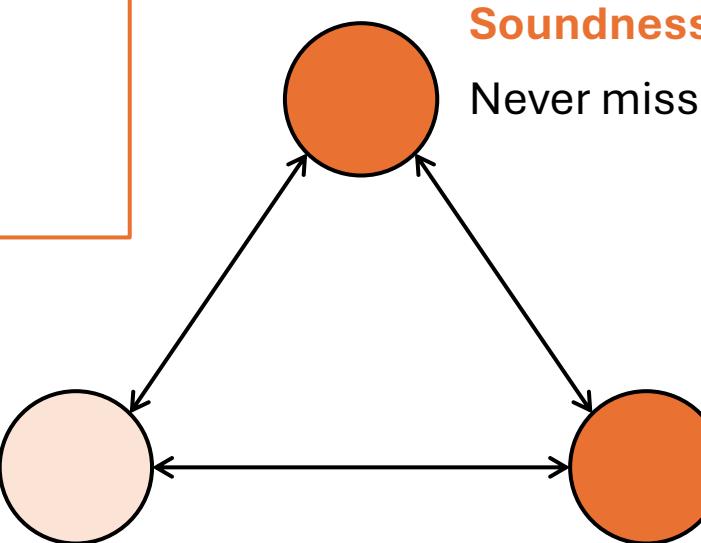
Taint Analysis

- (assuming correct specs)

Abstract Interpretation

- (assuming correct domains)

Completeness
Never raise a false alarm



Termination
Always finish the analysis within finite time

May report an alarm for developer confirmation, and that is rejected

Software Testing: Impossible Triangle

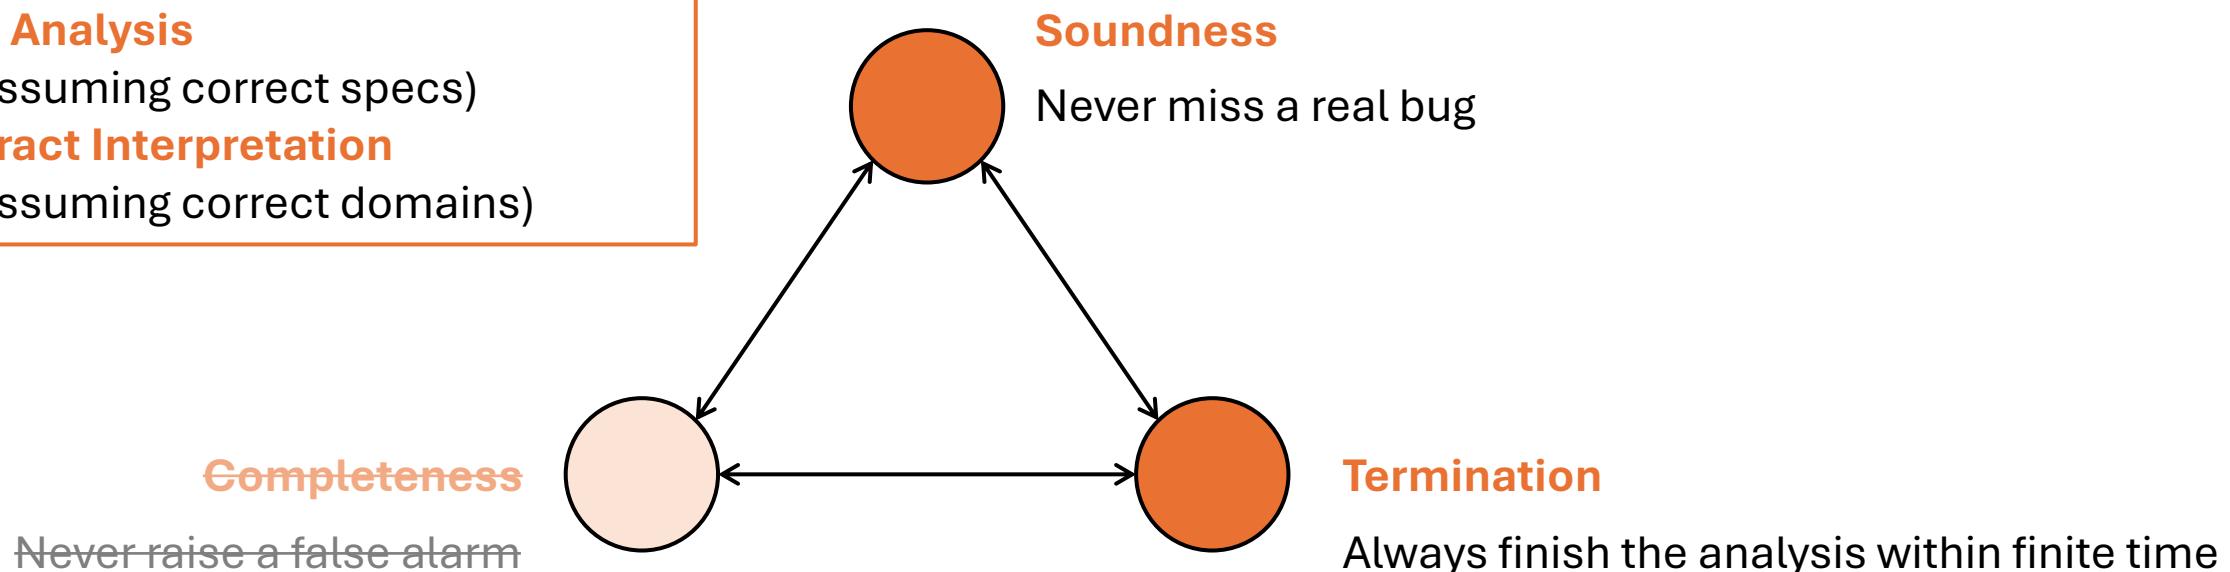
Static Analysis:

Taint Analysis

- (assuming correct specs)

Abstract Interpretation

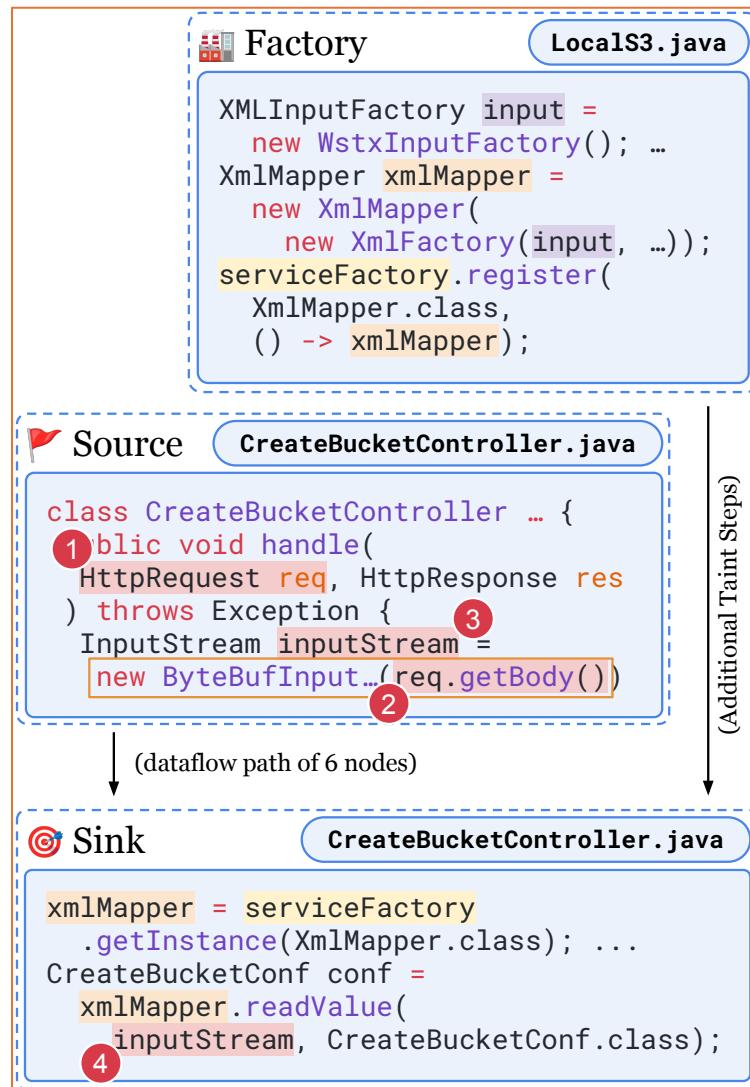
- (assuming correct domains)



Dataflow Analysis & Abstract Domain generalizes to entire space of input, but loses precision

Dataflow Analysis in Real Life: Taint Analysis

Dataflow Analysis in Real Life: Taint Analysis



[GitHub Advisory Database](#) / [GitHub Reviewed](#) / CVE-2025-27136

LocalS3 CreateBucketConfiguration Endpoint XML External Entity (XXE) Injection

Moderate severity GitHub Reviewed Published on Mar 8 in Robothy/local-s3 · Updated on Mar 14

Steps to Reproduce

1. Create an XML document that includes an external entity declaration pointing to the internal target:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE test [ <!ELEMENT xxе SYSTEM "http://internal-web/flag.txt"> ]>
<CreateBucketConfiguration>
    <LocationConstraint>&xxе;</LocationConstraint>
</CreateBucketConfiguration>
```

2. Send a PUT request to create a new bucket with this configuration:

```
curl -X PUT http://app/test-bucket-2 -H 'Content-Type: application/xml' -d @payload.xml
```

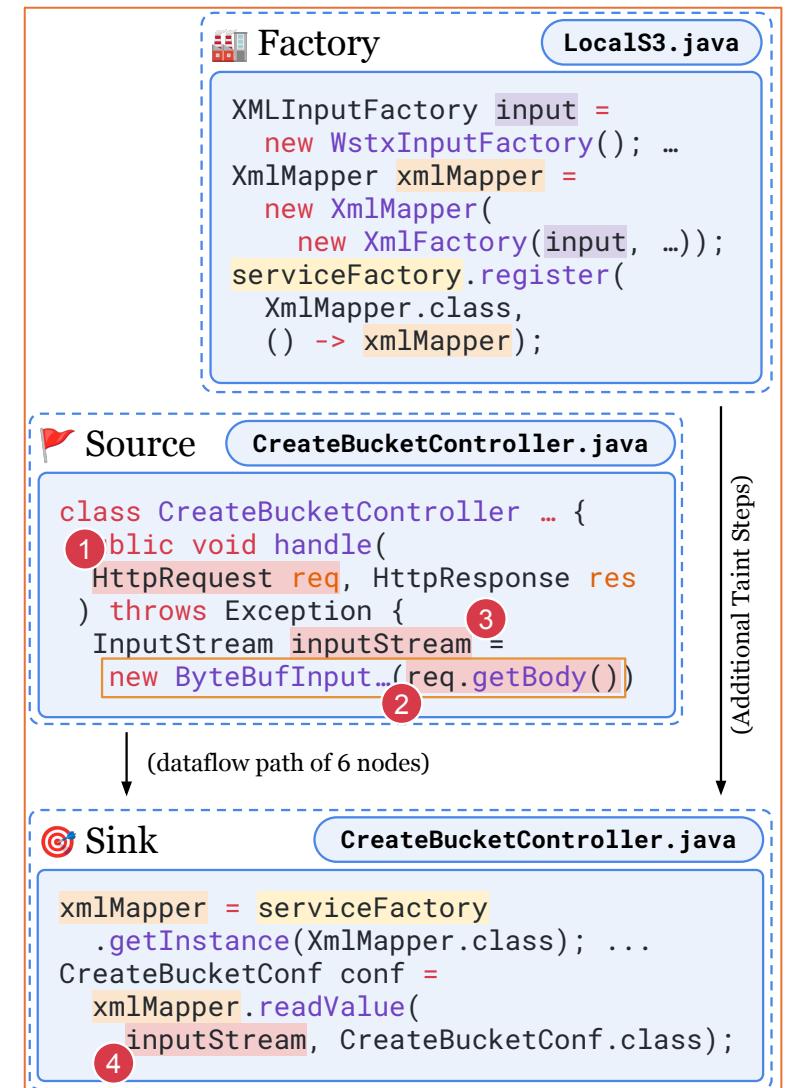
3. Retrieve the bucket location to see the resolved entity content:

```
curl http://app/test-bucket-2/?location
```

When these steps are executed, the server processes the XML, resolves the external entity by making a request to the internal URL, and includes the response in the bucket's location constraint. The attacker can then retrieve this information through the bucket location endpoint.

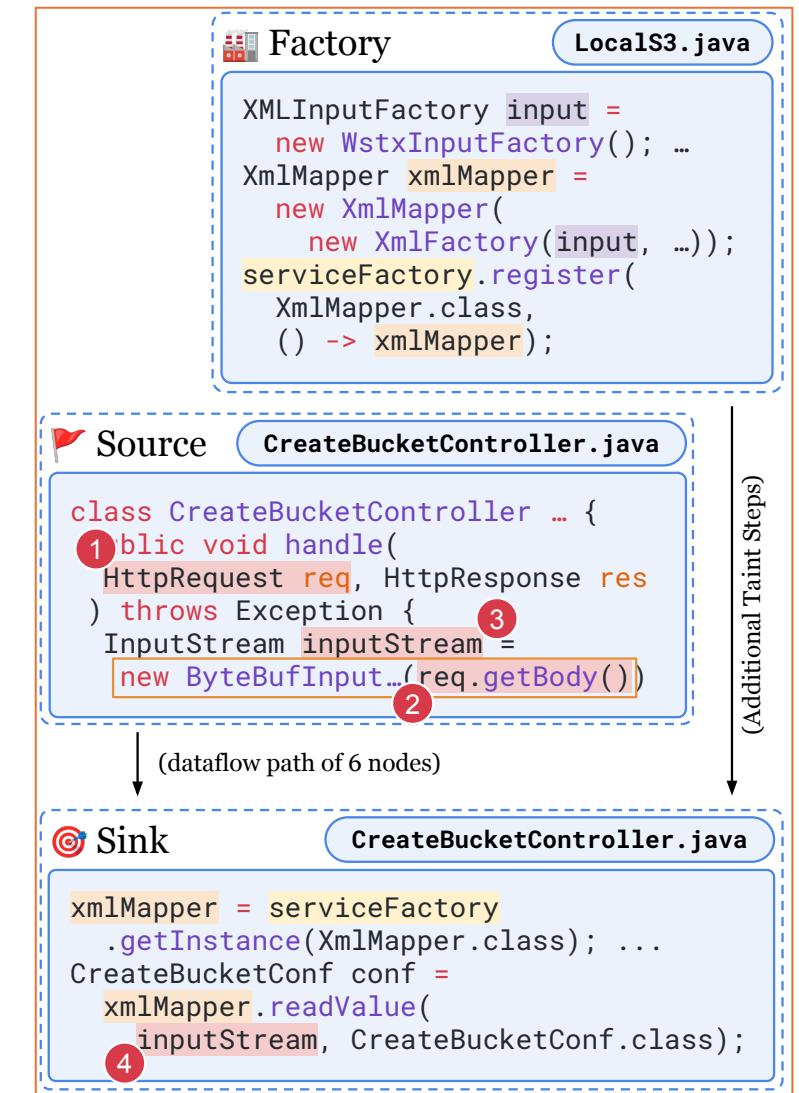
Dataflow Analysis in Real Life: Taint Analysis

- Goal:
 - See how a program point with “tainted information” may flow through the code
 - How is it “Manifested” into a real vulnerability



Dataflow Analysis in Real Life: Taint Analysis

- Goal:
 - See how a program point with “tainted information” may flow through the code
 - How is it “Manifested” into a real vulnerability
- Taint Specification:
 - **Source:** the `HttpRequest` parameter to a `handle` function of a Java Servlet
 - **Sink:** the `readValue` function of `XmlMapper`



Dataflow Analysis in Real Life: Taint Analysis

IRIS: LLM-ASSISTED STATIC ANALYSIS FOR DETECTING SECURITY VULNERABILITIES

Ziyang Li

University of Pennsylvania

liby99@cis.upenn.edu

Saikat Dutta

Cornell University

saikatd@cornell.edu

Mayur Naik

University of Pennsylvania

mhnaik@cis.upenn.edu

Dataflow Analysis in Real Life: Taint Analysis

IRIS · LLM-ASSISTED STATIC ANALYSIS FOR

D

z cronutils/validation/CronValidator.java

```
@override  
public boolean isValid(  
    String value, ConstraintValidatorContext context) {  
    try {  
        2 cronParser.parse(value).valida(5); // ...  
    } catch (IllegalArgumentException e) {  
        context  
            .buildConstraintViolationWithTemplate(e.getMessage())  
            .addConstraintViolation(); // ...  
    }  
}
```

Z U

1 cronutils/parser/CronParser.java

```
/** Parse string with cron expression. ... */  
public Cron parse(final String expression) {  
    try { /* ... */ } catch {  
        throw new IllegalArgumentException(  
            4 String.format("Failed to parse '%s'. %s",  
                expression, e.getMessage()), e);  
    }  
}
```

Figure 2: An example of Code Injection (CWE-94) vulnerability found in cron-utils (CVE-2021-41269) that CodeQL fails to detect. We number the program points of the vulnerable path.

Dataflow Analysis in Real Life: Taint Analysis

IRIS: LLM-ASSISTED STATIC ANALYSIS FOR DETECT

Ziyang Li
University of
liby99@ci

The figure shows two Java code snippets side-by-side:

CronValidator.java:

```
@Override
public boolean isValid(
    String value, ConstraintValidatorContext context) {
    try {
        cronParser.parse(value).valida⑤(); // ...
    } catch (IllegalArgumentException e) {
        context
            .buildConstraintViolationWithTemplate(e.getMessage())
            .addConstraintViolation(); // ...
    }
}
```

CronParser.java:

```
/** Parse string with cron expression. ... */
public Cron parse(final String expression) {
    try { /* ... */ } catch {
        throw new IllegalArgumentException(
            ④String.format("Failed to parse '%s'. %s",
                expression, e.getMessage()), e);
    }
}
```

Annotations (red circles numbered 1-6):

- ①: Source point in the `isValid` method's parameter `value`.
- ②: Intermediate point where the value is passed to `cronParser.parse`.
- ③: Sink point in the `parse` method's return statement.
- ④: Intermediate point where the error message is formatted.
- ⑤: Intermediate point where the validation result is returned.
- ⑥: Intermediate point where the constraint violation is built.

Figure 2: An example of Code Injection (CWE-94) vulnerability found in cron-utils (CVE-2021-41269) that CodeQL fails to detect. We number the program points of the vulnerable path.

Taint Specification:

Source: the `value` parameter to the `isValid` function of a `CronValidator`

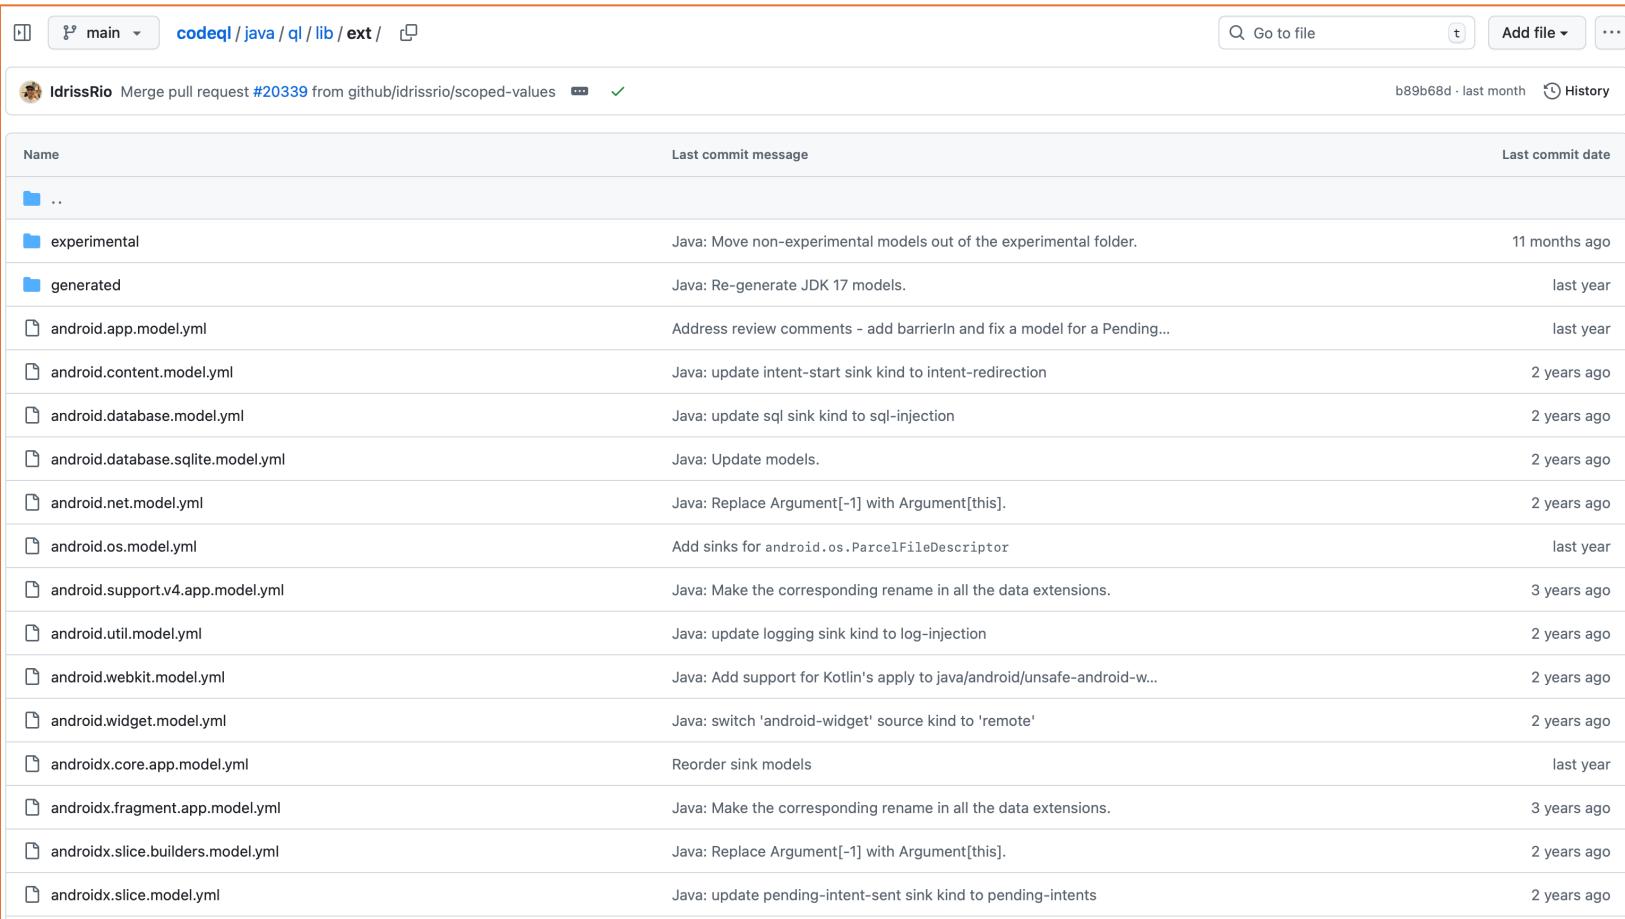
Sink: the `buildConstraintViolationWithTemplate` function of `ConstraintValidatorContext`

Dataflow Analysis in Real Life: Taint Analysis

- Where are those taint specifications coming from?

Dataflow Analysis in Real Life: Taint Analysis

- Where are those taint specifications coming from?

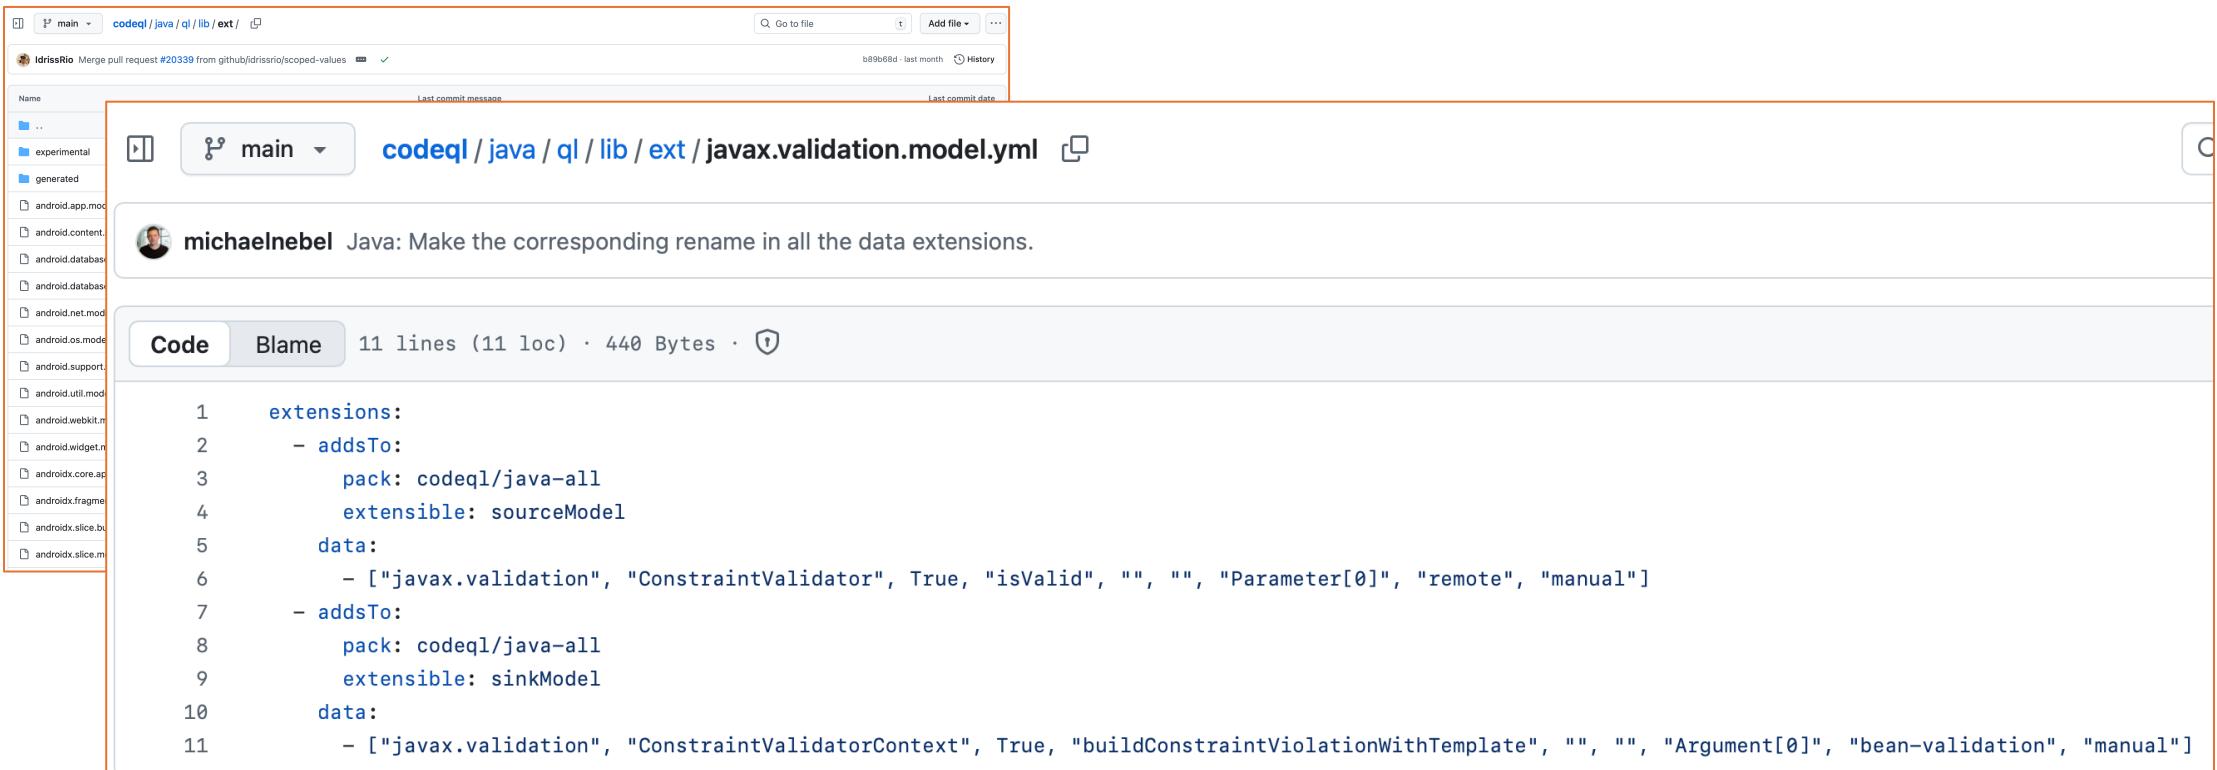


The screenshot shows a GitHub pull request history for a repository named 'codeql/java/ql/lib/ext'. The pull request is titled 'Merge pull request #20339 from github/idrissrio/scoped-values' and was merged by 'IdrissRio' last month. The history lists 21 commits, all of which are Java-related and involve moving or generating model files. The commits are dated from 'last month' to '3 years ago'. The commits are:

Name	Last commit message	Last commit date
...		
experimental	Java: Move non-experimental models out of the experimental folder.	11 months ago
generated	Java: Re-generate JDK 17 models.	last year
android.app.model.yml	Address review comments - add barrierIn and fix a model for a Pending...	last year
android.content.model.yml	Java: update intent-start sink kind to intent-redirection	2 years ago
android.database.model.yml	Java: update sql sink kind to sql-injection	2 years ago
android.database.sqlite.model.yml	Java: Update models.	2 years ago
android.net.model.yml	Java: Replace Argument[-1] with Argument[this].	2 years ago
android.os.model.yml	Add sinks for android.os.ParcelFileDescriptor	last year
android.support.v4.app.model.yml	Java: Make the corresponding rename in all the data extensions.	3 years ago
android.util.model.yml	Java: update logging sink kind to log-injection	2 years ago
android.webkit.model.yml	Java: Add support for Kotlin's apply to java/android/unsafe-android-w...	2 years ago
android.widget.model.yml	Java: switch 'android-widget' source kind to 'remote'	2 years ago
androidx.core.app.model.yml	Reorder sink models	last year
androidx.fragment.app.model.yml	Java: Make the corresponding rename in all the data extensions.	3 years ago
androidx.slice.builders.model.yml	Java: Replace Argument[-1] with Argument[this].	2 years ago
androidx.slice.model.yml	Java: update pending-intent-sent sink kind to pending-intents	2 years ago

Dataflow Analysis in Real Life: Taint Analysis

- Where are those taint specifications coming from?



The screenshot shows a GitHub pull request interface. The top bar indicates the repository is 'codeql/java/ql/lib/ext/' and the branch is 'main'. A commit by 'michaelnebel' is shown, with the message: "Java: Make the corresponding rename in all the data extensions." The file 'javax.validation.model.yml' is displayed, containing YAML code for data extensions. The code defines 'extensions' with 'addsTo' blocks for 'ConstraintValidator' and 'ConstraintValidatorContext' with specific parameters.

```
extensions:
  - addsTo:
      pack: codeql/java-all
      extensible: sourceModel
    data:
      - ["javax.validation", "ConstraintValidator", True, "isValid", "", "", "Parameter[0]", "remote", "manual"]
  - addsTo:
      pack: codeql/java-all
      extensible: sinkModel
    data:
      - ["javax.validation", "ConstraintValidatorContext", True, "buildConstraintViolationWithTemplate", "", "", "Argument[0]", "bean-validation", "manual"]
```

The screenshot shows a code editor interface with a pull request and a file diff.

Pull Request: IdrisRio Merge pull request #20339 from github/idrisrio/scoped-values

File Diff: `java.validation.model.yml`

Commit Message: Java: Make the corresponding rename in all the data extensions.

Code Blame: main · [Code](#) · [Blame](#) · 164 lines (164 loc) · 16.2 KB · [Protected](#)

Commit Author: Jami Cogswell and Jami Cogswell · Java: use AdditionalTaintStep

Code Content:

```
1  extensions:
2    - addsTo:
3      pack: codeql/java-all
4      extensible: sinkModel
5
6      data:
7        - ["java.io", "File", True, "canExecute", "()", "", "Argument[this]", "path-injection", "manual"]
8        - ["java.io", "File", True, "canRead", "()", "", "Argument[this]", "path-injection", "manual"]
9        - ["java.io", "File", True, "canWrite", "()", "", "Argument[this]", "path-injection", "manual"]
10       - ["java.io", "File", True, "createNewFile", "()", "", "Argument[this]", "path-injection", "ai-manual"]
11       - ["java.io", "File", True, "createTempFile", "(String,String,File)", "", "Argument[2]", "path-injection", "ai-manual"]
12       - ["java.io", "File", True, "delete", "()", "", "Argument[this]", "path-injection", "manual"]
13       - ["java.io", "File", True, "deleteOnExit", "()", "", "Argument[this]", "path-injection", "manual"]
14       - ["java.io", "File", True, "exists", "()", "", "Argument[this]", "path-injection", "manual"]
15       - ["java.io", "File", True, "isDirectory", "()", "", "Argument[this]", "path-injection", "manual"]
16       - ["java.io", "File", True, "isFile", "()", "", "Argument[this]", "path-injection", "manual"]
17       - ["java.io", "File", True, "isHidden", "()", "", "Argument[this]", "path-injection", "manual"]
18       - ["java.io", "File", True, "mkdir", "()", "", "Argument[this]", "path-injection", "manual"]
19       - ["java.io", "File", True, "mkdirs", "()", "", "Argument[this]", "path-injection", "manual"]
```

codeql / java / ql / lib / ext / javax.validation.model.yml

IdrisRio Merge pull request #20339 from github/driisrio/scoped-values

michaelnebel Java: Make the corresponding rename in all the data extensions.

Code Blame

main ext -

1 2 3 4 5 6 7 8 9 10 11

Code Blame

main

codeql / java / ql / lib / ext / java.io.model.yml

Jami Code

CWE Common Weakness Enumeration

A community-developed list of SW & HW weaknesses that can become vulnerabilities

Home > CWE List > CWE-22: Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal') (4.18)

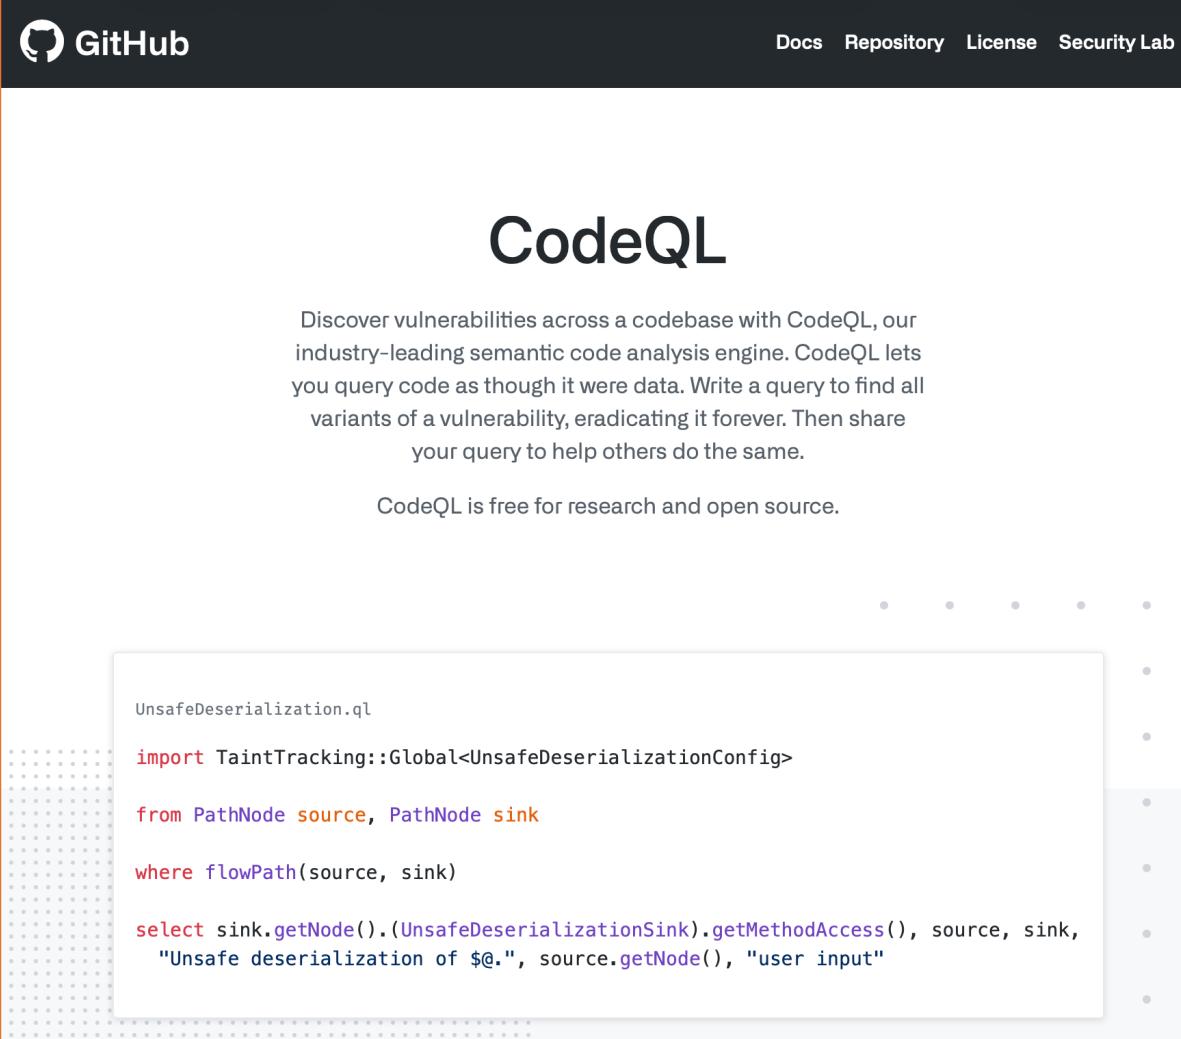
Home | About ▾ | Learn ▾ | Access Content ▾ | Community ▾ | Search ▾

CWE-22: Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal')

Weakness ID: 22
Vulnerability Mapping: ALLOWED
Abstraction: Base

```
data:
  - ["java.io", "File", True, "canExecute", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "canRead", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "canWrite", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "createNewFile", "()", "", "Argument[this]", "path-injection", "ai-manual"]
  - ["java.io", "File", True, "createTempFile", "(String,String,File)", "", "Argument[2]", "path-injection", "ai-manual"]
  - ["java.io", "File", True, "delete", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "deleteOnExit", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "exists", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "isDirectory", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "isFile", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "isHidden", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "mkdir", "()", "", "Argument[this]", "path-injection", "manual"]
  - ["java.io", "File", True, "mkdirs", "()", "", "Argument[this]", "path-injection", "manual"]
```

Dataflow Analysis in Real Life: CodeQL

A screenshot of the GitHub CodeQL landing page. The page has a dark header with the GitHub logo and navigation links for Docs, Repository, License, and Security Lab. The main title "CodeQL" is centered above a descriptive paragraph. Below the paragraph is a code snippet in CodeQL query language. A horizontal ellipsis at the bottom indicates more content.

CodeQL

Discover vulnerabilities across a codebase with CodeQL, our industry-leading semantic code analysis engine. CodeQL lets you query code as though it were data. Write a query to find all variants of a vulnerability, eradicating it forever. Then share your query to help others do the same.

CodeQL is free for research and open source.

```
UnsafeDeserialization.ql
import TaintTracking::Global<UnsafeDeserializationConfig>

from PathNode source, PathNode sink

where flowPath(source, sink)

select sink.getNode().(UnsafeDeserializationSink).getMethodAccess(), source, sink,
    "Unsafe deserialization of $@", source.getNode(), "user input"
```

Data



```
1  /**
2   * @name Uncontrolled data used in path expression
3   * @description Accessing paths influenced by users can allow an attacker to access un
4   * @kind path-problem
5   * @problem.severity error
6   * @security-severity 7.5
7   * @precision high
8   * @id java/path-injection
9   * @tags security
10  *      external/cwe/cwe-022
11  *      external/cwe/cwe-023
12  *      external/cwe/cwe-036
13  *      external/cwe/cwe-073
14  */
15
16 import java
17 import semmle.code.java.security.TaintedPathQuery
18 import TaintedPathFlow::PathGraph
19
20 from TaintedPathFlow::PathNode source, TaintedPathFlow::PathNode sink
21 where TaintedPathFlow::flowPath(source, sink)
22 select sink.getNode(), source, sink, "This path depends on a $@.", source.getNode(),
23       "user-provided value"
```

Dataflow Analysis in Real Life: CodeQL

- Question:
 - When we have a programming language for static analysis...
 - Can we Synthesize Program in this language?

Published as a conference paper at ICLR 2025

IRIS: LLM-ASSISTED STATIC ANALYSIS FOR DETECTING SECURITY VULNERABILITIES

Ziyang Li

University of Pennsylvania

liby99@cis.upenn.edu

Saikat Dutta

Cornell University

saikatd@cornell.edu

Mayur Naik

University of Pennsylvania

mhnaik@cis.upenn.edu

IRIS: LLM-ASSISTED STATIC ANALYSIS FOR

D

Z
U
1

Section 3.2 Candidate Extractor

External API Candidates

package	class	method
java.io	ObjectInputStream	readObject
javax.v...	ConstraintValida...	buildConstra...
		...

Internal Function Param Candidates

class	method	documentation
FieldExpression	accept	"Accept a visitor..."
CronParser	parse	"Parse string..."
		...

Section 3.3 Label Source/Sinks

API Labels

type	target
source	return-value
sink	argument[0]
	...

Function Param Labels

type	target
source	parameter[0]
sink	parameter[0]
	...

Section 3.4 Static Taint Analysis

Data-Flow Path

```

graph TD
    A(( )) --> B(( ))
    B --> C(( ))
    C --> D(( ))
    D --> E(( ))
    E --> F(( ))
    F --> G(( ))
    G --> H(( ))
    H --> I(( ))
    I --> J(( ))
    J --> K(( ))
    K --> L(( ))
    L --> M(( ))
    M --> N(( ))
    N --> O(( ))
    O --> P(( ))
    P --> Q(( ))
    Q --> R(( ))
    R --> S(( ))
    S --> T(( ))
    T --> U(( ))
    U --> V(( ))
    V --> W(( ))
    W --> X(( ))
    X --> Y(( ))
    Y --> Z(( ))
    Z --> A
  
```

boolean isValid(
 String value, ...)
 {
 Cron parse(
 String expr) {
 ...
 throw new IllegalA...
 ...
 context.buildConstra...
 intViolationWith
 Template(e.get...())

Section 3.5 Contextual Analysis

Vulnerability Analysis

Is Vulnerable: The source is user input being parsed and validated, which can lead to an exception with a controlled message. The sink is using the exception message in a context that could potentially be used for execution, satisfying the criteria for a Code Injection vulnerability (CWE-094).

IDE

ZU1

Section 3.2
Candidate Extractor

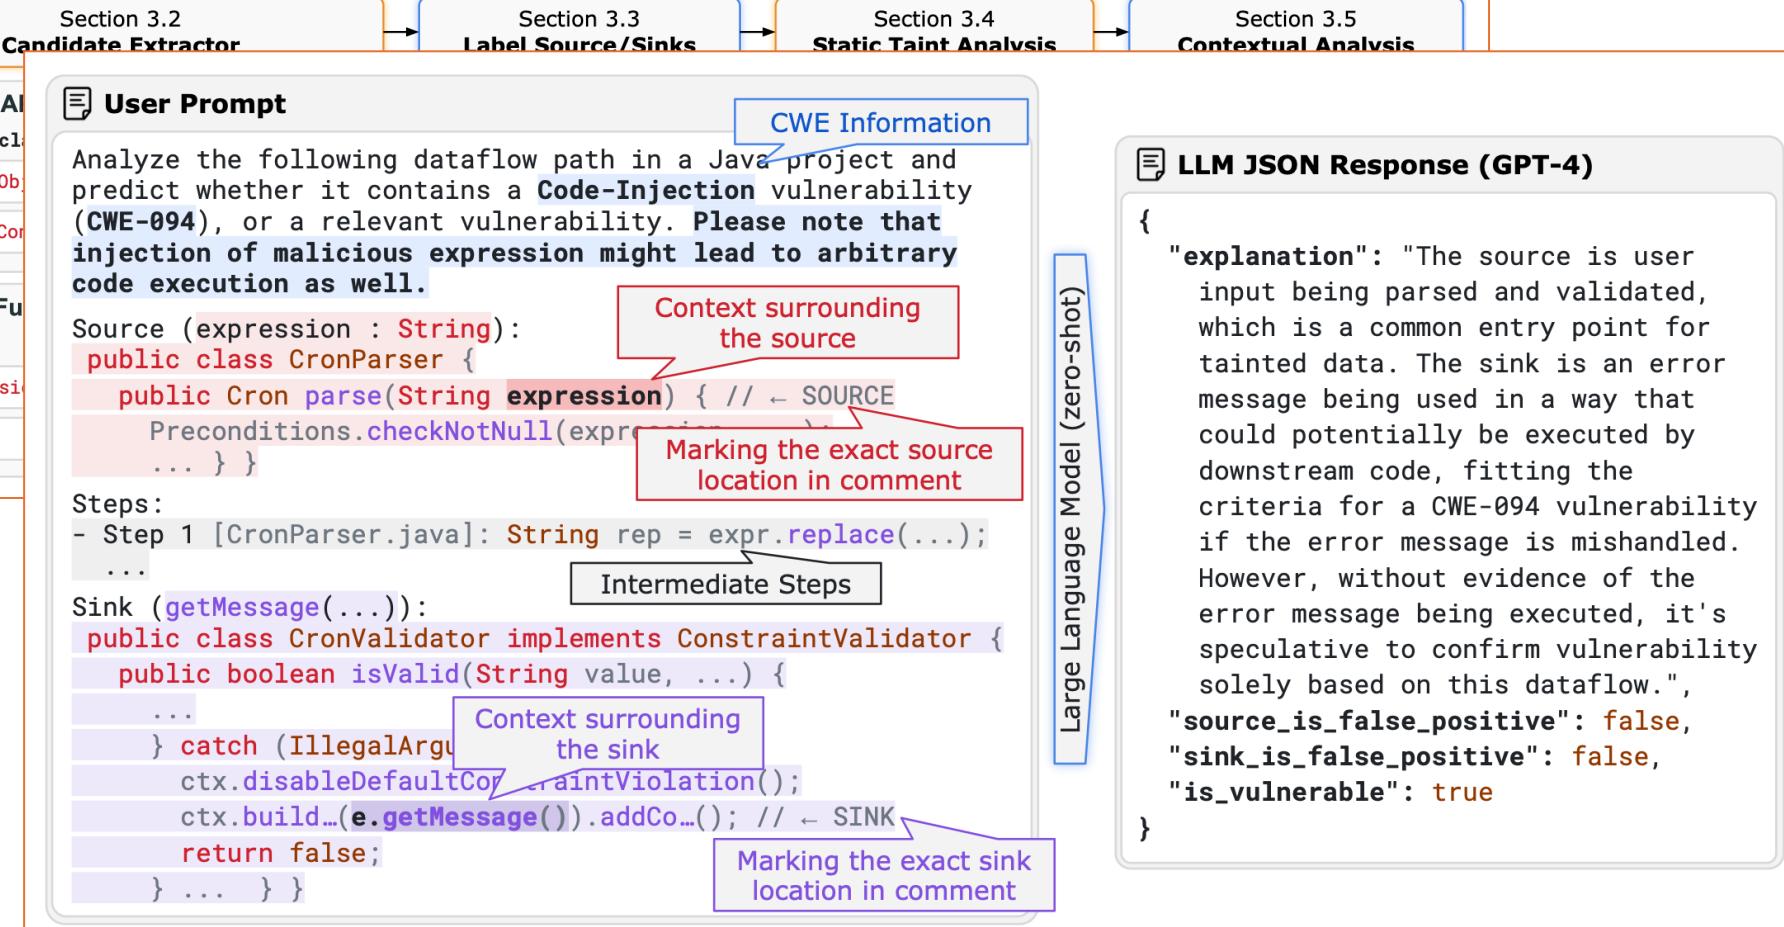
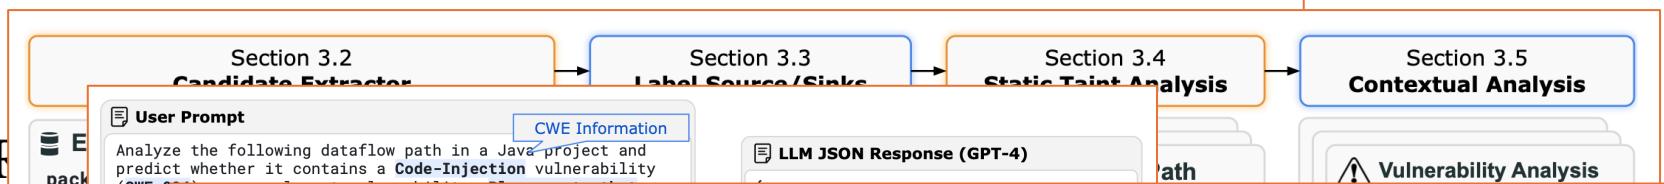


Table 1: Overall performance comparison of CodeQL vs IRIS on Detection Rate (\uparrow), Average FDR (\downarrow), and Average F1 (\uparrow). We present results of IRIS with LLMs including GPT-4 and GPT-3.5, L3 8B and 70B, Q2.5C 32B, G2 27B, and DSC 7B.

	Method	#Detected (/120)	Detection Rate (%)	Avg FDR (%)	Avg F1 Score
	CodeQL	27	22.50	90.03	0.076
IRIS +	GPT-4	55 (\uparrow 28)	45.83 (\uparrow 23.33)	84.82 (\downarrow 5.21)	0.177 (\uparrow 0.101)
	GPT-3.5	47 (\uparrow 20)	39.17 (\uparrow 16.67)	90.42 (\uparrow 0.39)	0.096 (\uparrow 0.020)
	L3 8B	41 (\uparrow 14)	34.17 (\uparrow 11.67)	95.55 (\uparrow 5.52)	0.058 (\downarrow 0.018)
	L3 70B	54 (\uparrow 27)	45.00 (\uparrow 22.50)	90.96 (\uparrow 0.93)	0.113 (\uparrow 0.037)
	Q2.5C 32B	47 (\uparrow 20)	39.17 (\uparrow 16.67)	92.38 (\uparrow 2.35)	0.097 (\uparrow 0.021)
	G2 27B	45 (\uparrow 18)	37.50 (\uparrow 15.00)	91.23 (\uparrow 1.20)	0.100 (\uparrow 0.024)
	DSC 7B	52 (\uparrow 25)	43.33 (\uparrow 20.83)	95.40 (\uparrow 5.37)	0.062 (\downarrow 0.014)



Software Testing: Impossible Triangle

Static Analysis:

Taint Analysis

- (assuming correct specs)

Abstract Interpretation

- (assuming correct domain)

Soundness

Never miss a real bug

Realistically, there is no “completely correct” specifications

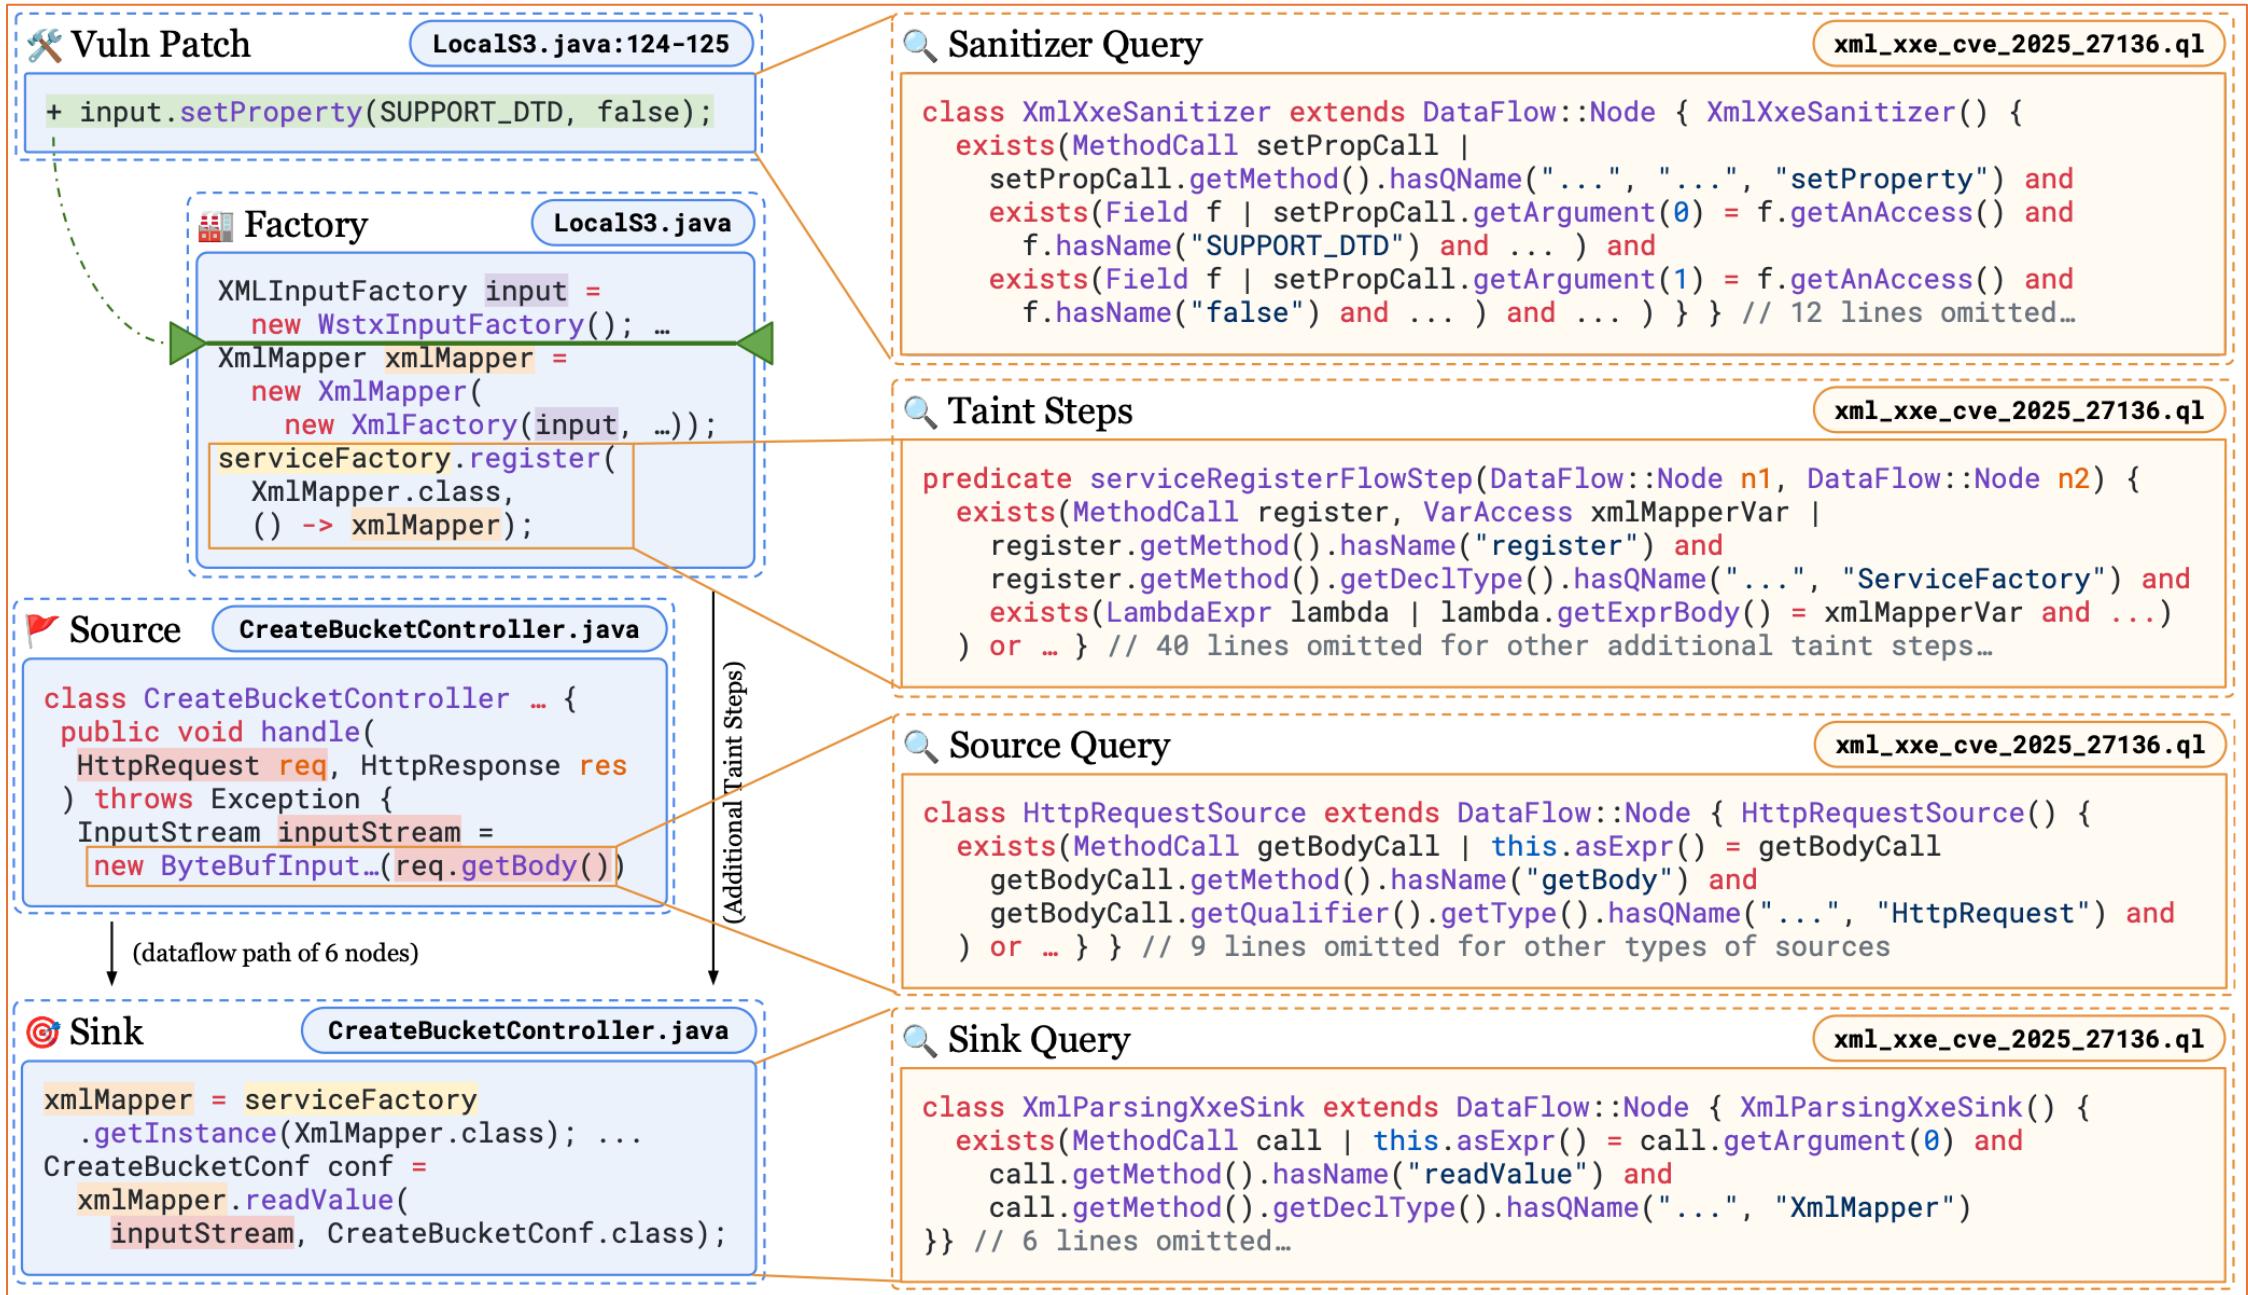
Completeness

Never raise a false alarm

Termination

Always finish the analysis within finite time

Dataflow Analysis & Abstract Domain generalizes to entire space of input, but loses precision



Automated Static Vulnerability Detection via a Holistic Neuro-symbolic Approach

Penghui Li

Columbia University

New York, NY, USA

pl2689@columbia.edu

Songchen Yao

Columbia University

New York, NY, USA

sy2743@columbia.edu

Josef Sarfati Korich

Columbia University

New York, NY, USA

jsk2266@columbia.edu

Changhua Luo

Wuhan University

Wuhan, Hubei, China

chlou2502@whu.edu.cn

Jianjia Yu

Johns Hopkins University

Baltimore, MD, USA

jyu122@jhu.edu

Yinzhi Cao

Johns Hopkins University

Baltimore, MD, USA

yinzhi.cao@jhu.edu

Junfeng Yang

Columbia University

New York, NY, USA

junfeng@cs.columbia.edu

Automated Static Vulnerability Detection via a Holistic Neuro-symbolic Approach

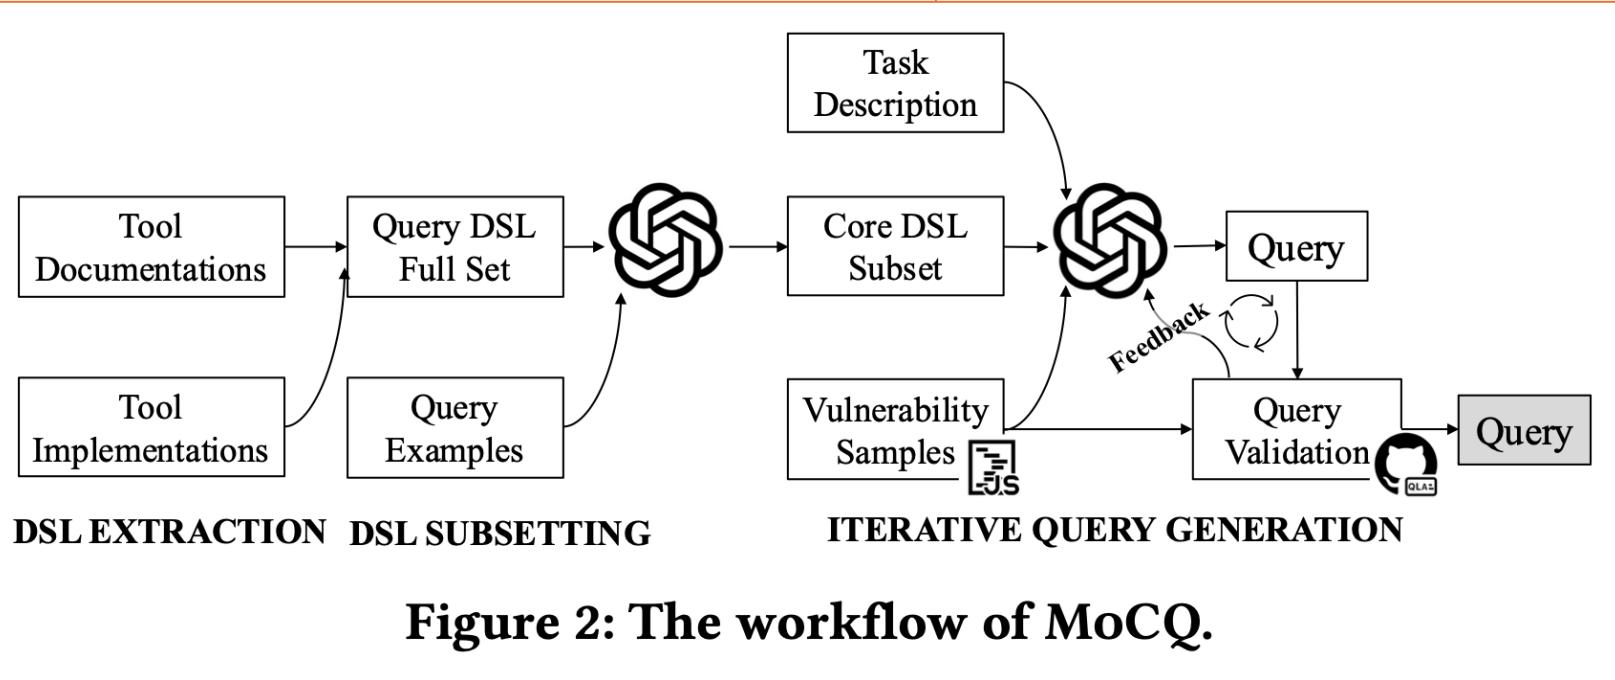
Penghui Li
Columbia University
New York, NY, USA
pl2689@columbia.edu

Songchen Yao
Columbia University
New York, NY, USA

Josef Sarfati Korich
Columbia University
New York, NY, USA

Changhua Luo
Wuhan University
Wuhan, Hubei, China

Jianj...
Johns Hopkins
Baltimore,
jyu122@...



Automated Static Vulnerability Detection via a Holistic Neuro-symbolic Approach

P
Column
New
pl2689

Tool
Documentations

Tool
Implementations

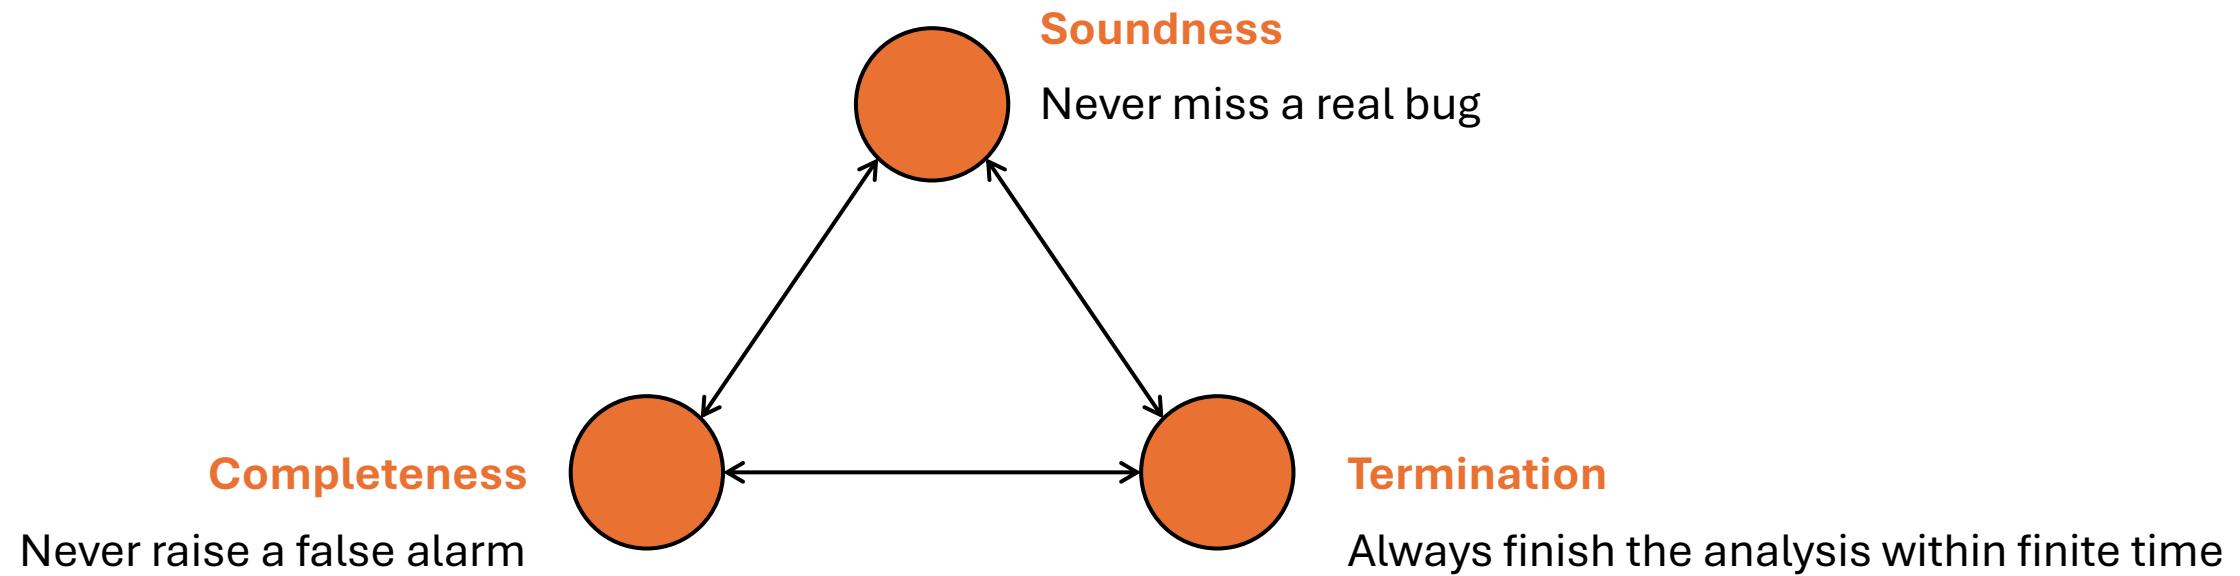
DSL EXTRACT

Task
Description

Table 1: New vulnerabilities discovered by MoCQ.

Project	Vul. Type	Experts'	MoCQ	Status
forkcms [42]	PHP SQLi	✓	✓	Reported
SyliusCmsPlugin [39]	PHP XSS	✗	✓	Reported
parse-static-imports [43]	JS Cmd.	✗	✓	Acknowledged
es2015-proxy [41]	JS Proto.	✗	✓	Acknowledged
web-worker [45]	JS Proto.	✓	✓	Reported
koa-send [46]	JS Proto.	✓	✓	Acknowledged
content-type [40]	JS Proto.	✗	✓	Reported

Software Testing: Impossible Triangle



Logistics – Week 9

- Assignment 3: Coding LLM Agents
 - <https://github.com/machine-programming/assignment-3>
 - Fully functional web-app agent. Due: Oct 23 (Thu)
- Forming groups for your final projects!
 - Form a group of 2-3 before This Thursday (~~Oct 19~~ Oct 23)