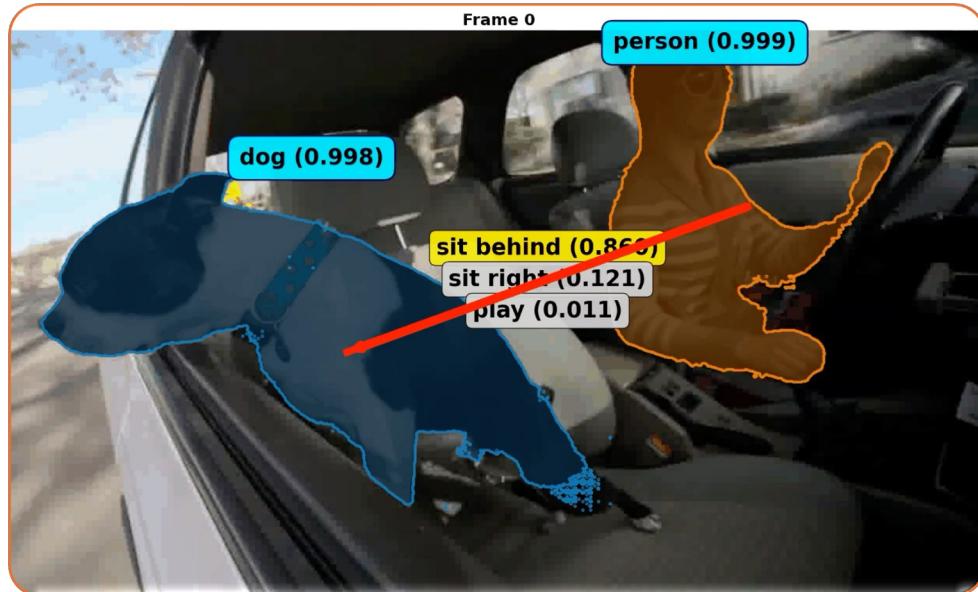


Machine Programming

Lecture 19 – Special Topic:
Neurosymbolic Programming & Application to Computer Vision

Spatio-Temporal Scene Graphs (STSGs)

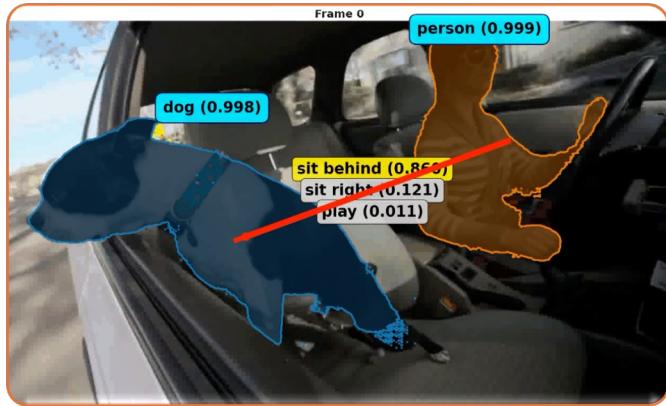


Scene Graphs



- Node: **Entity**
 - grounded by bounding-boxes or masks
- Self-loop: **Attribute**
 - entity class (e.g., person, dog, car)
 - physical properties (e.g., red, bending)
 - intransitive actions (e.g., moving, running)
- Edge: **Relations**
 - spatial relations (e.g., behind of, above)
 - interactions (e.g., moving towards)

Spatio-Temporal Scene Graphs



Evolving over time

Evolving probabilistic score is assigned to each element (node, edge, etc.) in the graph.

- Node: **Evolving Entity**
 - grounded by bounding-boxes or masks
- Self-loop: **Evolving Attribute**
 - entity class (e.g., person, dog, car)
 - physical properties (e.g., red, bending)
 - intransitive actions (e.g., moving, running)
- Edge: **Evolving Relations**
 - spatial relations (e.g., behind of, above)
 - interactions (e.g., moving towards)

Spatio-Temporal Scene Graphs: Applications



Identifying that the green cloth is “traversable” meaning that the model needs to understand physical properties of scene entities.

Spatio-Temporal Scene Graphs: Applications

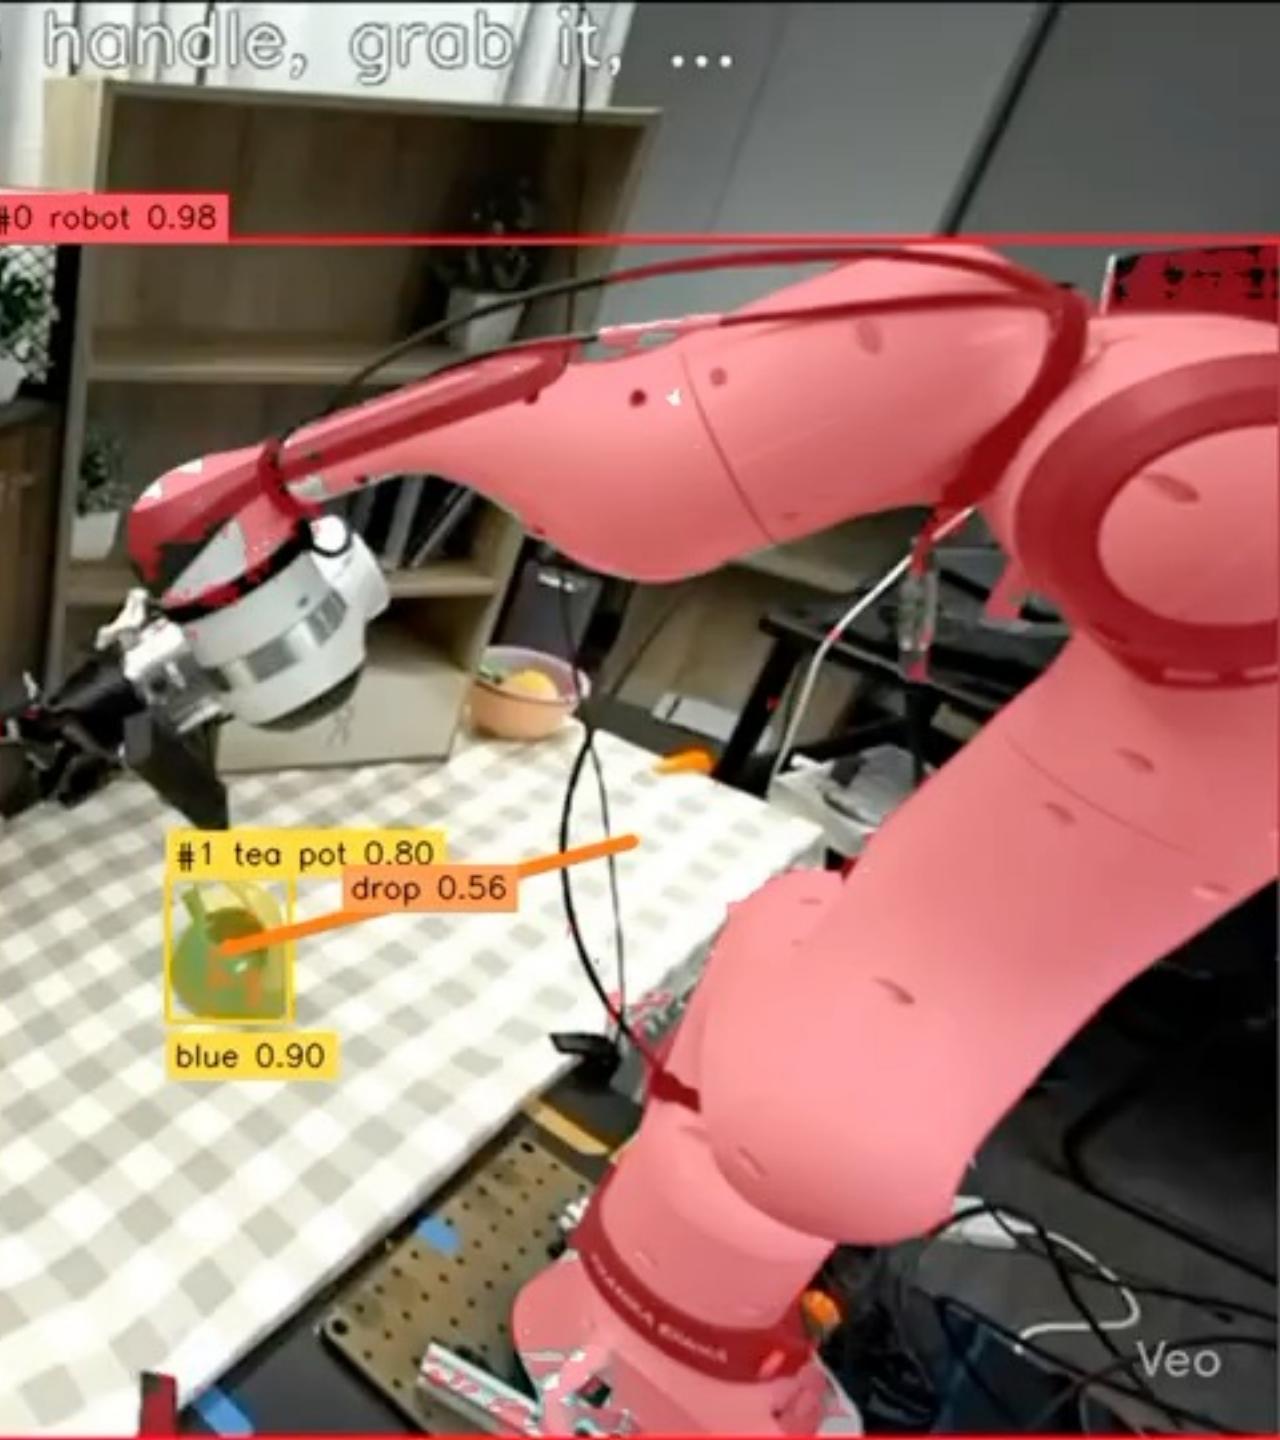
```
behavior achieve_at_3(g: Pose):  
    goal: at(g)  
    body:  
        bind x: Pose  
        bind t: Trajectory where:  
            connect(x, g, t)  
    [promotable:  
        achieve at(x)  
        forall o: Object:  
            achieve not blocking(o, t)  
    do exec_traj(t)
```



```
(:action pick-up  
:parameters (?p-paper)  
:precondition (and  
    (at ?l)  
    (isHomeBase ?l)  
    (unpacked ?p))  
:effect (and  
    (not (unpacked ?p))  
    (carrying ?p)))
```

Being able to **programmatically plan**
requires structured scene understanding
such as “**blocking(o, t)**”.

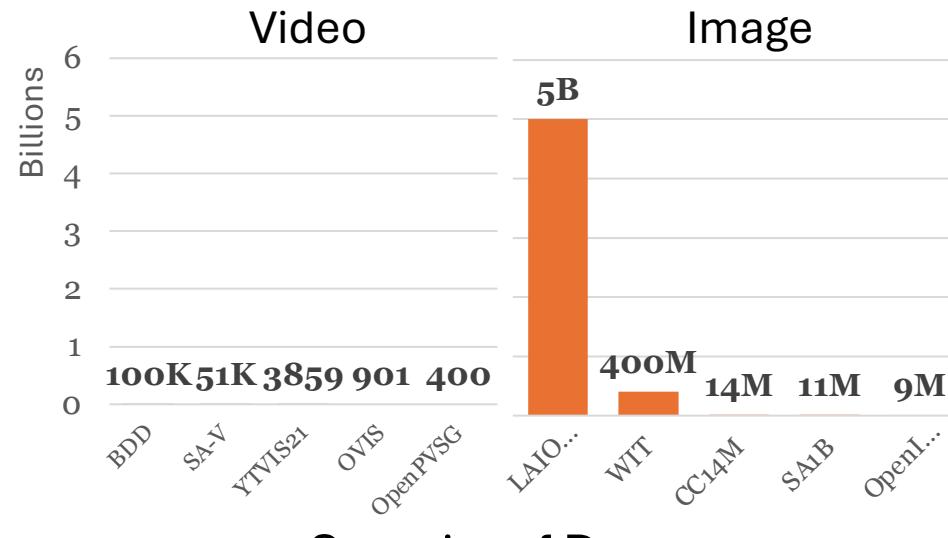
go to the teapot, go down to its handle, grab it, ...



Spatio-Temporal Scene Graphs: Challenges



Unknown Granularity



Scarcity of Dataset



Open-Vocabulary

Neurosymbolic to the rescue!

- Scene graph is a **symbolic data-structure**
 - Can be regularized by specifications
 - Symbolic algorithm can align scene graphs and specifications
- **Specifications** can be extracted from **natural scene description**
 - Vision-language models can generate natural language captions
 - Captions can be semantically parsed into
- **Neurosymbolic framework** can facilitate training
 - Scallop programming language is differentiable
 - Allowing to program alignment checker for neurosymbolic learning

Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning Spatio-temporal Scene Graphs with Weak Supervision

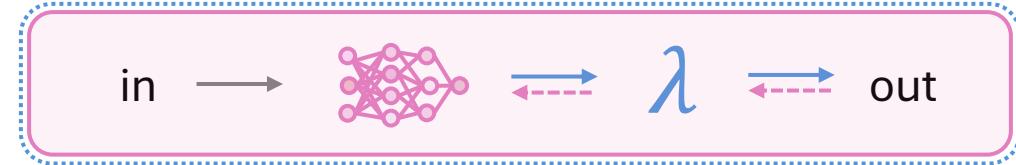
Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]



Scallop: A Language for Neurosymbolic Programming



[Algorithmic Supervision]

Illustrative Example: PacMan-Maze



Environment: a 5x5 grid of PacMan Arena

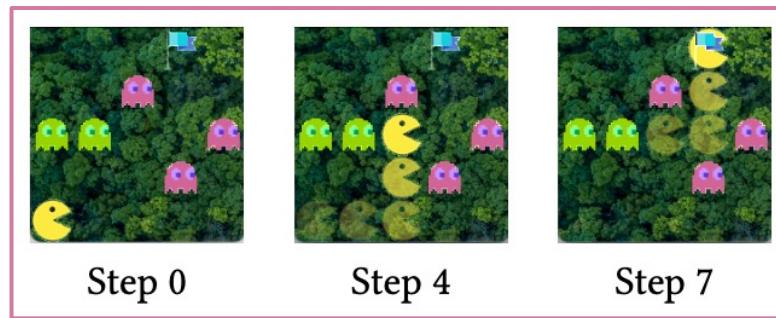
State: 200x200 colored image

Action: Up/Down/Left/Right

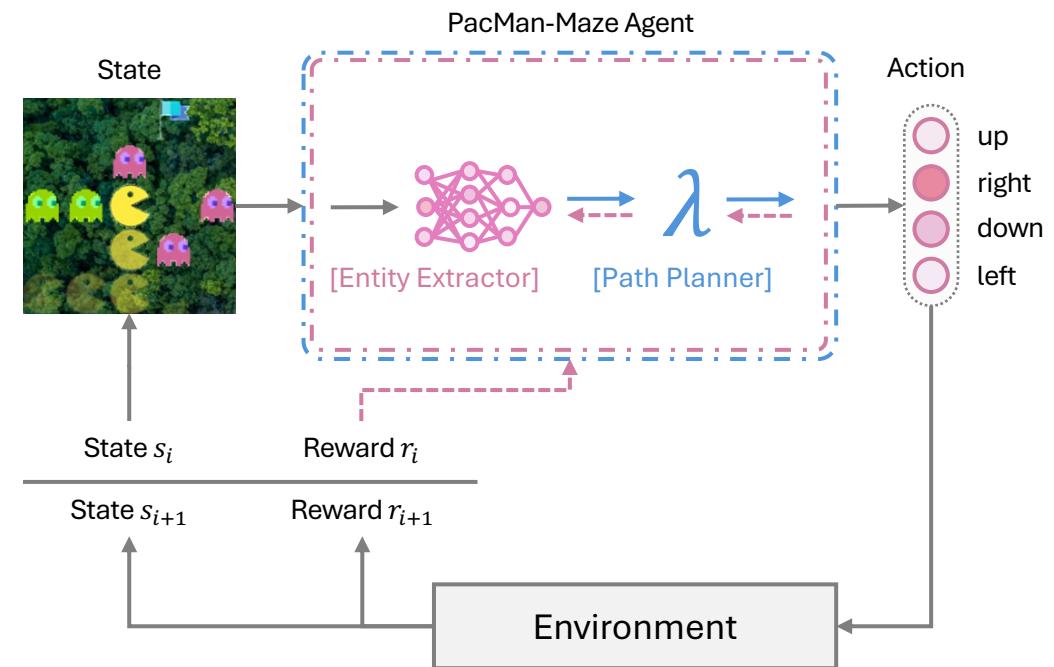
Goal: Reach the flag without touching enemy

(Environments are randomly generated)

Illustrative Example: PacMan-Maze



Environment: a 5x5 grid of PacMan Arena
State: 200x200 colored image
Action: Up/Down/Left/Right
Goal: Reach the flag without touching enemy
(Environments are randomly generated)



Illustrative Example

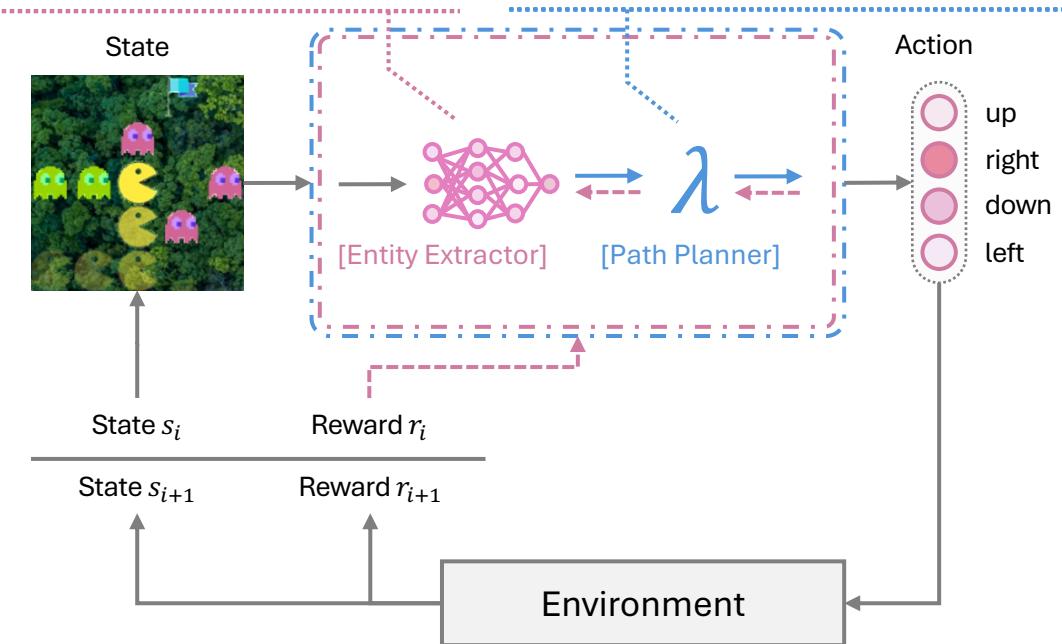


Entity Extractor

Simple **Convolutional Neural Network (CNN)** based model. Initialized randomly and **needs to be trained** to recognize entities.

Path Planner

Given the entities, **plan the best path** to reach the goal. Planner also needs to propagate **learning signal** to train the entity extractor.



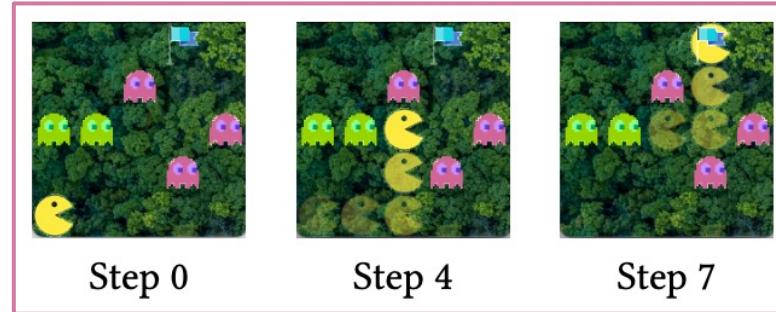
Environment: a 5x5 grid of PacMan Arena

State: 200x200 colored image

Action: Up/Down/Left/Right

Goal: Reach the flag without touching enemy
(Environments are randomly generated)

PacMan-Maze: Results



Environment: a 5x5 grid of PacMan Arena

State: 200x200 colored image

Action: Up/Down/Left/Right

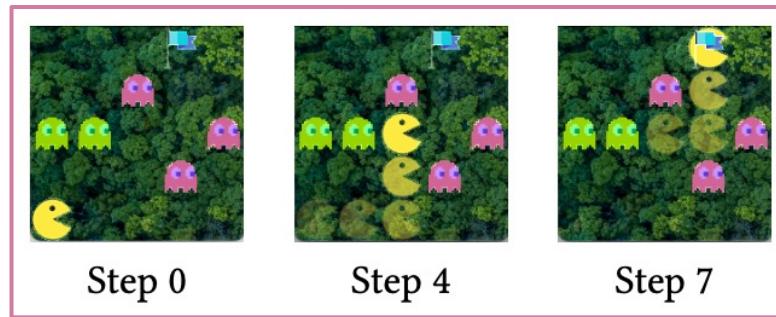
Goal: Reach the flag without touching enemy

(Environments are randomly generated)

Method	Success Rate / #Training Episodes
in → → out [Deep-Q-Network]	84.9% / 50,000
in → → λ → out [Scallop+CNN Q-Learning]	99.4% / 50

Success rate: whether the flag is reached within 50 steps.
#Training Episodes: # of training episodes needed to achieve the success rate.

Illustrative Example: PacMan-Maze

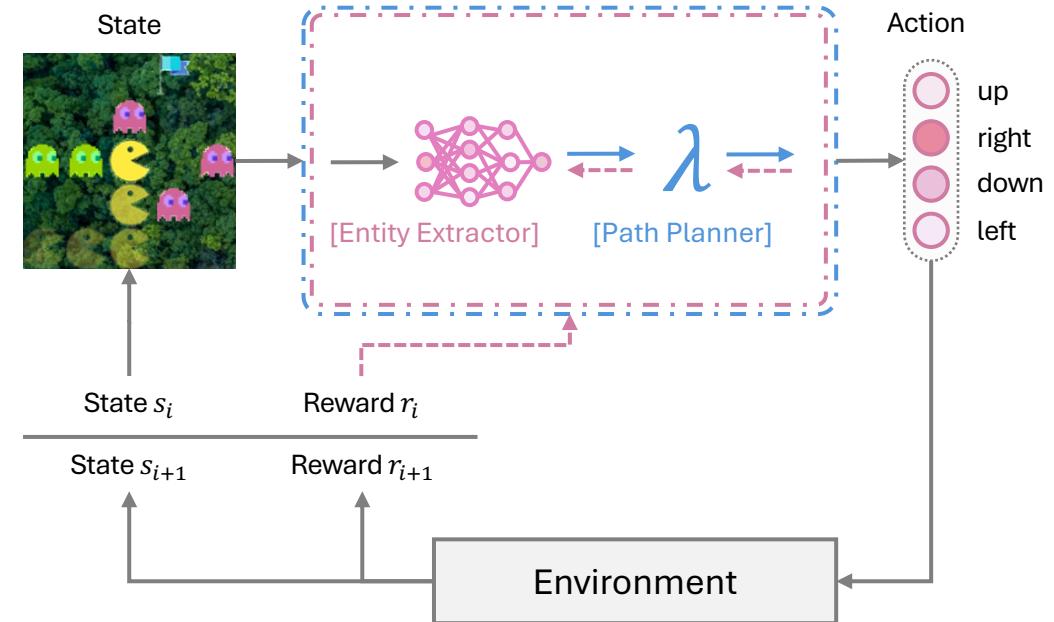


Environment: a 5x5 grid of PacMan Arena

State: 200x200 colored image

Action: Up/Down/Left/Right

Goal: Reach the flag without touching enemy
(Environments are randomly generated)



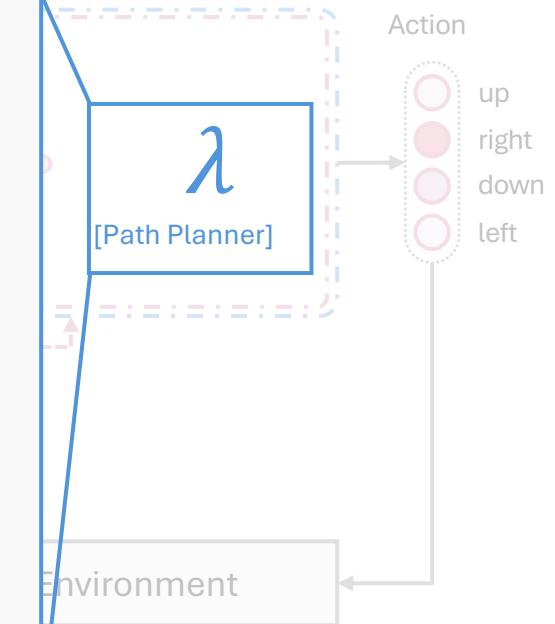
Illustrative Example: PacMan-Maze

```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32), flag(i: i32, j: i32), ghost(i: i32, j: i32)

// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and safe_cell(i, j + 1) //...

// Find safe paths recursively
rel path(i, j, i, j) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3, _)

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
                    path(m, n, g, h) and flag(g, h)
```



Illustrative Example: PacMan-Maze

```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32) = Pr(pacmanij |  $\theta$ )
// Entity extractor outputs a probability matrix
rel pacman(i: i32, j: i32) {
    | i   | j | |
    | 0.00 | 2   | 1   |
    | 0.92 | 2   | 2   |
    | 0.02 | 3   | 4   |
    | ...  | ... | ... |
}
// Compute the action that can lead to the next state
rel take_action(a) = pacman(i, j) and path(m, n, g, h) and flag(g, h)
```

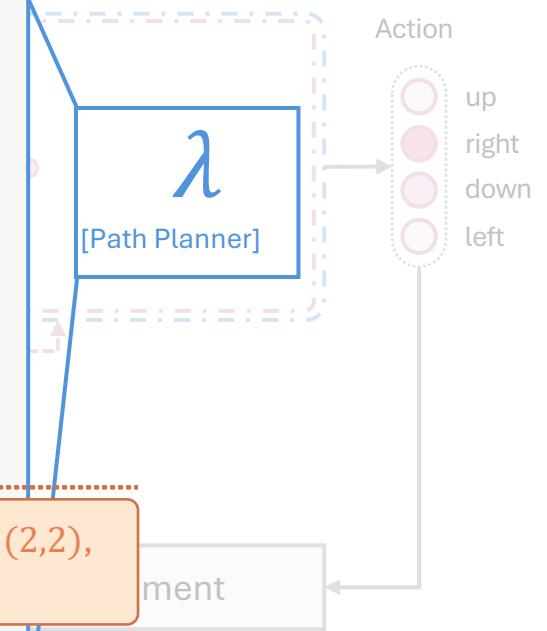
pacman(i: i32, j: i32)

	i	j
0.00	2	1
0.92	2	2
0.02	3	4
...

not ghost(i, j) and safe_cell(i, j + 1) //...

path(i₁, j₁, i₂, j₂) = path(i₁, j₁, i₂, j₂) and edge(i₂, j₂, i₃, j₃, ...)

Estimates the probability of PacMan being at the location (2,2), given the neural network parameters θ .



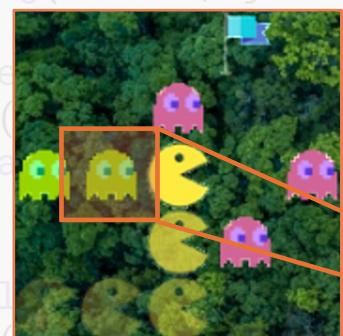
Illustrative Example: PacMan-Maze

```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32), flag(i: i32, j: i32), ghost(i: i32, j: i32)

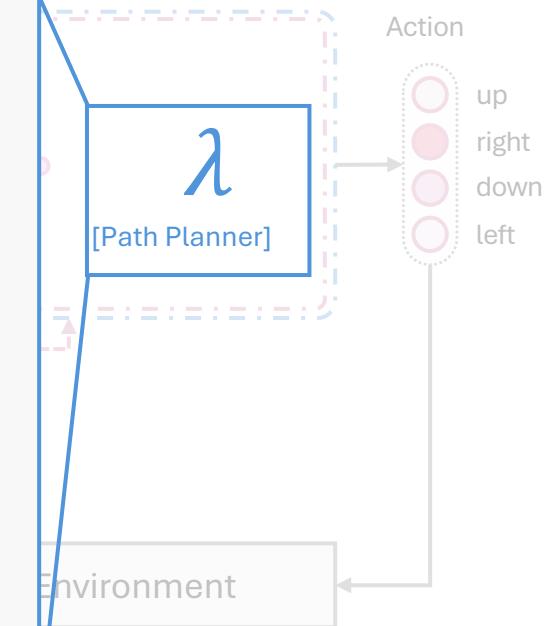
// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and not ghost(i, j + 1, j)
rel edge(i, j, i + 1, j, LEFT) = safe_cell(i, j) and not ghost(i + 1, j, j)
rel edge(i, j, i - 1, j, RIGHT) = safe_cell(i, j) and not ghost(i - 1, j, j)
rel edge(i, j, i, j - 1, DOWN) = safe_cell(i, j) and not ghost(i, j - 1, j)

// Find safe paths recursively
rel path(i, j, i, j) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3, _)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3, _)

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
                    path(m, n, g, h) and flag(g, h)
```



$\Pr(\text{ghost}_{ij} \theta)$	i	j
0.00	0	0
0.94	1	2
0.85	2	1
...



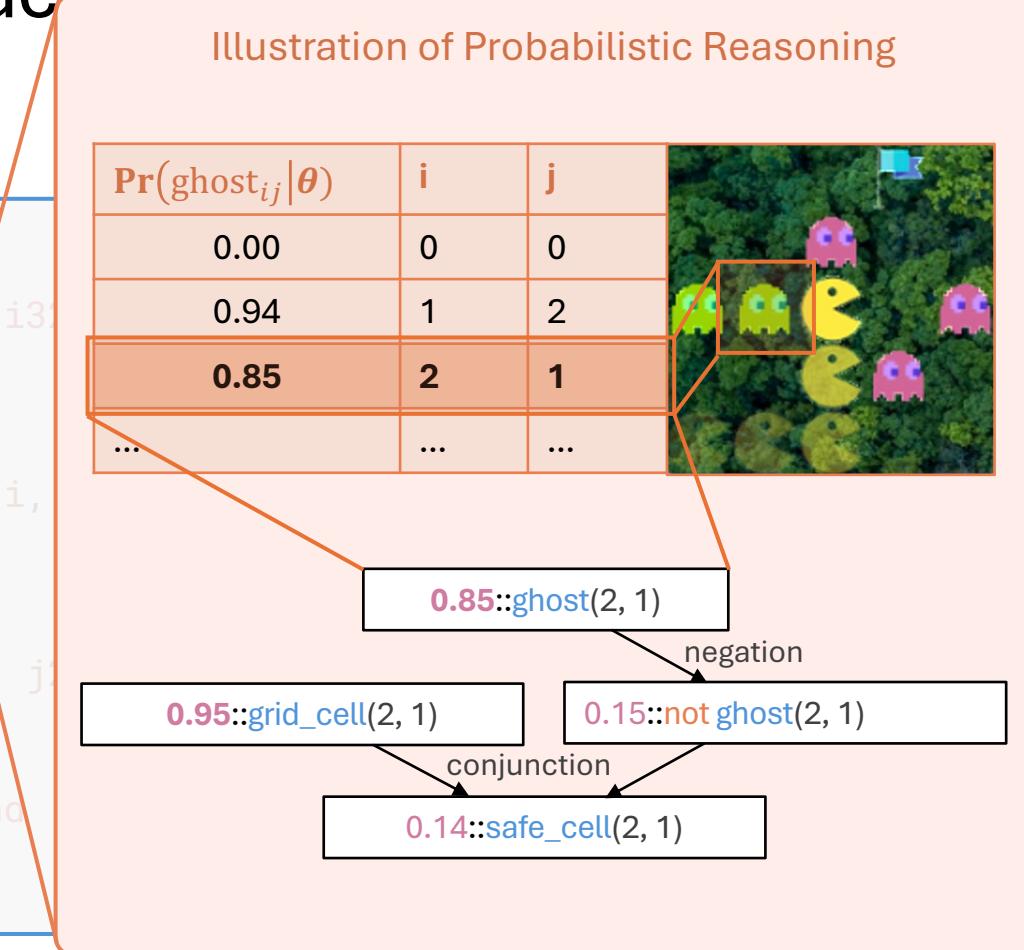
Illustrative Example: PacMan-Maze

```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32), flag(i: i32, j: i32), ghost(i: i32, j: i32)

// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i1, j1) = safe_cell(i, j) and safe_cell(i1, j1) and (i, j ≠ i1, j1) and (i, j ≠ i1, j or i, j ≠ i, j1) and (i, j ≠ i, j1)

// Find safe paths recursively
rel path(i, j, i1, j1) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3) and (i1, j1 ≠ i2, j2) and (i1, j1 ≠ i3, j3) and (i2, j2 ≠ i3, j3)

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
    path(m, n, g, h) and flag(g, h)
```



Illustrative Example: PacMan-Maze

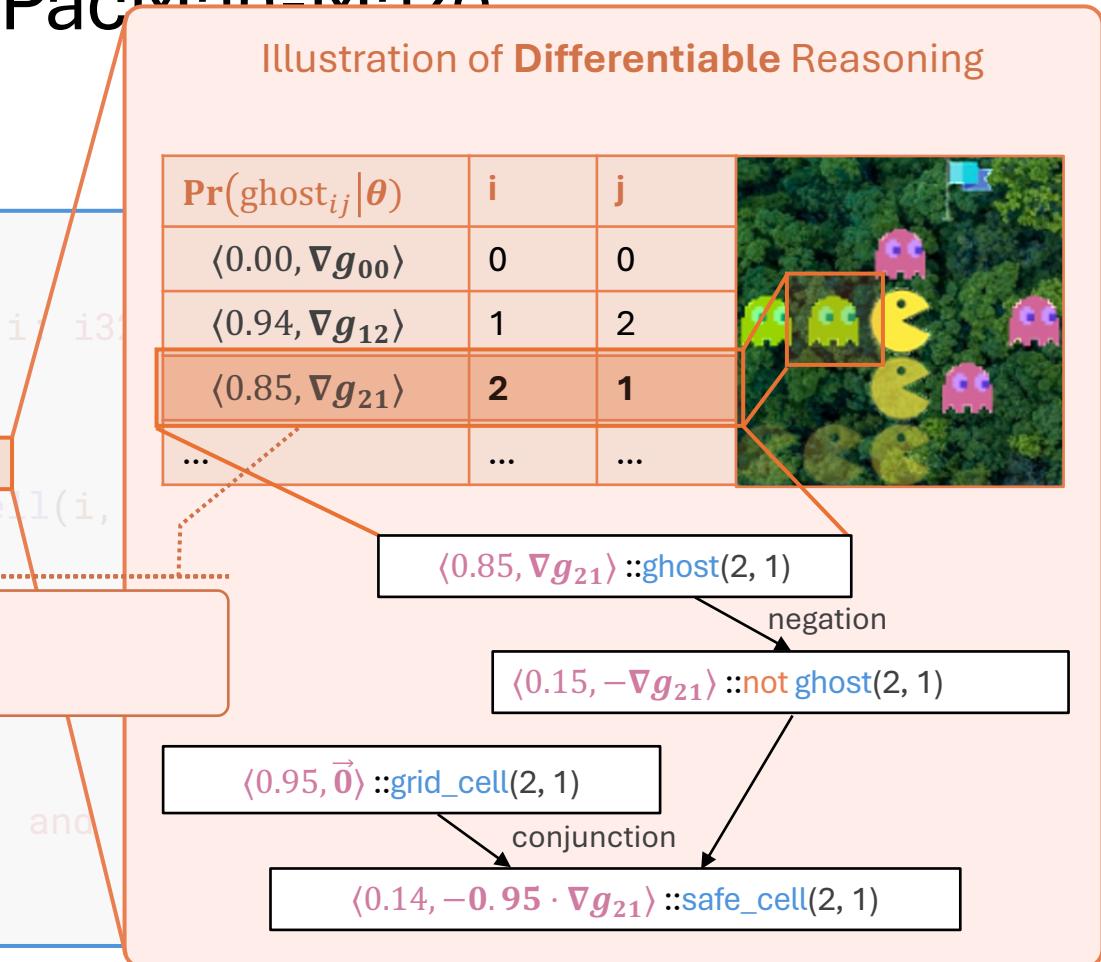
```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32), flag(i: i32, j: i32), ghost(i: i32, j: i32)

// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and safe_cell(i, j + 1)

// Find safe paths
rel path(i, j, i, j, a) = ...
rel path(i1, j1, i2, j2, a) = ...

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
                    path(m, n, g, h) and flag(g, h)
```

$\frac{\partial g_{21}}{\partial \theta}$ The gradient carries **partial derivative** with respect to neural-network parameters!



Illustrative Example: PacMan-Maze

```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32)

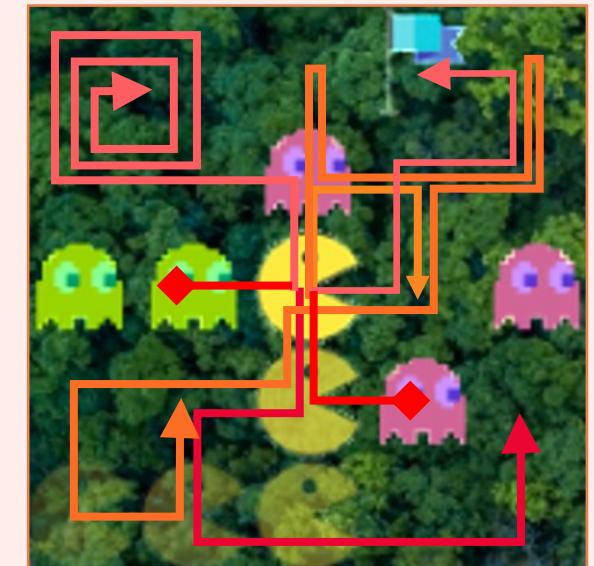
// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and safe_cell(i, j + 1) //...
rel edge(i, j, i + 1, j, RIGHT) = safe_cell(i, j) and safe_cell(i + 1, j) //...
rel edge(i, j, i - 1, j, LEFT) = safe_cell(i, j) and safe_cell(i - 1, j) //...
rel edge(i, j, i, j - 1, DOWN) = safe_cell(i, j) and safe_cell(i, j - 1) //...

// Find safe paths recursively
rel path(i, j, i, j) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3, _)

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
    path(m, n, g, h) and flag(g, h)
```

Recursively building safe paths by appending safe edges to existing safe paths.

What would happen when reasoning becomes more complicated?



Non-Termination

Complexity of Exact Probability Estimation

Illustrative Example

```
// Inputs from the Environment
type pacman(i: i32, j: i32)
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and
    safe_cell(i, j + 1)
```

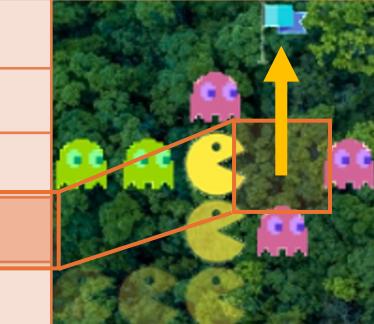
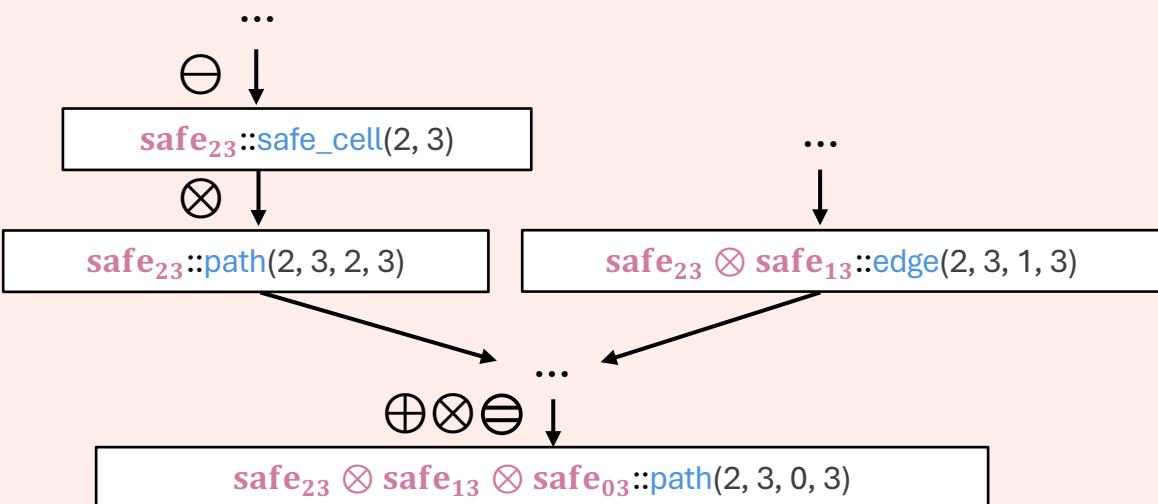
```
// Find safe paths recursively
rel path(i, j, i, j) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and
    path(i2, j2, i3, j3)
```

```
// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, i, j + 1, UP) and
    path(m, n, g, h) and flag(g, h)
```

(Tag Space) $t \in T$
 (False) $0 \in T$
 (True) $1 \in T$
 (Disjunction) $\oplus : T \times T \rightarrow T$
 (Conjunction) $\otimes : T \times T \rightarrow T$
 (Negation) $\ominus : T \rightarrow T$
 (Saturation) $\equiv : T \times T \rightarrow \text{Bool}$

Illustration of Abstracted Tagged-Semantics

Tag: T	i	j
safe ₀₀	0	0
safe ₁₃	1	3
safe₂₃	2	3
...

Built-in Library of Provenance Structures

Kind	Provenance	T	$\mathbf{0}$	$\mathbf{1}$	\oplus	\otimes	\ominus	\ominus	τ	ρ
Discrete	unit	$\{\mathbf{0}\}$	$\mathbf{0}$	$\mathbf{0}$	$\lambda t_1, t_2.()$	$\lambda t_1, t_2.()$	$\lambda a.\text{FAIL}$	$==$	$\lambda i.()$	$\lambda t.()$
	bool	$\{\top, \perp\}$	\perp	\top	\vee	\wedge	\neg	$==$	id	id
	natural	\mathbb{N}	0	1	$+$	\times				
Probabilistic	max-min-prob	$[0, 1]$	0	1	max	min	(Tag Space)	$t \in T$		
	add-mult-prob	$[0, 1]$	0	1	$\lambda t_1, t_2.\text{clamp}(t_1 + t_2)$	$\lambda t_1, t_2.(t_1 \cdot t_2)$	(False)	$0 \in T$		
	nand-min-prob	$[0, 1]$	0	1	$\lambda t_1, t_2. - (1 - t_1)(1 - t_2)$	min	(True)	$1 \in T$		
	nand-mult-prob	$[0, 1]$	0	1	$\lambda t_1, t_2. - (1 - t_1)(1 - t_2)$	$\lambda t_1, t_2.t_1 \cdot t_2$	(Disjunction)	$\oplus : T \times T \rightarrow T$		
	top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$	(Conjunction)	$\otimes : T \times T \rightarrow T$		
	sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$	(Negation)	$\ominus : T \rightarrow T$		
Differentiable	diff-max-min-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	max	min	(Saturation)	$\ominus : T \times T \rightarrow \text{Bool}$		
	diff-add-mult-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2.\text{clamp}(\hat{t}_1 + \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2.\hat{t}_1 \cdot \hat{t}_2$				
	diff-nand-min-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. - (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	min				
	diffnand-mult-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. - (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2.\hat{t}_1 \cdot \hat{t}_2$				
	diff-top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$				
	diff-sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$				
...

Inspired by Provenance Semiring (Green et. al. PODS 2007)

Approximated Probabilistic Provenance

Top-k Proofs

T_{tkp} : Boolean Formula in **DNF**

\otimes_{tkp} : Perform the logical “conjunction”, “disjunction”, and “negation”,

\oplus_{tkp} : while **preserving only the k most significant conjunctive clauses**.

\ominus_{tkp} :

Top-Bottom-k Clauses

T_{tbkc} : Boolean Formula in **DNF** or **CNF**

\otimes_{tbkc} : Perform the logical “conjunction”, “disjunction”, and “negation”. For

\oplus_{tbkc} : DNF, keep the **top- k conjunctive clauses**, while for CNF, keep the **bottom- k disjunctive clauses**.

\ominus_{tbkc} :

Optimal-k Proofs

T_{okp} : Boolean Formula in **DNF**

\otimes_{okp} : Employs a **greedy** algorithm to preserve the **k conjunctive clauses** that has a joint probability the **closest to the exact probability**.

\oplus_{okp} :

Inspired by [Renkens et. al. 2012]

Performs **Weighted Model Counting (WMC)** to recover probabilities from Boolean formula ϕ_y .

$$\Pr(y \mid \theta) = \text{WMC}(\phi_y, \{x \mapsto \Pr(x \mid \theta)\})$$

(Tag Space)	t	\in	T
(False)	0	\in	T
(True)	1	\in	T
(Disjunction)	\oplus	$:$	$T \times T \rightarrow T$
(Conjunction)	\otimes	$:$	$T \times T \rightarrow T$
(Negation)	\ominus	$:$	$T \rightarrow T$
(Saturation)	\ominus	$:$	$T \times T \rightarrow \text{Bool}$

Approximated Differentiable Provenance

Diff. Top-k Proofs

- T_{tkp} : Boolean Formula in DNF
- \otimes_{tkp} : Perform the logical “conjunction”, “disjunction”, and “negation”, while
- \oplus_{tkp} : **preserving only the k most significant conjunctive clauses.**
- \ominus_{tkp} : **significance conjunctive clauses.**

Diff. Top-Bottom-k Clauses

- T_{tbkc} : Boolean Formula in DNF or CNF
- \otimes_{tbkc} : Perform the logical “conjunction”, “disjunction”, and “negation”. For DNF, keep the **top- k conjunctive clauses**, while for CNF, keep the **bottom- k disjunctive clauses**.
- \oplus_{tbkc} :
- \ominus_{tbkc} :

Diff. Optimal-k Proofs

- T_{okp} : Boolean Formula in DNF
- \otimes_{okp} : Employs a **greedy** algorithm to preserve the **k conjunctive clauses** that has a joint probability the **closest to the exact probability**.
- \oplus_{okp} :
- \ominus_{okp} : Inspired by [Renkens et. al. 2012]

Performs **Weighted Model Counting (WMC)** to recover probabilities **with gradients** from Boolean formula ϕ_y .

$$\langle \Pr(y | \theta), \nabla \Pr(y | \theta) \rangle = \text{WMC}(\phi_y, \{x \mapsto \langle \Pr(x | \theta), \nabla \Pr(x | \theta) \rangle\})$$

(Tag Space)	t	\in	T
(False)	0	\in	T
(True)	1	\in	T
(Disjunction)	\oplus	$:$	$T \times T \rightarrow T$
(Conjunction)	\otimes	$:$	$T \times T \rightarrow T$
(Negation)	\ominus	$:$	$T \rightarrow T$
(Saturation)	\ominus	$:$	$T \times T \rightarrow \text{Bool}$

Built-in Library of Provenance Structures

Kind	Provenance	T	$\mathbf{0}$	$\mathbf{1}$	\oplus	\otimes	\ominus	\ominus	τ	ρ
Discrete	unit	$\{\mathbf{0}\}$	$\mathbf{0}$	$\mathbf{0}$	$\lambda t_1, t_2.()$	$\lambda t_1, t_2.()$	$\lambda a.\text{FAIL}$	$==$	$\lambda i.()$	$\lambda t.()$
	bool	$\{\top, \perp\}$	\perp	\top	\vee	\wedge	\neg	$==$	id	id
	natural	\mathbb{N}	0	1	$+$	\times				
Probabilistic	max-min-prob	$[0, 1]$	0	1	max	min	(Tag Space)	$t \in T$		
	add-mult-prob	$[0, 1]$	0	1	$\lambda t_1, t_2.\text{clamp}(t_1 + t_2)$	$\lambda t_1, t_2.(t_1 \cdot t_2)$	(False)	$0 \in T$		
	nand-min-prob	$[0, 1]$	0	1	$\lambda t_1, t_2. - (1 - t_1)(1 - t_2)$	min	(True)	$1 \in T$		
	nand-mult-prob	$[0, 1]$	0	1	$\lambda t_1, t_2. - (1 - t_1)(1 - t_2)$	$\lambda t_1, t_2.t_1 \cdot t_2$	(Disjunction)	$\oplus : T \times T \rightarrow T$		
	top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$	(Conjunction)	$\otimes : T \times T \rightarrow T$		
	sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$	(Negation)	$\ominus : T \rightarrow T$		
Differentiable	diff-max-min-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	max	min	(Saturation)	$\ominus : T \times T \rightarrow \text{Bool}$		
	diff-add-mult-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2.\text{clamp}(\hat{t}_1 + \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2.\hat{t}_1 \cdot \hat{t}_2$				
	diff-nand-min-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. - (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	min				
	diffnand-mult-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. - (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2.\hat{t}_1 \cdot \hat{t}_2$				
	diff-top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$				
	diff-sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$				
...

Inspired by Provenance Semiring (Green et. al. PODS 2007)

Built-in Library of Provenance Structures

Kind	Provenance	T	$\mathbf{0}$	$\mathbf{1}$	\oplus	\otimes	\ominus	\ominus	τ	ρ
Discrete	unit	{0}	0	0	$\lambda t_1, t_2. ()$	$\lambda t_1, t_2. ()$	$\lambda a.\text{FAIL}$	\equiv	$\lambda i.()$	$\lambda t.()$
	bool						\neg	\equiv	id	id
	natural						$\lambda n. \mathbb{1}[n > 0]$	\equiv	id	id
Probabilistic	max-min-prob						$\lambda t. 1 - t$	\equiv	id	id
	add-mult-prob						$\lambda t. 1 - t$	$\lambda t. \top$	id	id
	nand-min-prob	[0, 1]	0	1	$\lambda t_1, t_2. \neg (1 - t_1)(1 - t_2)$	min	$\lambda t. 1 - t$	$\lambda t. \top$	id	id
	nand-mult-prob	[0, 1]	0	1	$\lambda t_1, t_2. \neg (1 - t_1)(1 - t_2)$	$\lambda t_1, t_2. t_1 \cdot t_2$	$\lambda t. 1 - t$	$\lambda t. \top$	id	id
	top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$	$\neg_{\text{top-}k}$	\equiv	$\lambda p_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \Gamma)$
Differentiable	sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$	$\neg_{\text{sample-}k}$	\equiv	$\lambda p_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \Gamma)$
	diff-max-min-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	max	min	$\lambda \hat{t}. \hat{1} - \hat{t}$	\equiv	id	id
	diff-add-mult-prob	\mathbb{D}	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. \text{clamp}(\hat{t}_1 + \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2. \hat{t}_1 \cdot \hat{t}_2$	$\lambda \hat{t}. \hat{1} - \hat{t}$	$\lambda t. \top$	id	id
	diff-nand-min-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. \neg (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	min	$\lambda \hat{t}. \hat{1} - \hat{t}$	$\lambda t. \top$	id	id
	diffnand-mult-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. \neg (\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2. \hat{t}_1 \cdot \hat{t}_2$	$\lambda \hat{t}. \hat{1} - \hat{t}$	$\lambda t. \top$	id	id
	diff-top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$	$\neg_{\text{top-}k}$	\equiv	$\lambda \hat{p}_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \hat{\Gamma})$
	diff-sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$	$\neg_{\text{sample-}k}$	\equiv	$\lambda \hat{p}_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \Gamma)$
...

Employing **approximation** algorithms helps in scaling the probabilistic and differentiable reasoning modules.

Allowing extensions such as uncertainty intervals and embeddings

Built-in Library of Provenance Structures

Kind	Provenance	T	0	1	\oplus	\otimes	\ominus	\ominus	τ	ρ
Discrete	unit	$\{()\}$	$()$	$()$	$\lambda t_1, t_2.()$	$\lambda t_1, t_2.()$	$\lambda a.\text{FAIL}$	$==$	$\lambda i.()$	$\lambda t.()$
	bool	$\{\top, \perp\}$	\perp	\top	\vee	\wedge	\neg	$==$	id	id
	natural	\mathbb{N}	0	1	$+$	\times	$\lambda n. \mathbb{1}[n > 0]$	$==$	id	id
Probabilistic	max-min-prob	$[0, 1]$	0	1	max	min	$\lambda t. 1 - t$	$==$	id	id

<p style="text-align: center;">While the syntax of Scallop program remains familiar, employing different tags allows discrete, probabilistic, and differentiable modes of reasoning.</p>										
Differentiable	discrete	\mathbb{D}	0	1	$\lambda t_1, t_2.\text{clamp}(t_1 + t_2)$	$\lambda t_1, t_2. t_1 \cdot t_2$	$\lambda t. 1 - t$	$\lambda t. \top$	id	id
	diff-add-mult-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. -(\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	min	$\lambda \hat{t}. \hat{1} - \hat{t}$	$\lambda \hat{t}. \top$	id	id
	diff-nand-min-prob	$[\hat{0}, \hat{1}]$	$\hat{0}$	$\hat{1}$	$\lambda \hat{t}_1, \hat{t}_2. -(\hat{1} - \hat{t}_1)(\hat{1} - \hat{t}_2)$	$\lambda \hat{t}_1, \hat{t}_2. \hat{t}_1 \cdot \hat{t}_2$	$\lambda \hat{t}. \hat{1} - \hat{t}$	$\lambda \hat{t}. \top$	id	id
	diff-top-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{top-}k}$	$\wedge_{\text{top-}k}$	$\neg_{\text{top-}k}$	$==$	$\lambda \hat{p}_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \hat{\Gamma})$
	diff-sample-k-proofs	Φ	\emptyset	$\{\emptyset\}$	$\vee_{\text{sample-}k}$	$\wedge_{\text{sample-}k}$	$\neg_{\text{sample-}k}$	$==$	$\lambda \hat{p}_i. \{\{\text{pos}(i)\}\}$	$\lambda \varphi. \text{WMC}(\varphi, \hat{\Gamma})$
...

Illustrative Example

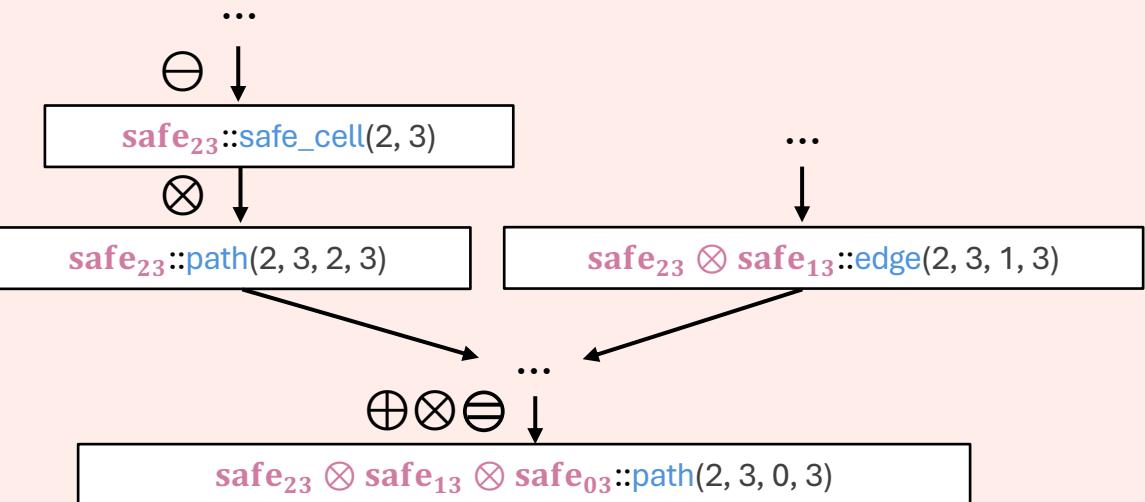
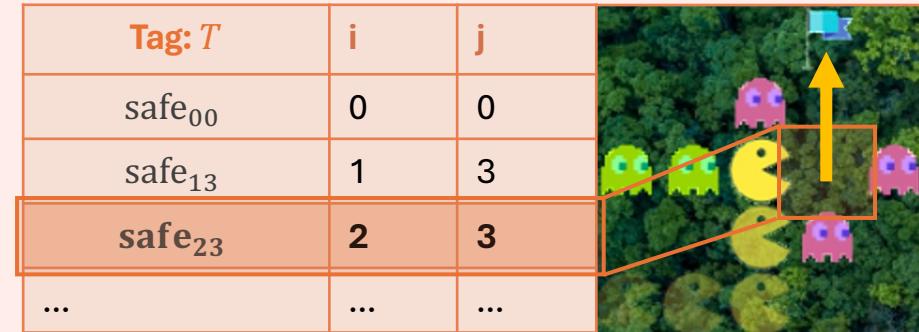
```
// Inputs from the Entity Extractor (Neural Network)
type pacman(i: i32, j: i32), flag(i: i32, j: i32)

// Find safe cells and safe edges
rel safe_cell(i, j) = grid_cell(i, j) and not ghost(i, j)
rel edge(i, j, i, j + 1, UP) = safe_cell(i, j) and
    safe_cell(i, j + 1) and not ghost(i, j + 1, UP)

// Find safe paths recursively
rel path(i, j, i, j) = safe_cell(i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and
    path(i2, j2, i3, j3)
    ... and ...

// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, i, j + 1, UP) and
    path(m, n, g, h) and flag(g, h)
```

Illustration of Abstracted Tagged-Semantics



Illustrative Example: PacM

```
// Inputs from the Entity Store
type pacman(i: i32, j: i32)
type ghost(i: i32, j: i32)
type edge(i: i32, j: i32, m: i32, n: i32, a: action)
type path(i: i32, j: i32, i1: i32, j1: i32, i3: i32, j3: i32)
type flag(i: i32, j: i32)

// Find safe cells and edges
rel safe_cell(i, j) = ghost(i, j) = false
rel edge(i, j, i, j) = false
rel edge(i, j, i, j + 1, UP)
rel edge(i, j, i + 1, j, DOWN)
rel edge(i, j, i, j - 1, LEFT)
rel edge(i, j, i - 1, j, RIGHT)

// Find safe paths recursively
rel path(i, j, i1, j1) = safe_cell(i, j) and path(i1, j1, i, j)
rel path(i1, j1, i3, j3) = path(i1, j1, i2, j2) and edge(i2, j2, i3, j3, a) and path(i3, j3, i, j)

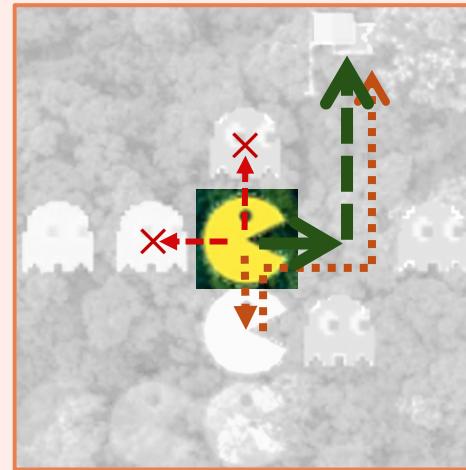
// Compute the action that can lead to the goal
rel take_action(a) = pacman(i, j) and edge(i, j, m, n, a) and
path(m, n, g, h) and flag(g, h)
```

Estimates the **probability of success** after taking the action a , conditioned on neural network parameters θ .

`take_action(RIGHT) <- pacman(2,2) and
not ghost(2,3) and
not ghost(1,3) and
flag(0,3) or ...`

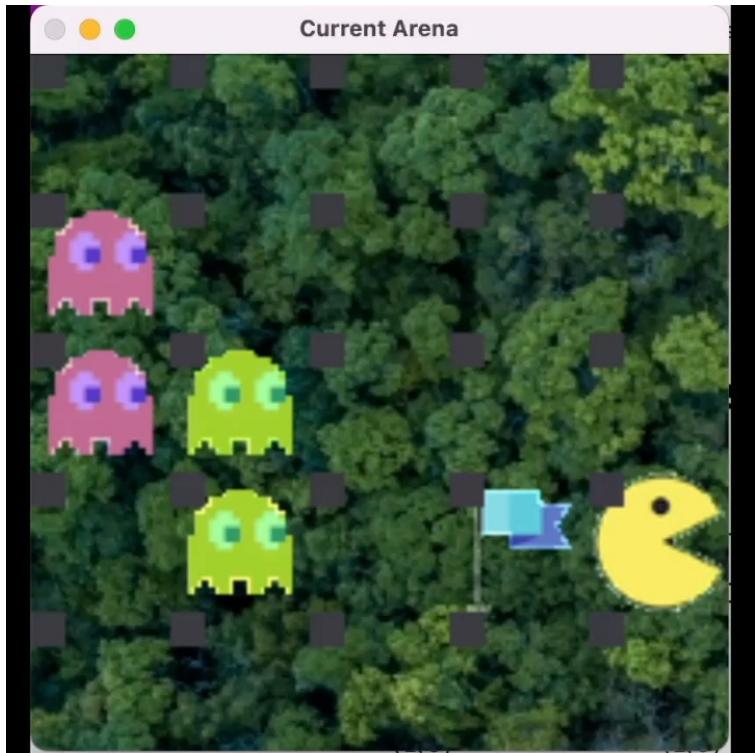
$\Pr(\text{success} \mid \text{action} = \text{RIGHT}) = \Pr(\text{pacman}_{22}) \times$
 $(1 - \Pr(\text{ghost}_{23})) \times$
 $(1 - \Pr(\text{ghost}_{13})) \times$
 $\Pr(\text{flag}_{03}) + \dots$

Applying Top-k-Proofs Provenance



$\widehat{\Pr}(\text{success} \mid \text{action} = a, \theta)$	a
$\langle 0.01, \nabla \text{UP} \rangle$	UP
$\langle 0.92, \nabla \text{RIGHT} \rangle$	RIGHT
$\langle 0.79, \nabla \text{DOWN} \rangle$	DOWN
$\langle 0.02, \nabla \text{LEFT} \rangle$	LEFT

Demo: Training PacMan-Maze Agent

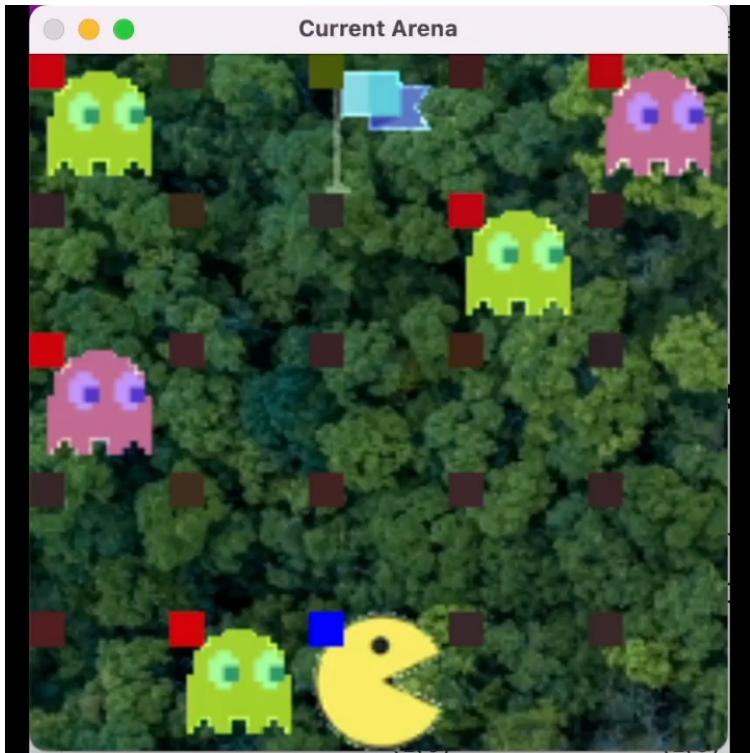


Top-left marker indicates neural network prediction
on what the cell contains.

- Ghost (Red)
- PacMan (Blue)
- Flag (Green)



Demo: Testing PacMan-Maze Agent



Top-left marker indicates neural network prediction
on what the cell contains.

- Ghost (Red)
- PacMan (Blue)
- Flag (Green)



Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning
Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

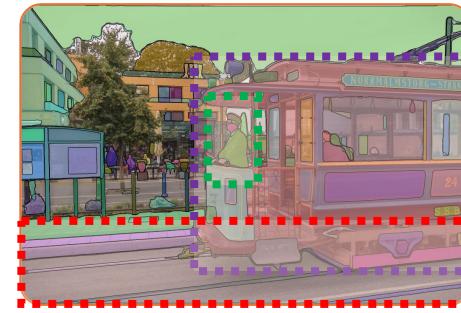
ESCA: Contextualizing Embodied Agents via
Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

Neurosymbolic Multi-modal Alignment



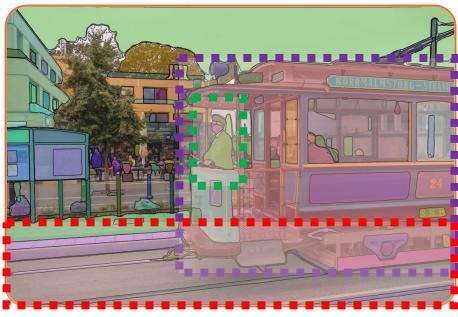
Unknown Granularity



A driver is driving the
trolley towards the left"

Known Granularity

Neurosymbolic Multi-modal Alignment



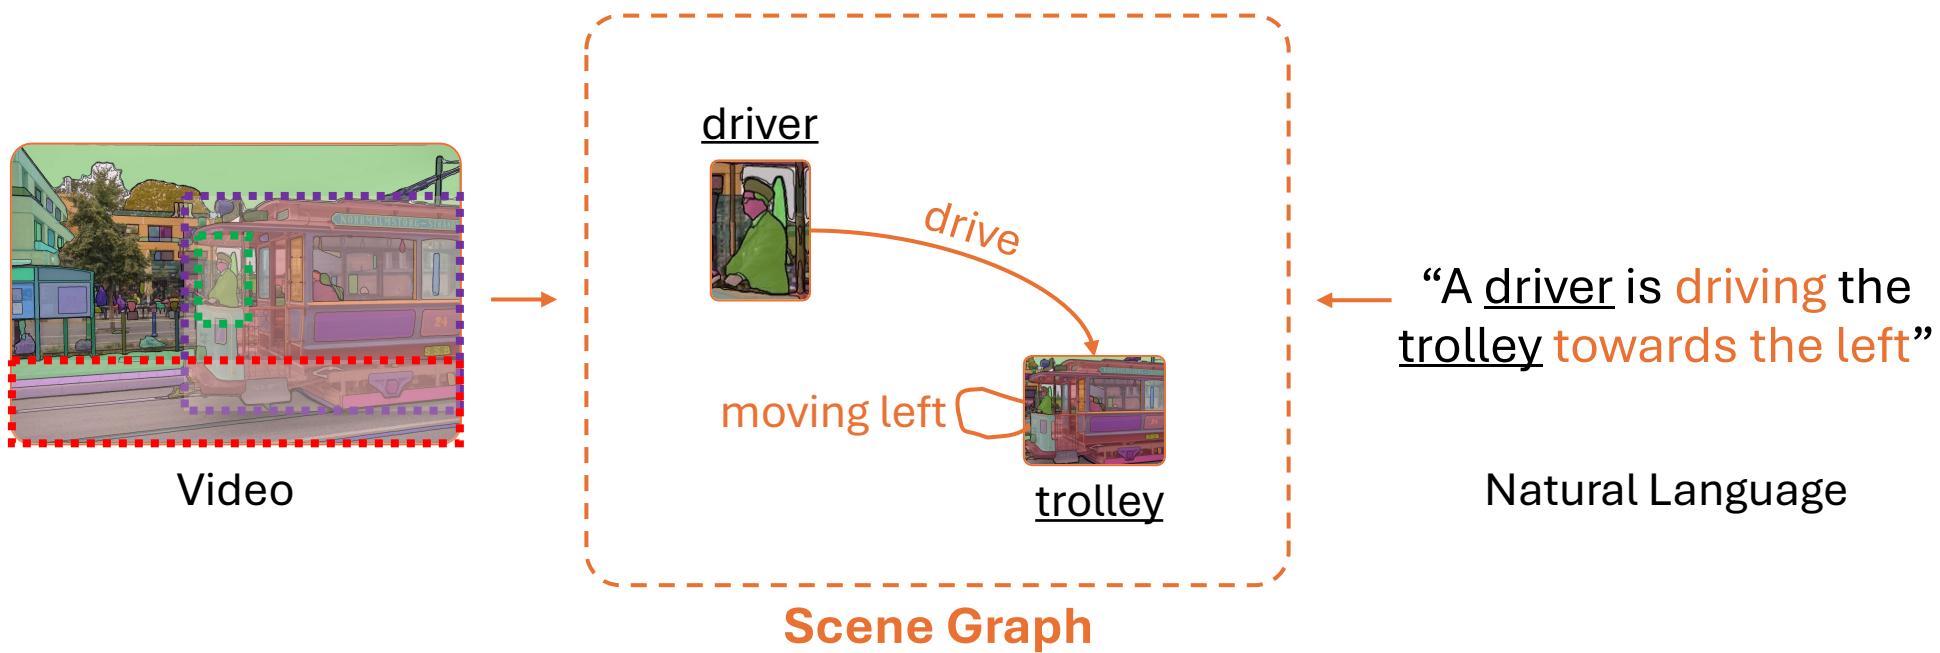
Video

Spatio-Temporal
Scene Graph
Alignment

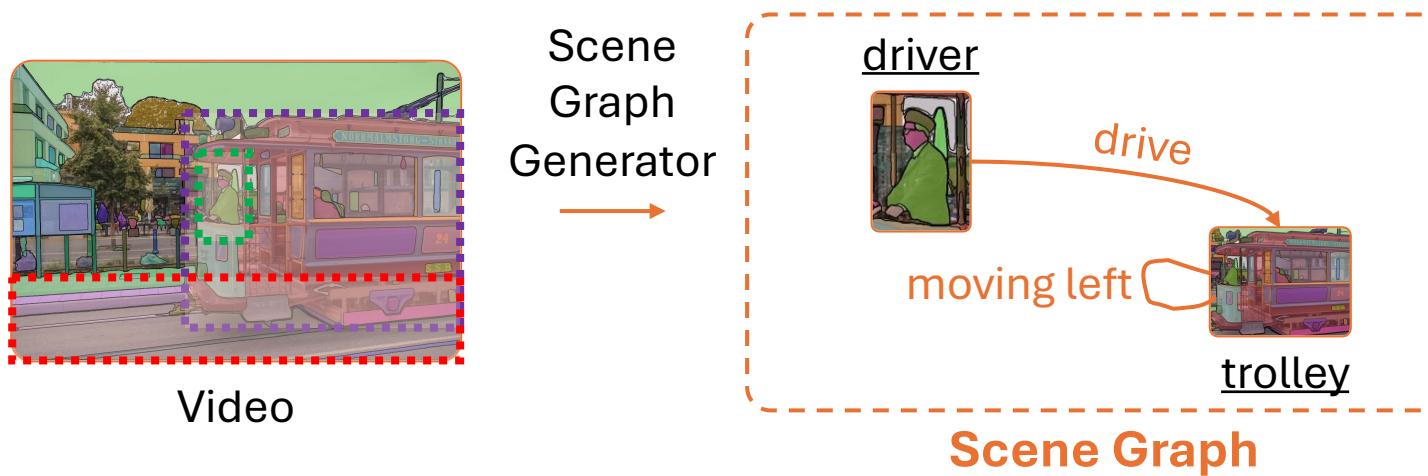
“A driver is driving the
trolley towards the left”

Natural Language

Neurosymbolic Multi-modal Alignment



Aligning Symbolic Representations



“A driver is driving the trolley towards the left”

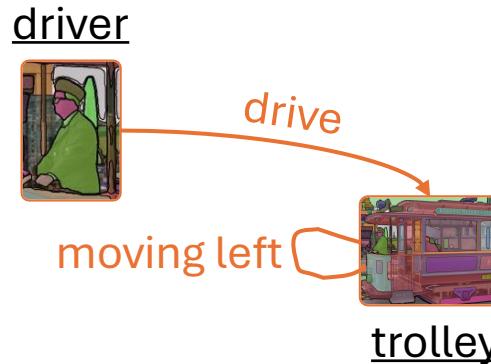
Natural Language

Aligning Symbolic Representations



Video

Scene
Graph
Generator
→



Scene Graph

"A driver is driving the
trolley towards the left"

Language
Model
→

Natural Language

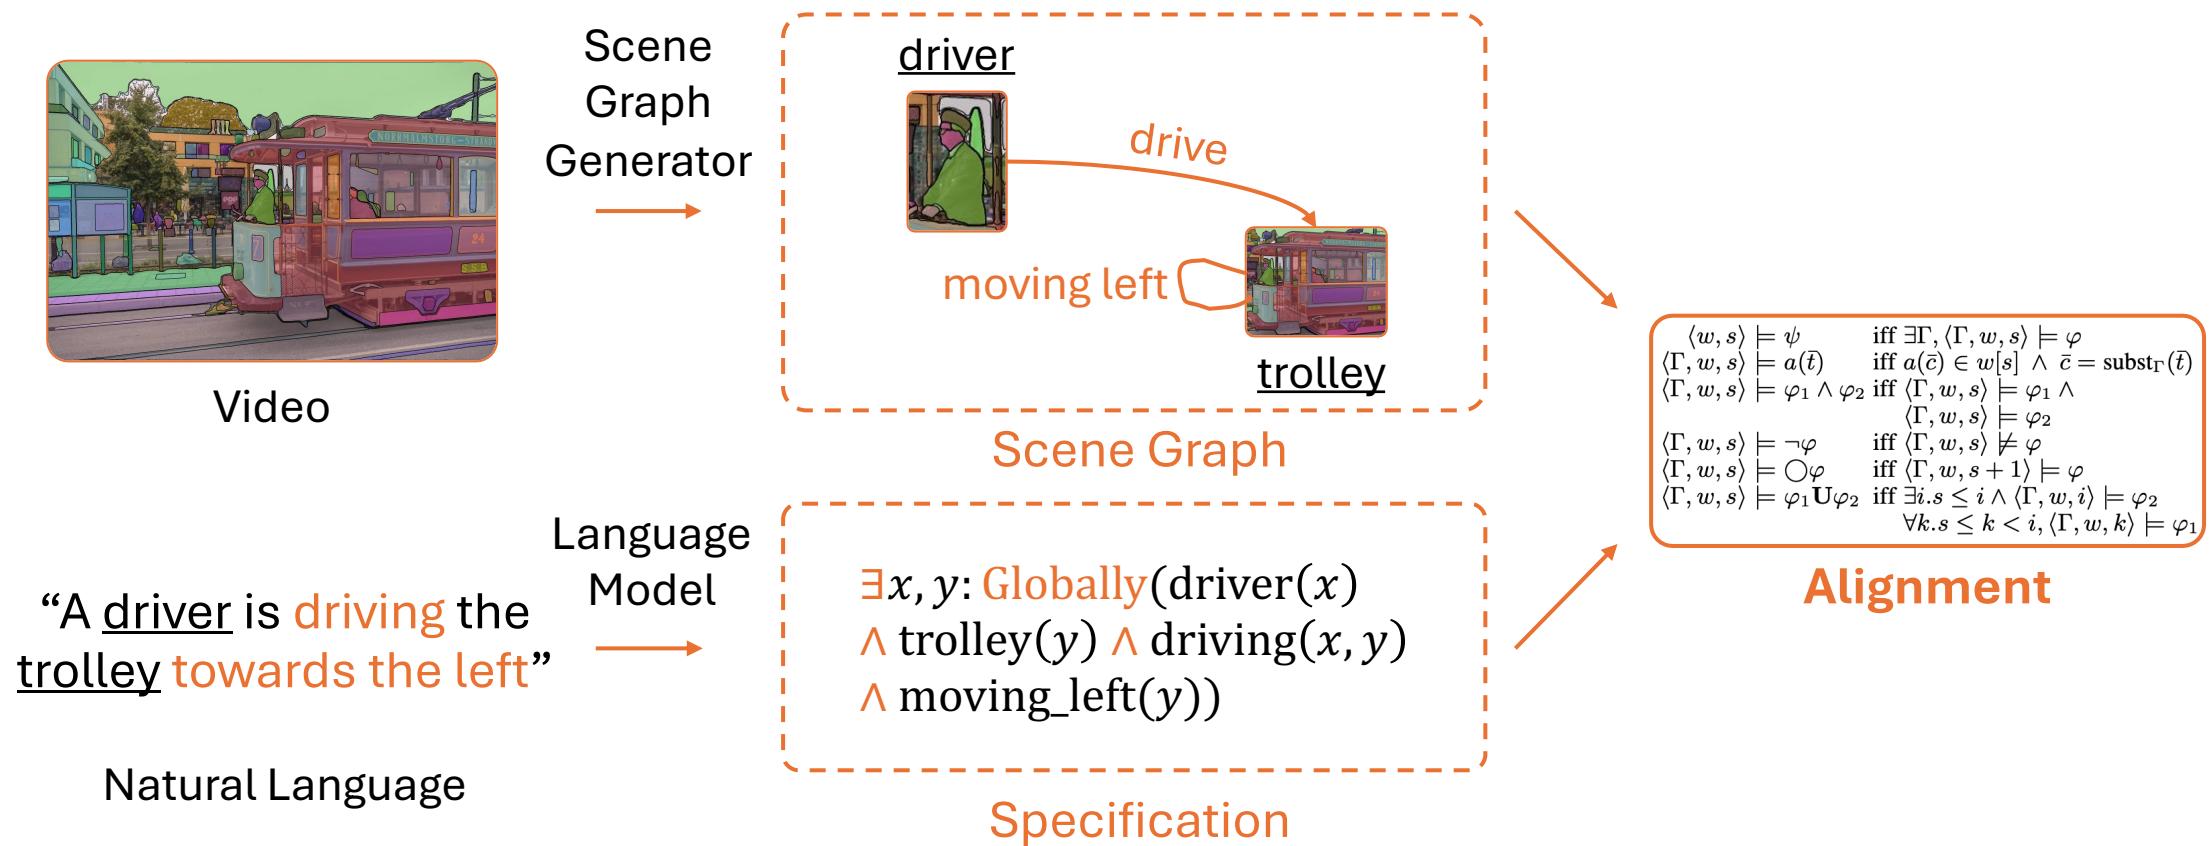
$\exists x, y: \text{Globally}(\text{driver}(x) \wedge \text{trolley}(y) \wedge \text{driving}(x, y) \wedge \text{moving_left}(y))$

Specification

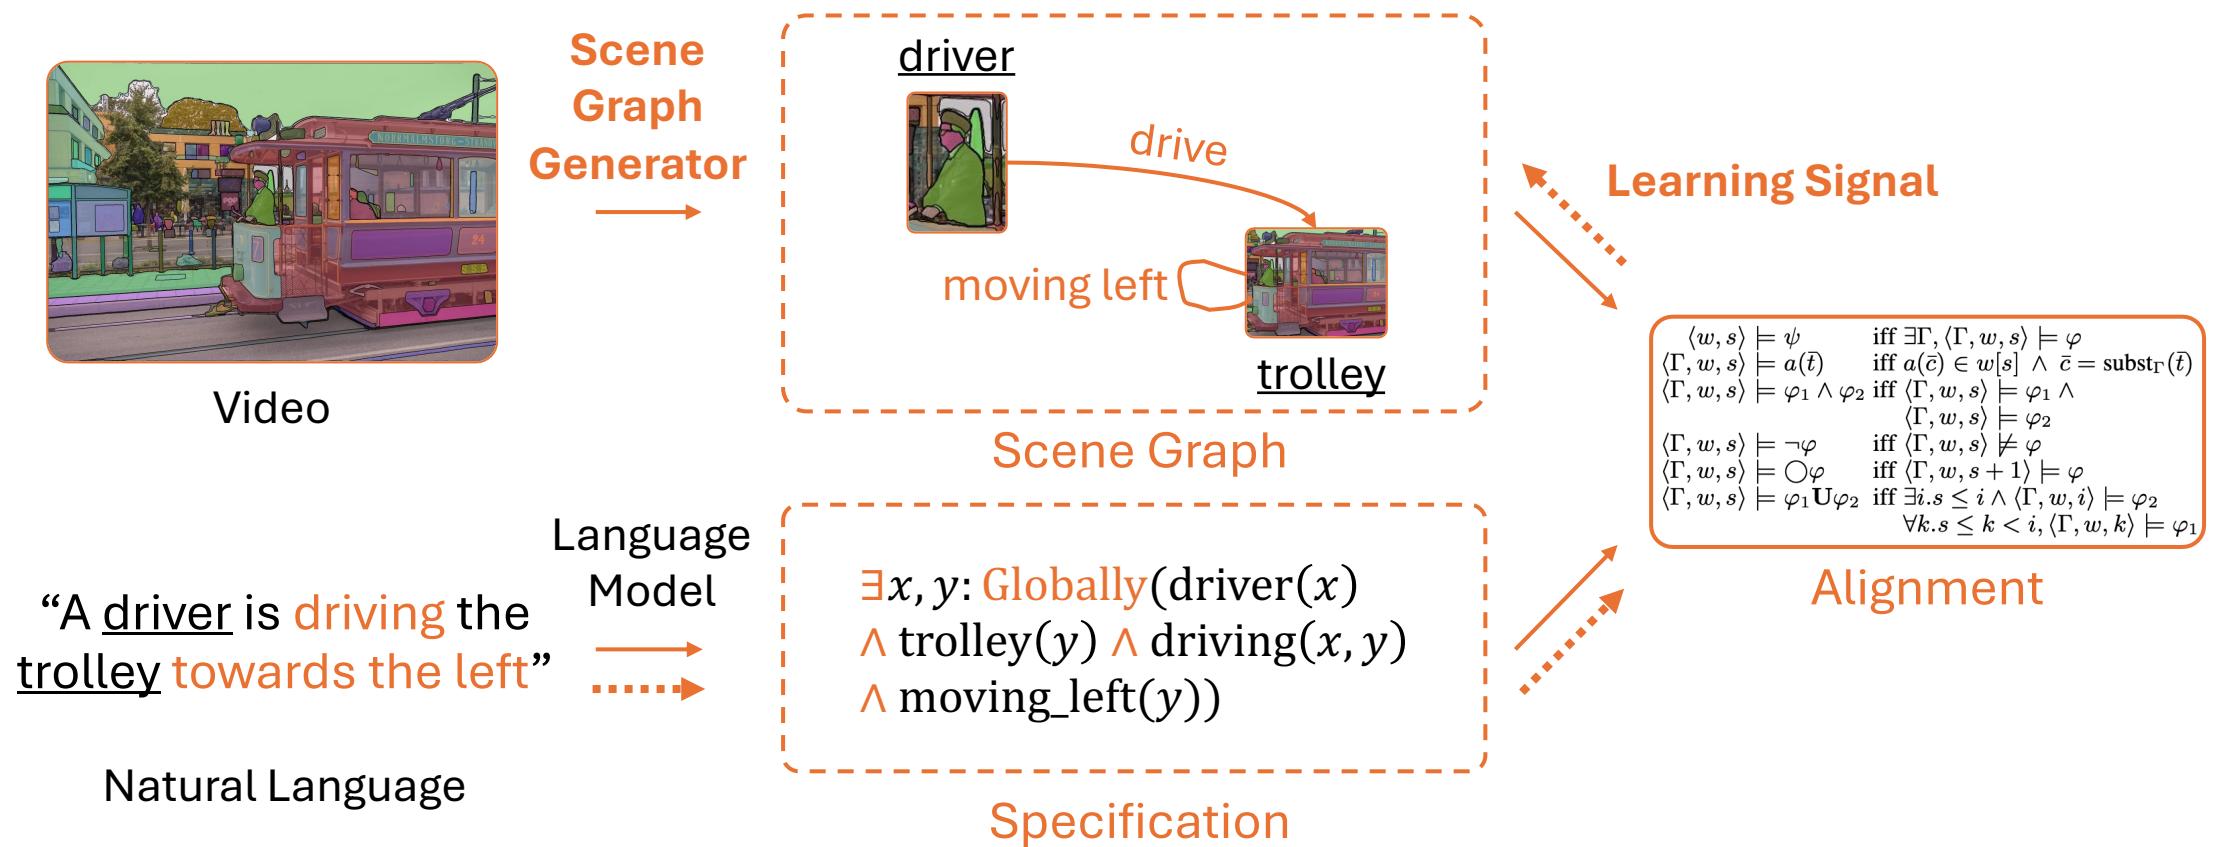
(Formula) $\varphi ::= a(\bar{t}) \mid \varphi_1 \wedge \varphi_2 \mid \varphi_1 \vee \varphi_2 \mid \neg \varphi \mid \bigcirc \varphi \mid \varphi_1 \mathbf{U} \varphi_2 \mid \Box \varphi \mid \Diamond \varphi$
(Specification) $\psi ::= \exists v_1, \dots, v_k, \text{s.t. } \varphi$

Specification Language
Based on Linear Temporal Logic

Aligning Symbolic Representations



Training by Aligning Symbolic Representations



Training by Aligning Symbolic Representations



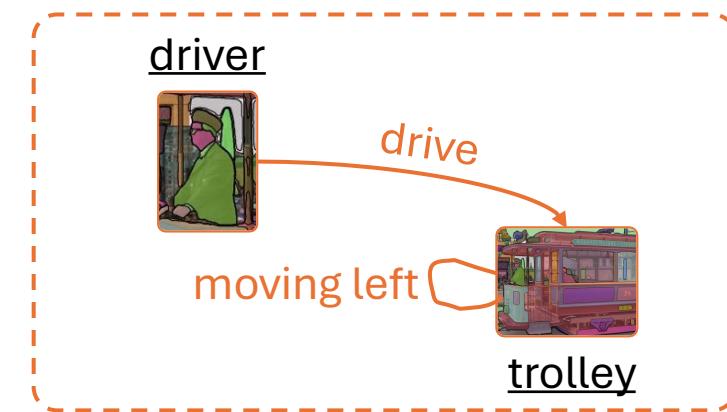
Video

A driver is driving the trolley towards the left

Natural Language

Scene
Graph
Generator

Language
Model



$\exists x, y: \text{Globally}(\text{driver}(x) \wedge \text{trolley}(y) \wedge \text{driving}(x, y) \wedge \text{moving_left}(y))$

Specification

Learning Signal

```
type align(psi: Formula, s: Time)
rel align(Global(p1), s) = max_time(n) and align_all(p1, s, n)
rel align(Finally(p1), s) = align_once(p1, s)
rel align(Until(p1, p2), s) = time(t + 1) and s < (t + 1) and
    align_all(p1, s, t) and align(p2, t + 1)
rel align(Until(p1, p2), s) = align(p2, s)
rel align(Next(p1), s) = align(p1, s + 1) and time(s)
rel align(And(p1, p2), s) = align(p1, s) and align(p2, s)
```

Scallop
Programming
Language



Differentiable Alignment Checker

Specification Language and Alignment



(Term) $t ::= c \mid v$
(Formula) $\varphi ::= a(\bar{t}) \mid \neg a(\bar{t})$
 $\quad \mid \varphi_1 \wedge \varphi_2 \mid \varphi_1 \vee \varphi_2$
 $\quad \mid \Diamond \varphi \mid \varphi_1 \mathbf{U} \varphi_2 \mid \Box \varphi \mid \Diamond \varphi$
(Specification) $\psi ::= \exists v_1, \dots, v_k, \text{s.t. } \varphi$

```
type Term = Constant(String) | Var(VarId)  
type Atom = UnaryAtom(String, Term) | BinaryAtom(String, Term, Term)  
type Spec = Atom(Atom) | NotAtom(Atom)  
           | And(Spec, Spec) | Or(Spec, Spec)  
           | Global(Spec) | Until(Spec, Spec) | Next(Spec) | Finally(Spec)
```

Specification Language and Alignment

$$\langle \Gamma, w, [i, j] \rangle \models \varphi$$

type align(phi: Spec, start: Time, end: Time)



Specification Languages and Alignment

Evaluation Context

Context contains variable assignments that is stored in Scallop's intentional database (IDB)

Spatio-Temporal Scene Graph

NN predicted scene graph is encoded as probabilistic facts within Scallop's extensional database (EDB)

$$\langle \Gamma, w, [i, j] \rangle \models \varphi$$

`type align(phi: Spec, start: Time, end: Time)`



Specification Language and Alignment

$$\langle \Gamma, w, [i, j] \rangle \models a(\bar{t}) \quad \text{iff } \forall f \in [i, j], a(\bar{c}) \in w[f] \wedge \bar{c} = \text{substitute}_\Gamma(\bar{t})$$

```
rel align(UnaryAtom(a, t), i, i)      = sg_unary(i, a, c) and subst_term(t, c)
rel align(BinaryAtom(a, t1, t2), i, i) = sg_binary(i, a, c1, c2) and
                                         subst_term(t1, c1) and subst_term(t2, c2)
rel align(phi, i, k)                  = align(phi, i, j) and align(phi, j + 1, k)
```


$$\langle \Gamma, w, [i, j] \rangle \models \varphi_1 \mathbf{U} \varphi_2 \text{ iff } \exists k, \langle \Gamma, w, [i, k] \rangle \models \varphi_1 \wedge \langle \Gamma, w, [k, j] \rangle \models \varphi_2$$

```
rel align(Until(phi_1, phi_2), i, j) = align(phi_1, i, k) and align(phi_2, k + 1, j)
```



Probabilistic and Differentiable Alignment



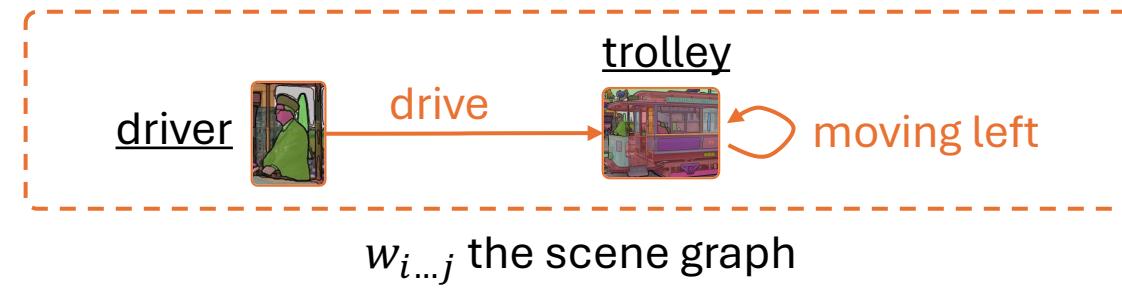
$$\langle w, [i, j] \rangle \models \psi \rightarrow \Pr(\langle w, [i, j] \rangle \models \psi) \rightarrow \langle \Pr(\langle w, [i, j] \rangle \models \psi), \nabla \Pr(\langle w, [i, j] \rangle \models \psi) \rangle$$

Discrete Alignment

Probabilistic Alignment

Differentiable Alignment

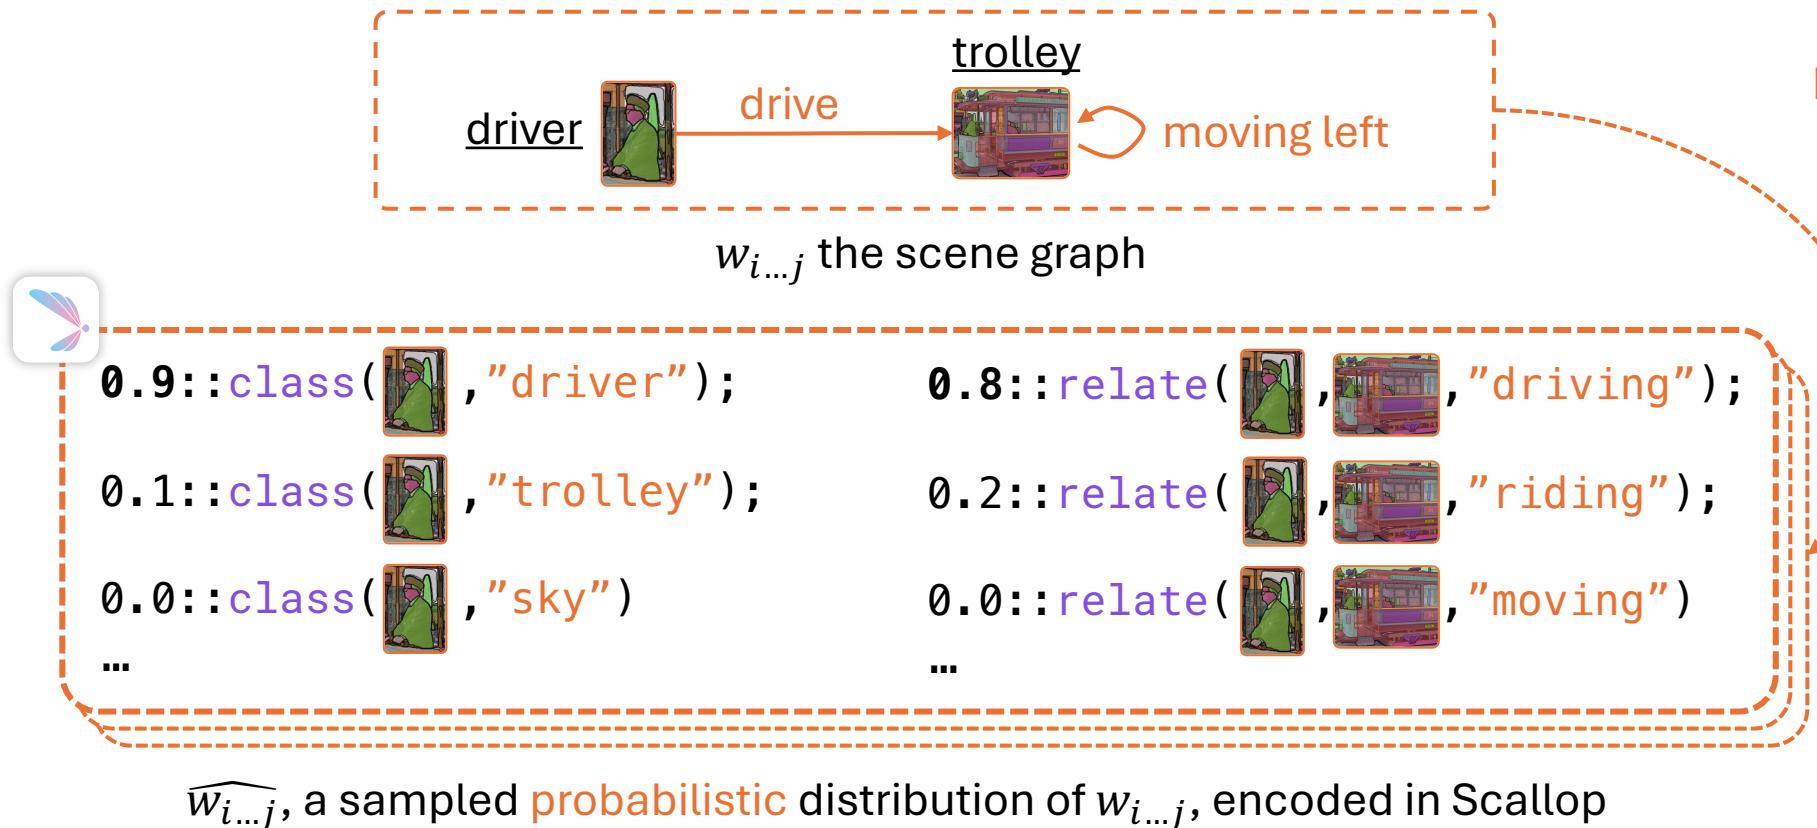
Probabilistic Alignment



$$\exists d, t. \bigcirc \text{driver}(d) \wedge \text{trolley}(t) \wedge \text{driving}(d, t)$$

ψ : throughout the video, driver is driving the trolley

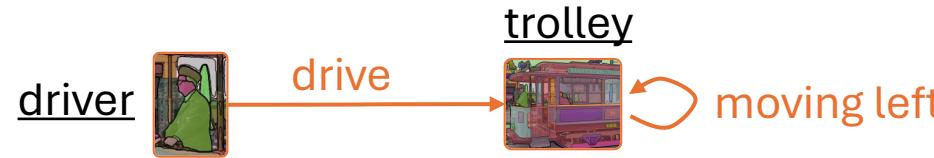
Probabilistic Alignment



$$\exists d, t. \bigcirc \text{driver}(d) \wedge \text{trolley}(t) \wedge \text{driving}(d, t)$$

ψ : throughout the video, driver is driving the trolley

Probabilistic Alignment



$w_{i \dots j}$ the scene graph

0.9::class(, "driver"); 0.8::relate(, , "driving");
0.1::class(, "trolley"); 0.2::relate(, , "riding");
0.0::class(, "sky") 0.0::relate(, , "moving")
...
...

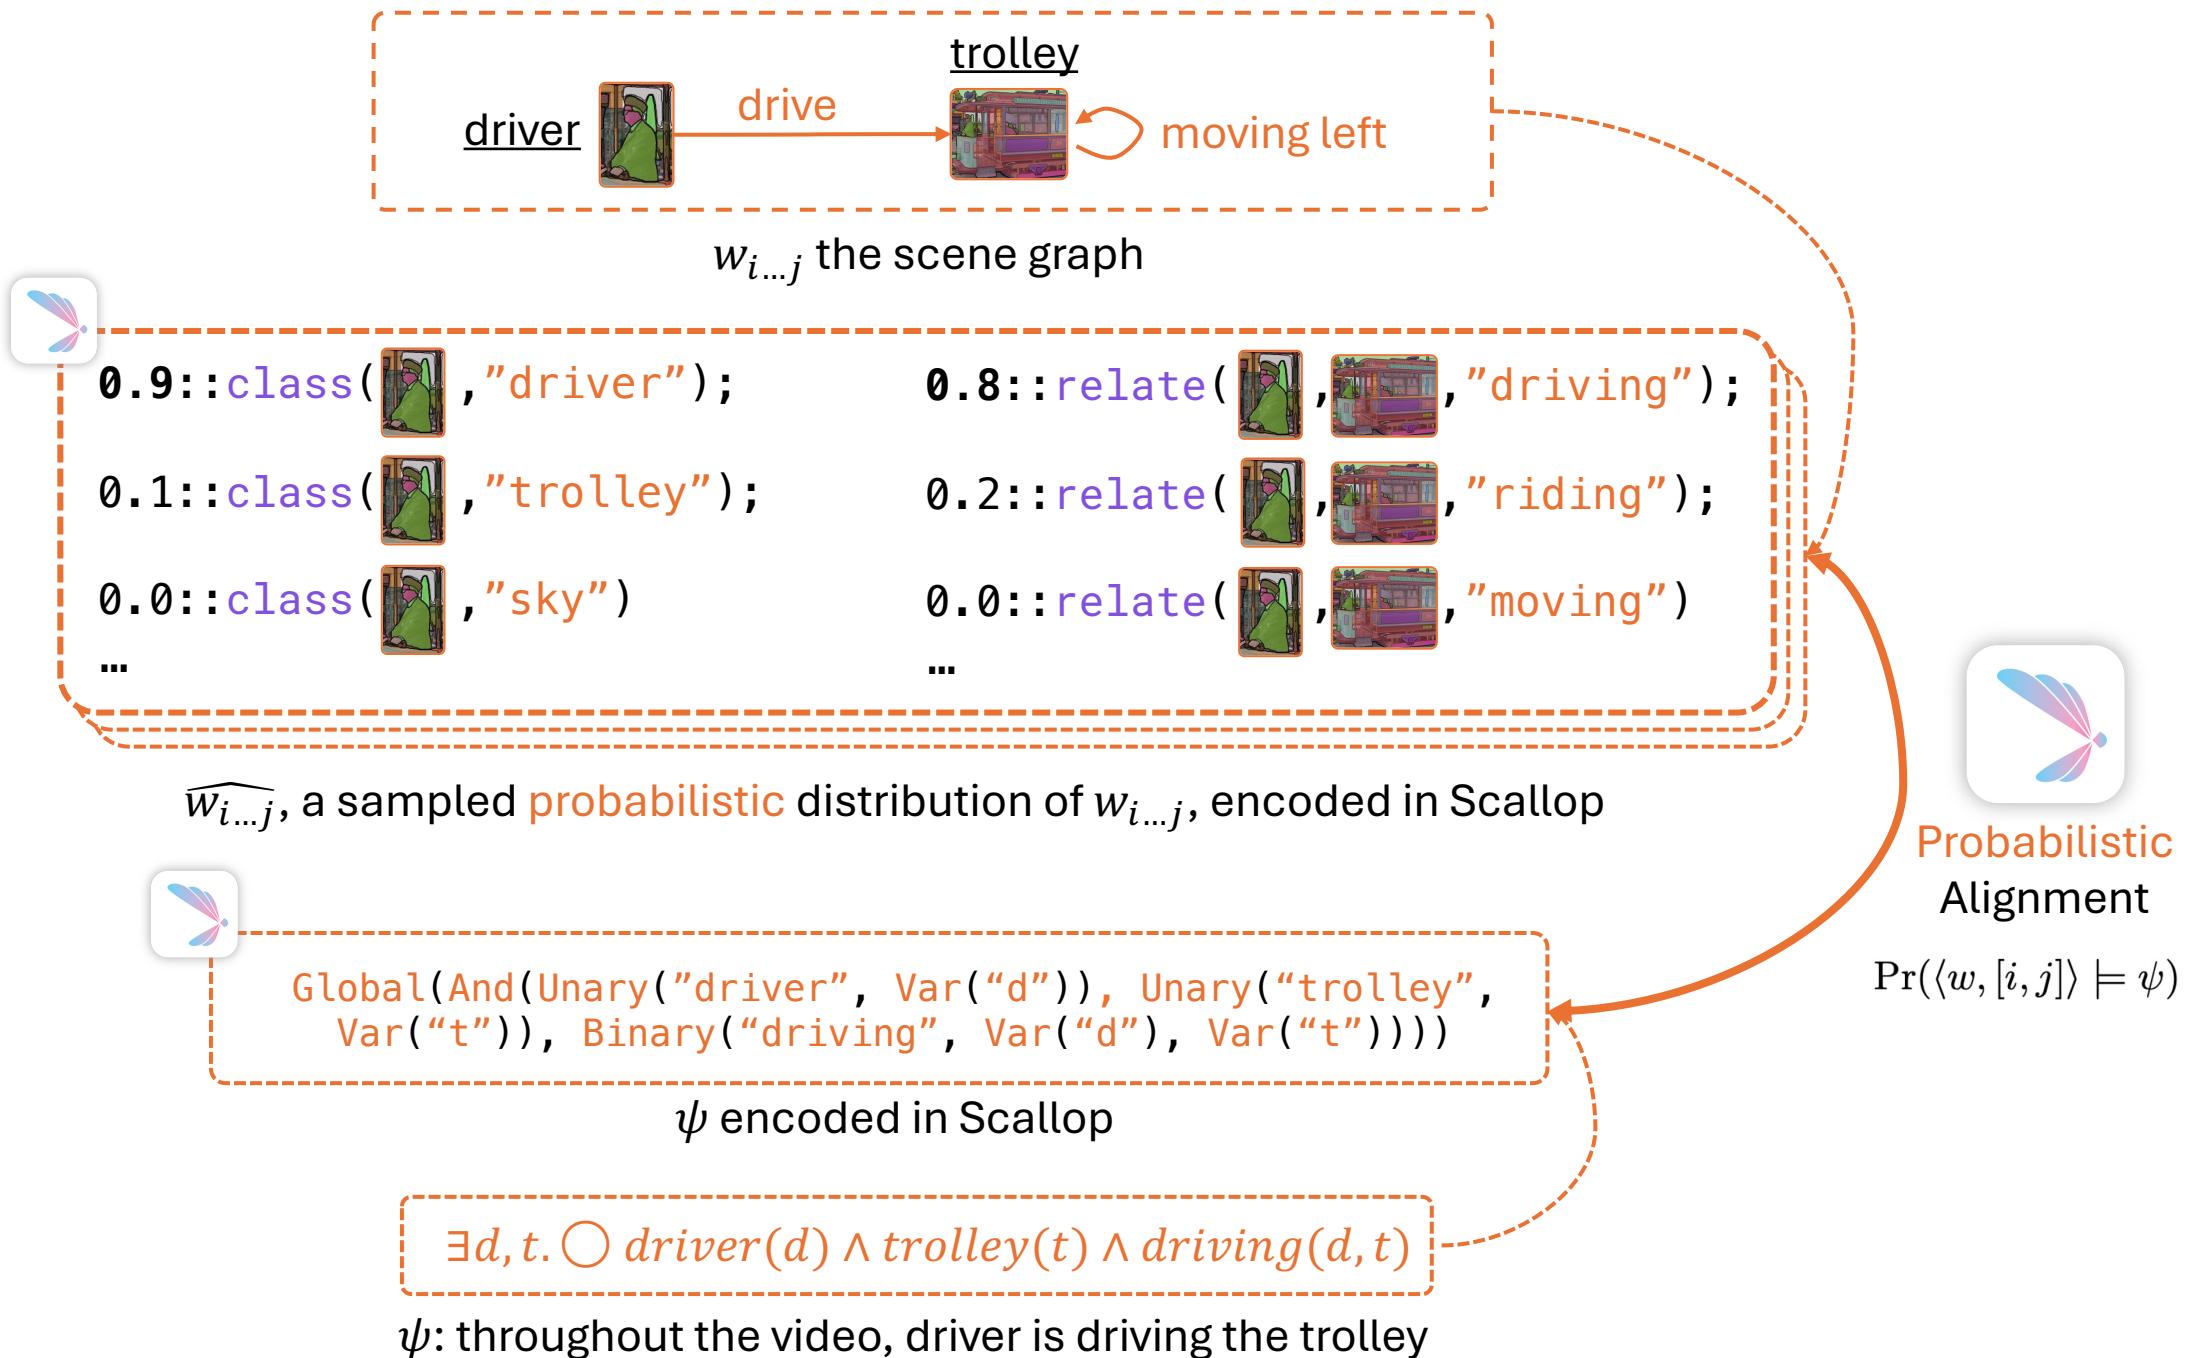
$\widehat{w}_{i \dots j}$, a sampled probabilistic distribution of $w_{i \dots j}$, encoded in Scallop

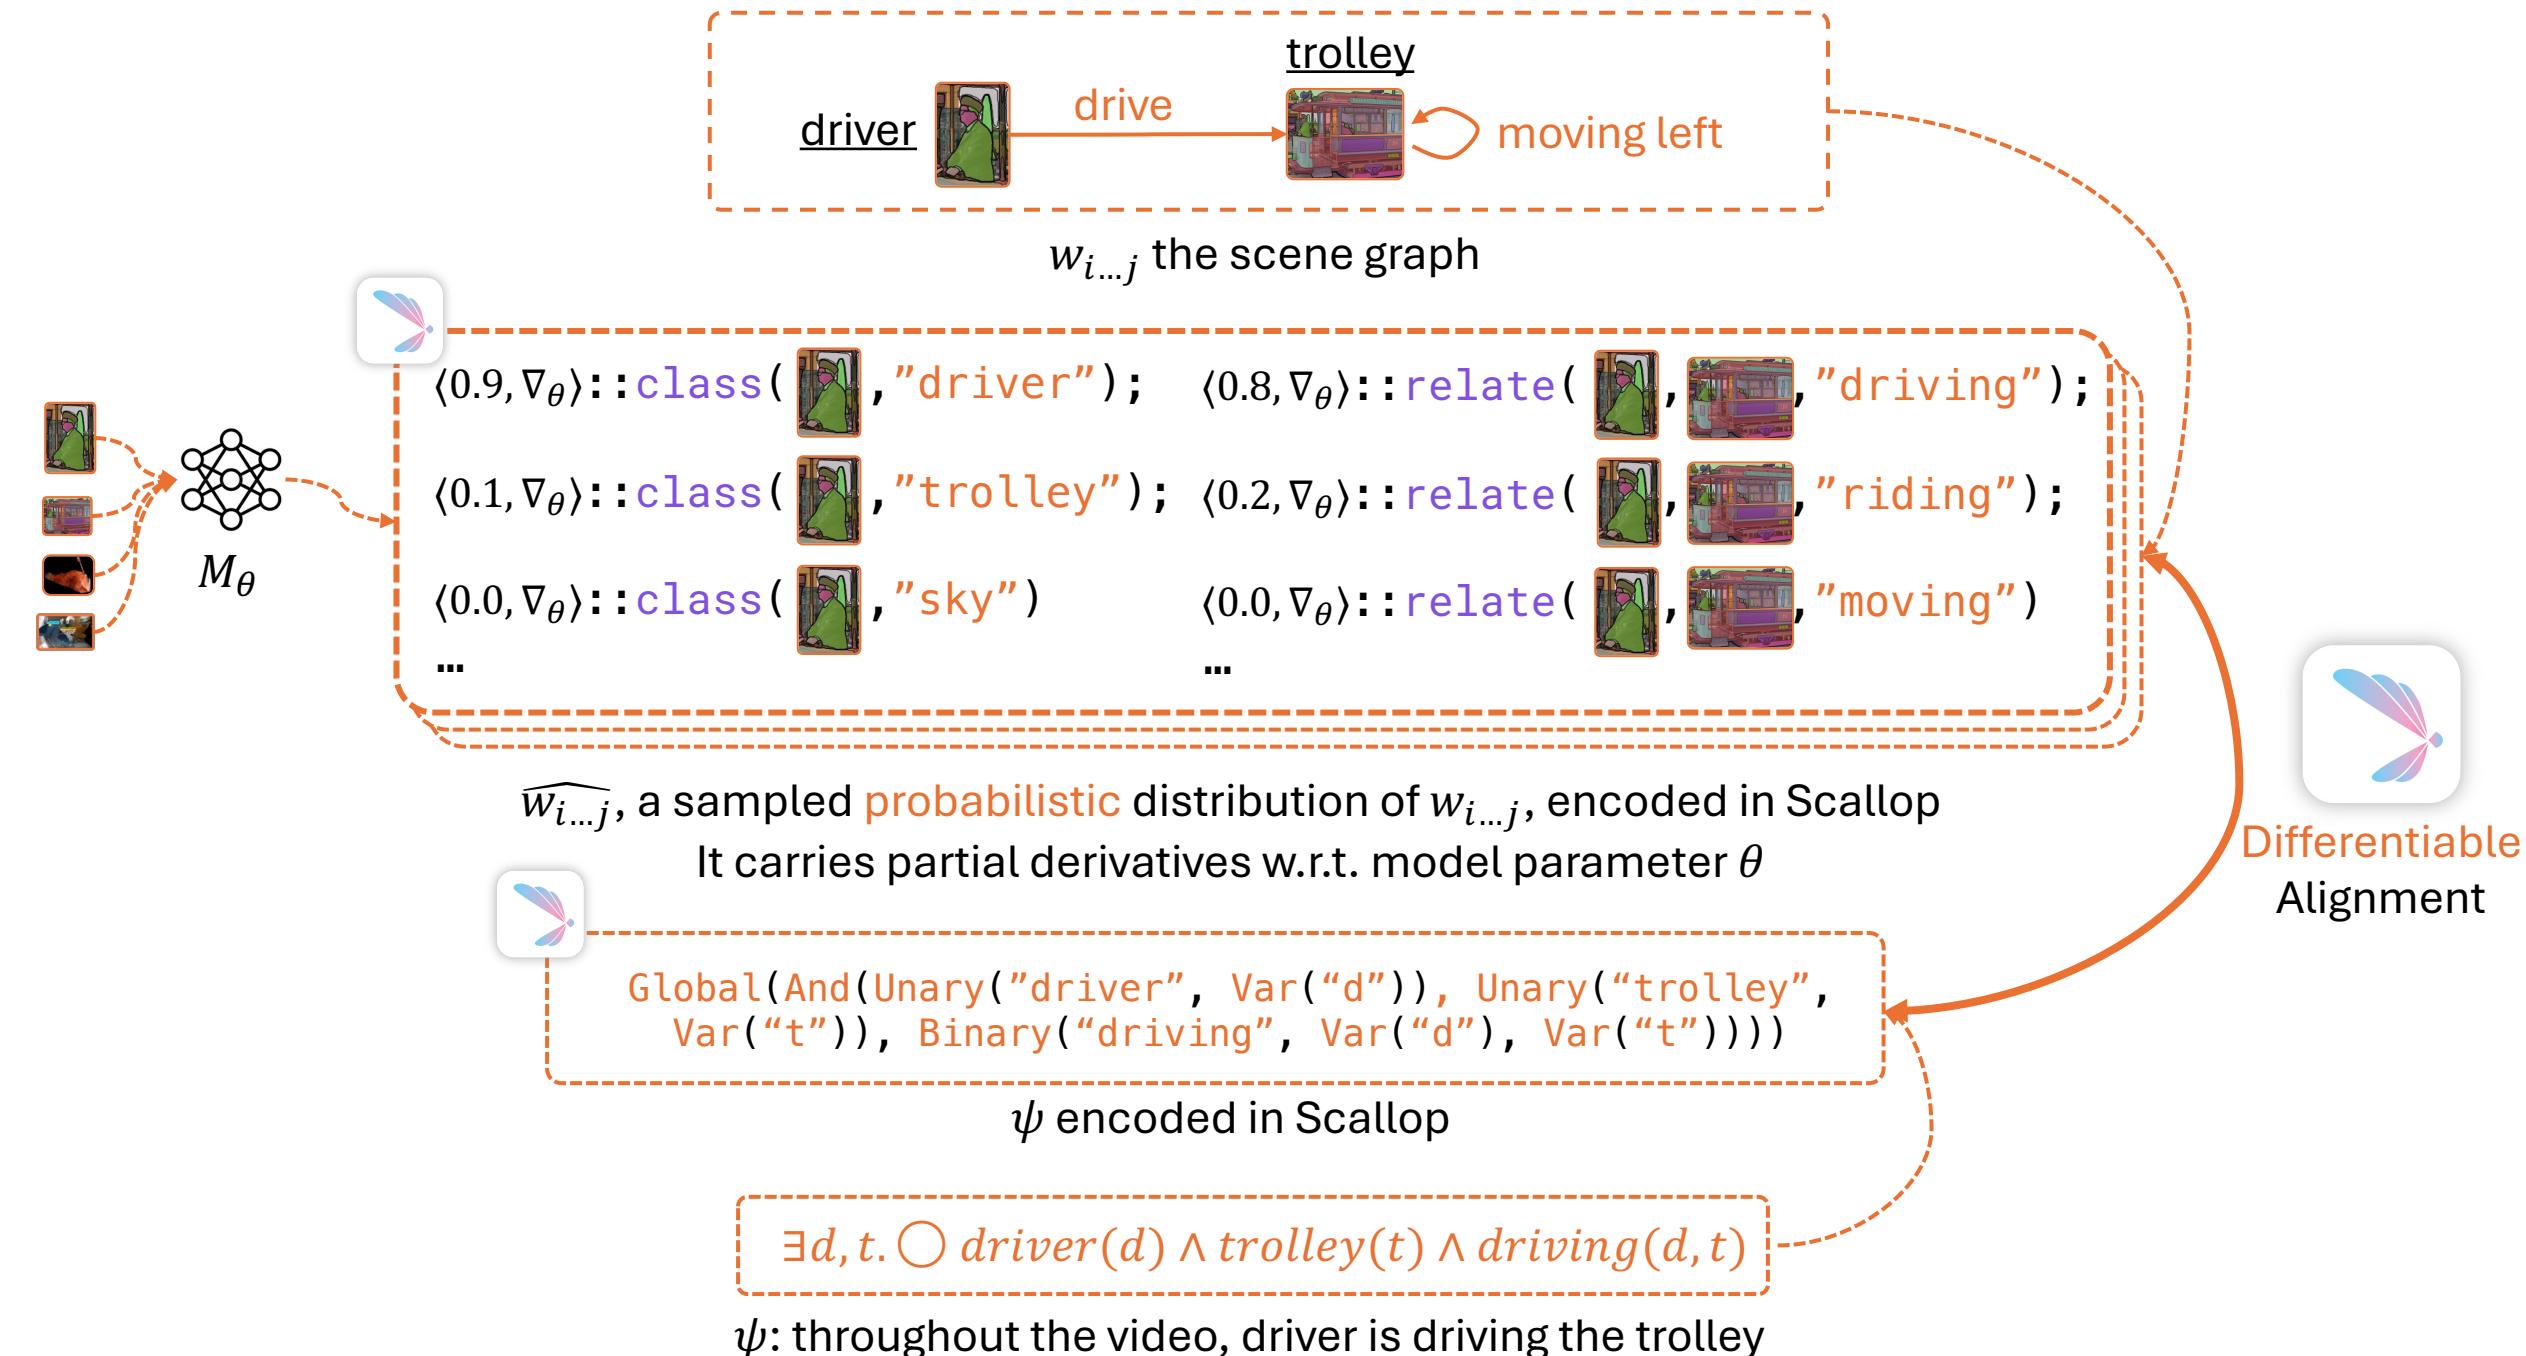
Global(And(Unary("driver", Var("d")), Unary("trolley", Var("t")), Binary("driving", Var("d"), Var("t"))))

ψ encoded in Scallop

$\exists d, t. \bigcirc \text{driver}(d) \wedge \text{trolley}(t) \wedge \text{driving}(d, t)$

ψ : throughout the video, driver is driving the trolley





Probabilistic and Differentiable Alignment



$$\langle w, [i, j] \rangle \models \psi \rightarrow \Pr(\langle w, [i, j] \rangle \models \psi) \rightarrow \langle \Pr(\langle w, [i, j] \rangle \models \psi), \nabla \Pr(\langle w, [i, j] \rangle \models \psi) \rangle$$

Discrete Alignment

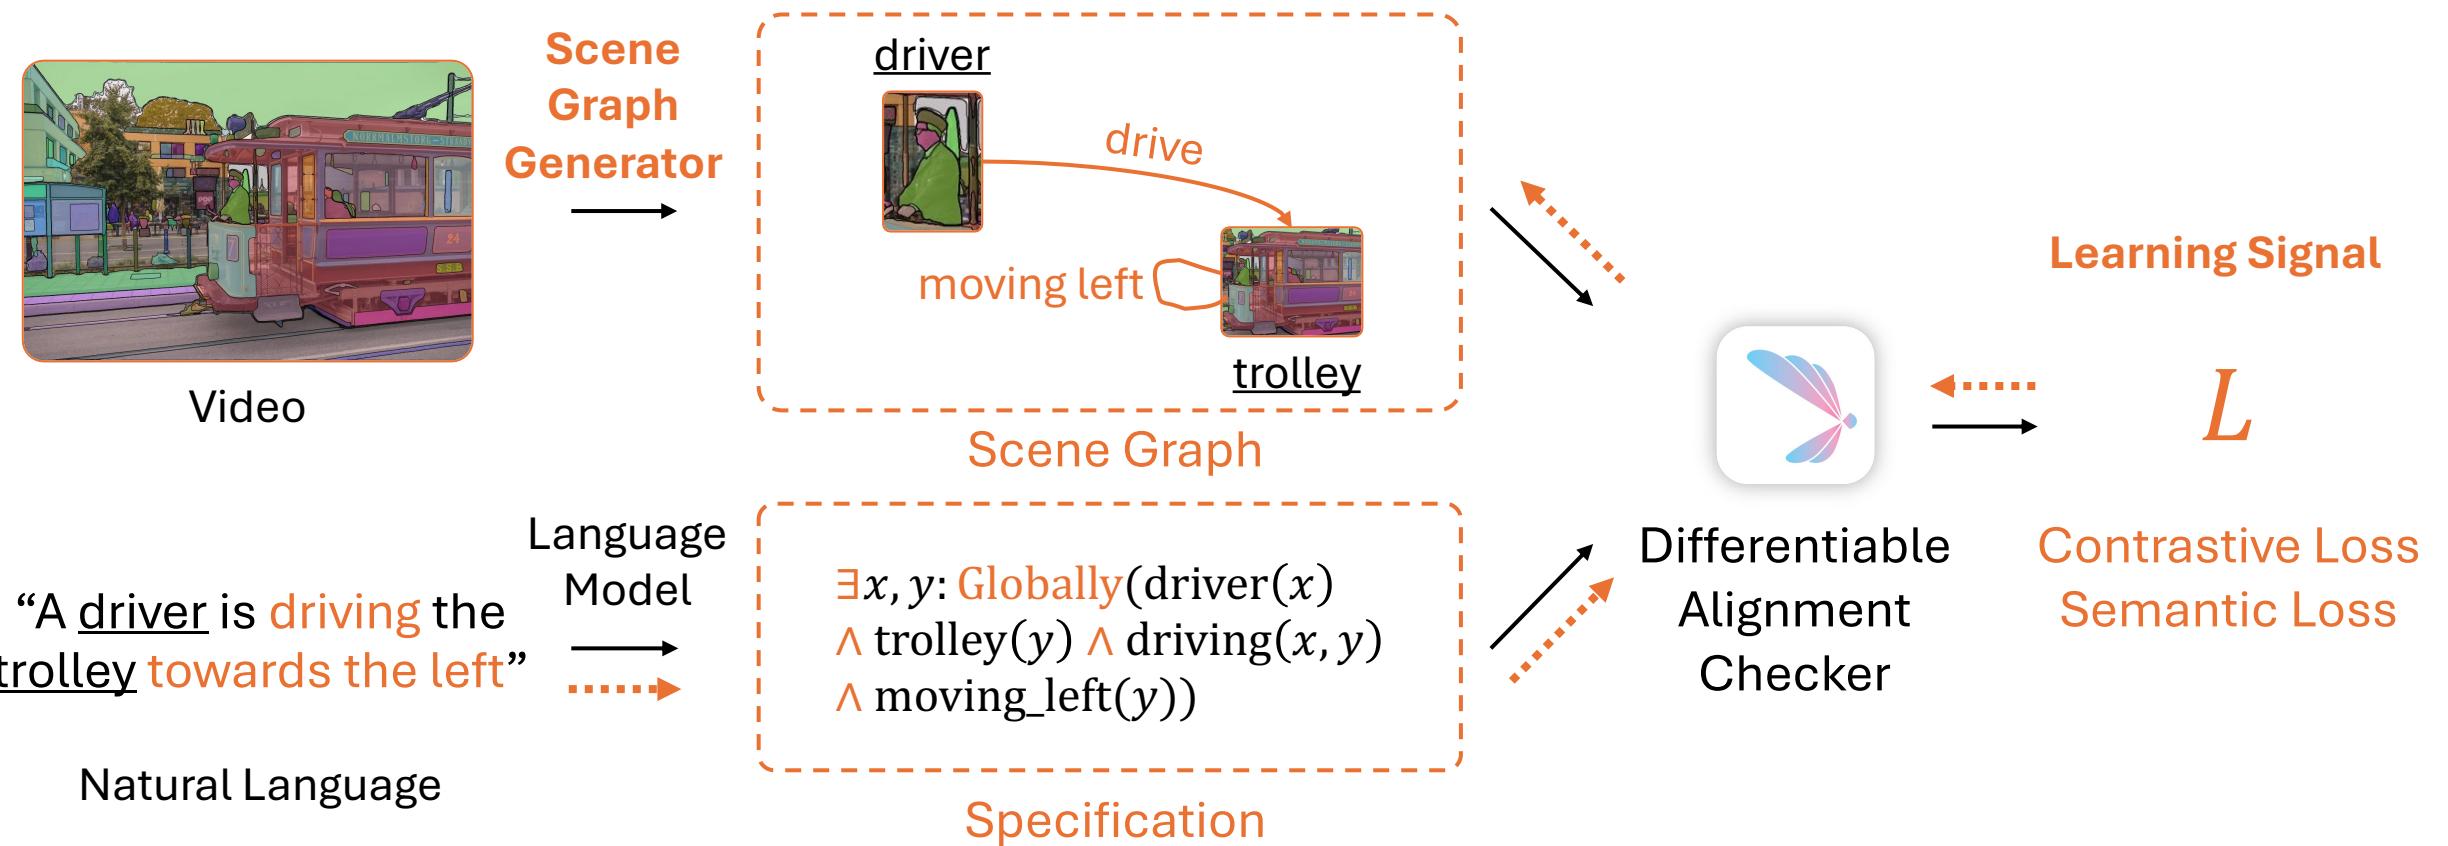
Probabilistic Alignment

Differentiable Alignment

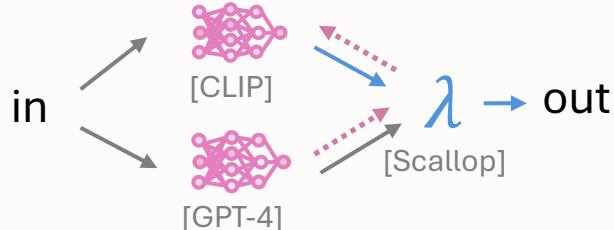
(Tag Space)	$t \in T$
(False)	$\mathbf{0} \in T$
(True)	$\mathbf{1} \in T$
(Disjunction)	$\oplus : T \times T \rightarrow T$
(Conjunction)	$\otimes : T \times T \rightarrow T$
(Negation)	$\ominus : T \rightarrow T$
(Saturation)	$\Theta : T \times T \rightarrow \text{Bool}$

Provenance Framework

Training by Aligning Symbolic Representations



Results

Method	Unary Acc. / Binary R@10
 <p>[IPS-Conv/VPS-Conv]</p>	15.1% / 23.4%
 <p>[CLIP] [GPT-4]</p> <p>λ</p> <p>[Scallop]</p>	27.8% / 54.0%

Trained and Evaluated on Dataset *OpenPVSG*.

Results

Method (with LASER)		Unary			Binary		
		R@1	R@5	R@10	R@1	R@5	R@10
VIOLET	Base	0.0660	0.1855	0.2983	0.0460	0.1307	0.2636
	Fine-tuned	0.0878	0.2574	0.3463	0.0501	0.2028	0.3451
	Incr.	↑ 0.0218	↑ 0.0719	↑ 0.0480	↑ 0.0041	↑ 0.0721	↑ 0.0815
SigLIP	Base	0.0000	0.0179	0.0483	0.0000	0.0362	0.1667
	Fine-tuned	0.1467	0.2627	0.3152	0.0347	0.1624	0.3012
	Incr.	↑ 0.1467	↑ 0.2448	↑ 0.2669	↑ 0.0347	↑ 0.1262	↑ 0.1345
CLIP	Base	0.1633	0.3381	0.4404	0.0197	0.0673	0.0988
	Fine-tuned	0.2778	0.5231	0.6402	0.1482	0.4214	0.5398
	Incr.	↑ 0.1145	↑ 0.1850	↑ 0.1998	↑ 0.1284	↑ 0.3540	↑ 0.4410

Table 1: We show the performance improvements of base backbone models and their fine-tuned version, on the R@k metrics of unary and binary predicate prediction. As shown by the increments, LASER’s weak supervisory learning framework significantly enhances all three models’ performance on the STSG extraction tasks.

Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning
Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

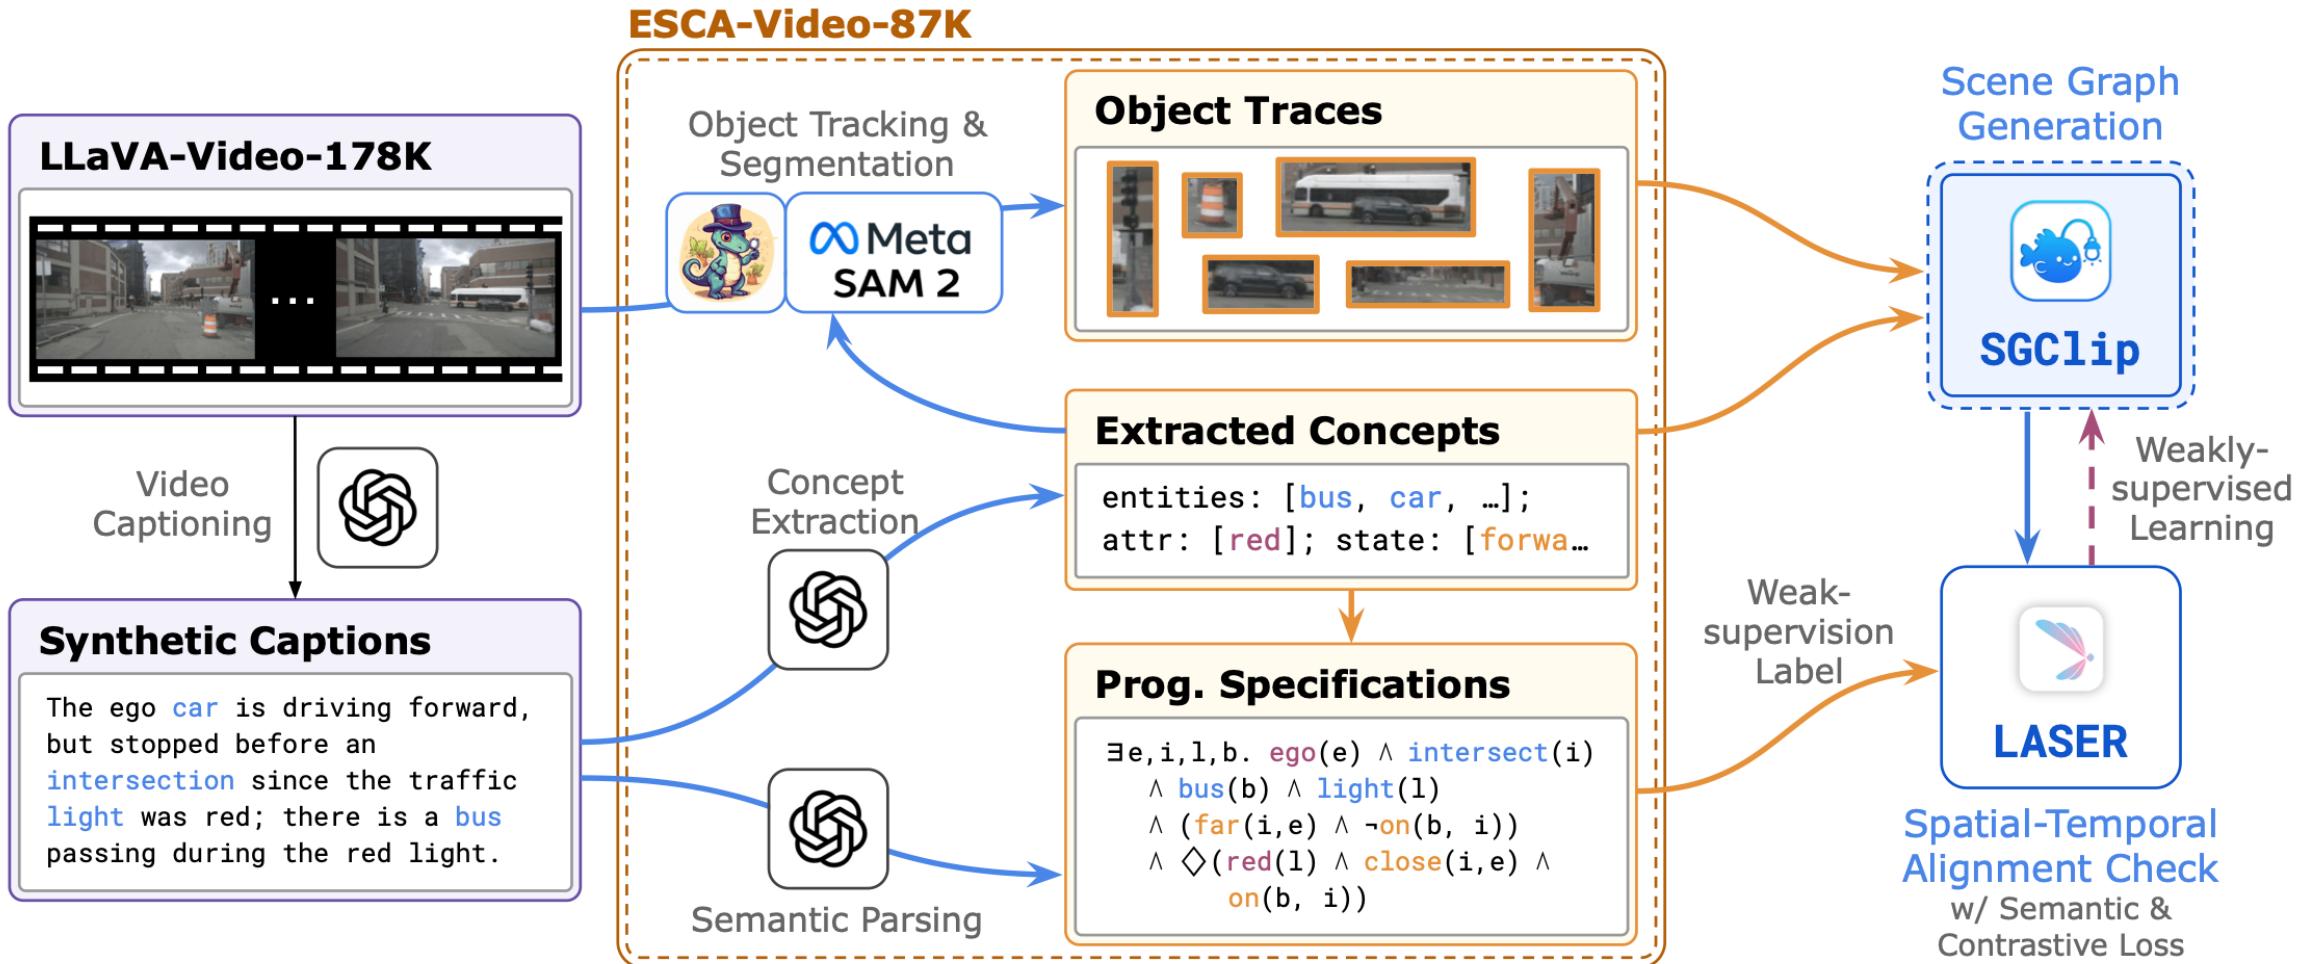
ESCA: Contextualizing Embodied Agents via
Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

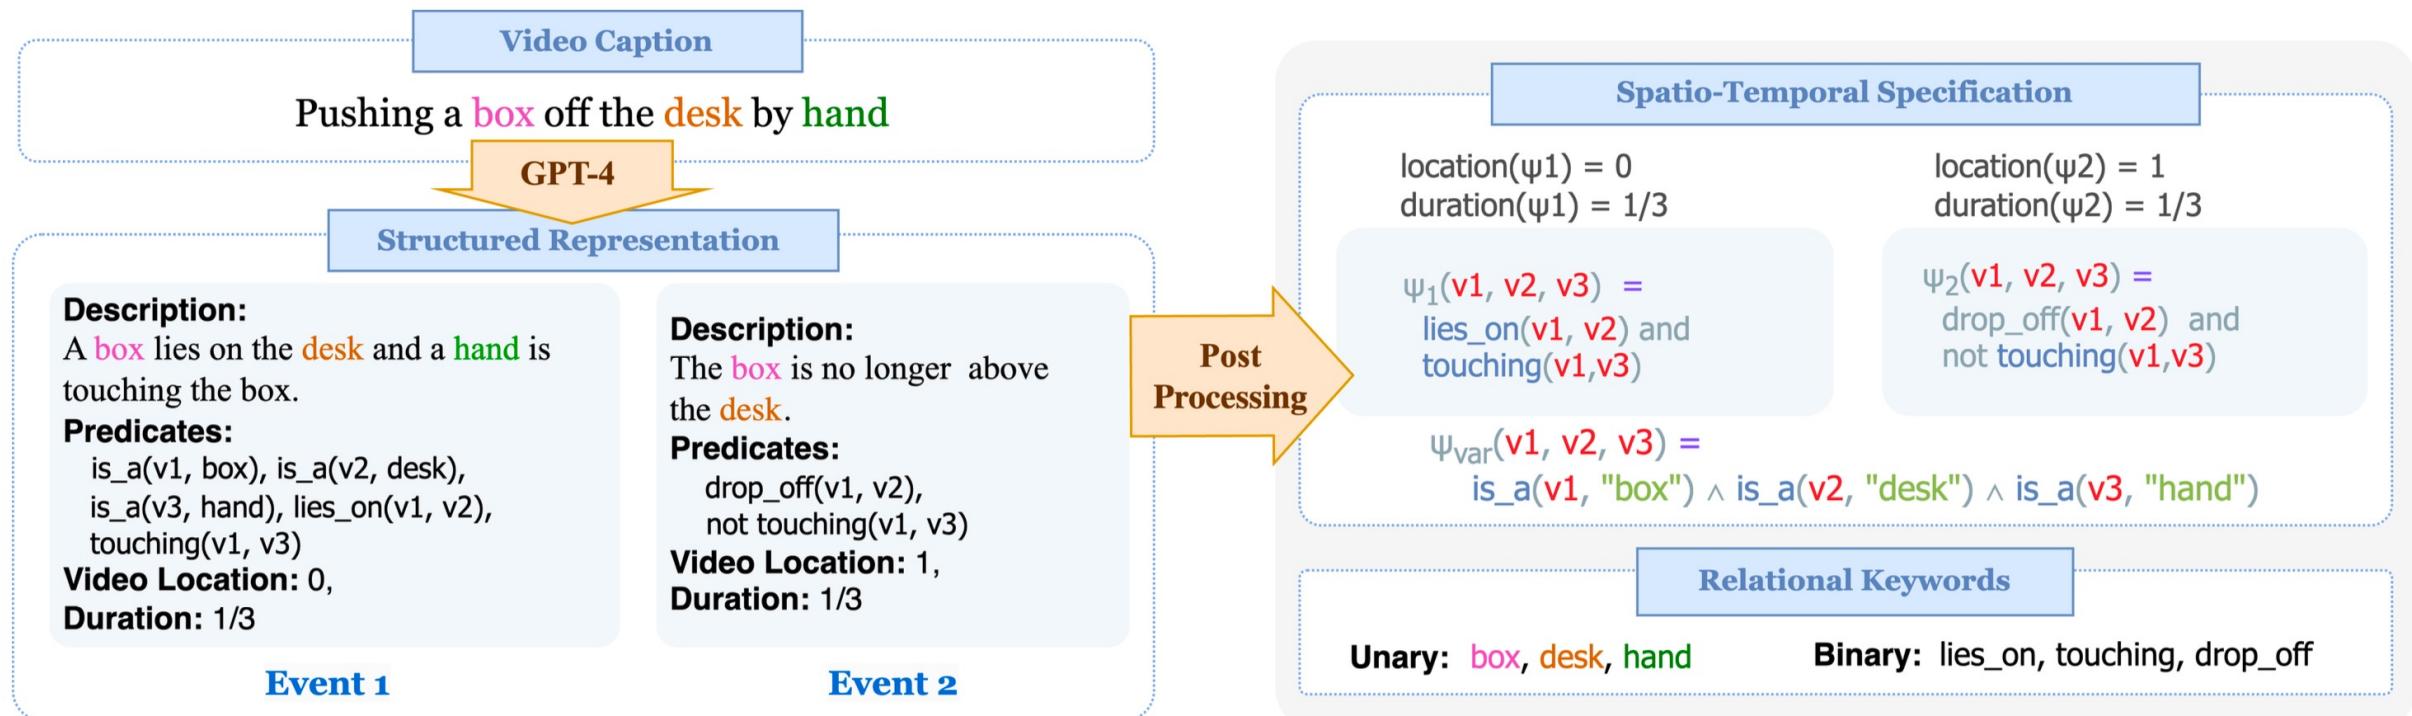


ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

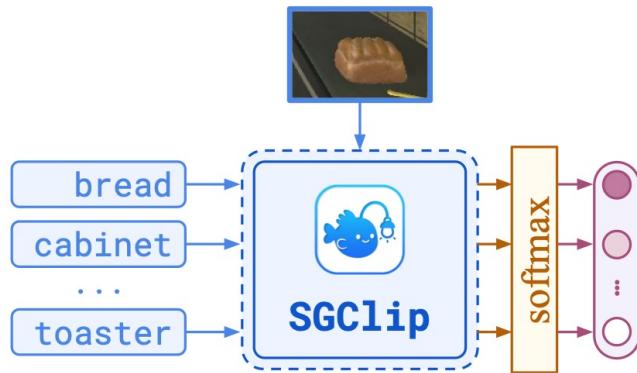
Scaling Up Alignment Training Pipeline



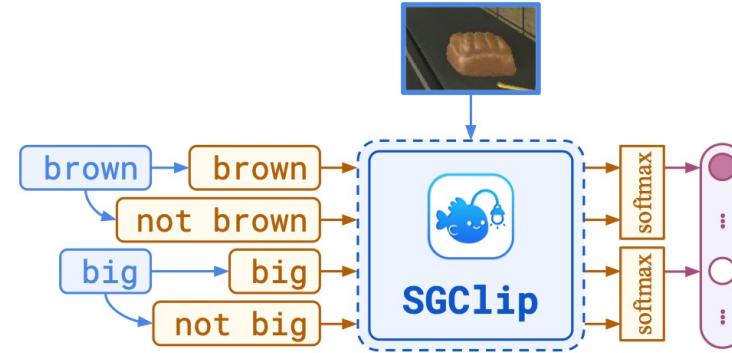
Scaling Up Alignment Training Pipeline



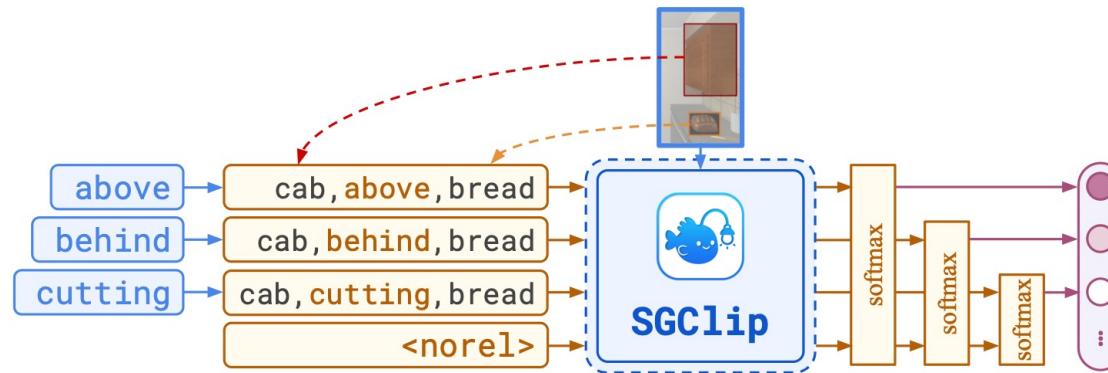
Unified Open-Vocabulary Interface



(a) Entity classes



(b) Attributes

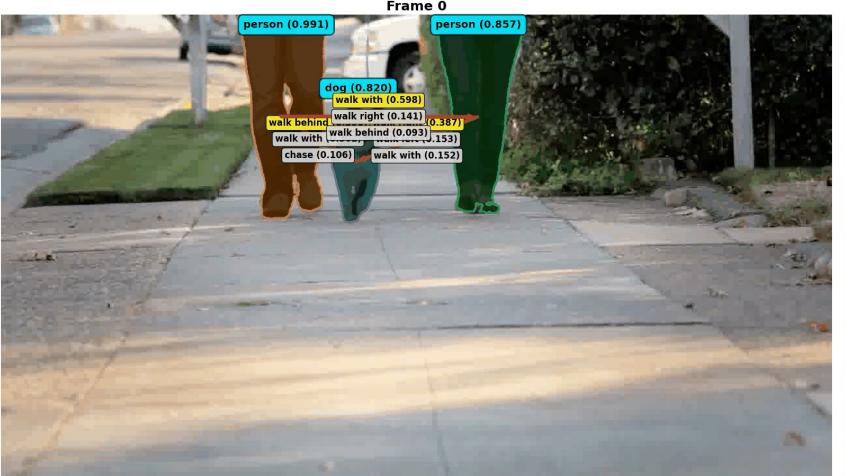


(c) Binary relations

Promptable STSG Generation Pipeline

- Model-driven self-supervised training
 - Eliminating the need for fine-grained STSG label
 - Scaling up training for SGClip
- Unified open-vocabulary interface
 - Can simultaneously handle entity classes, attributes, actions, spatial relations, interactions
 - Zero-shot generalizable to diverse scenarios such as action localization
- Controllable scene graph generation granularity
 - Granularity is controlled via prompting the caption generation and concept extraction

A couple walking their dog



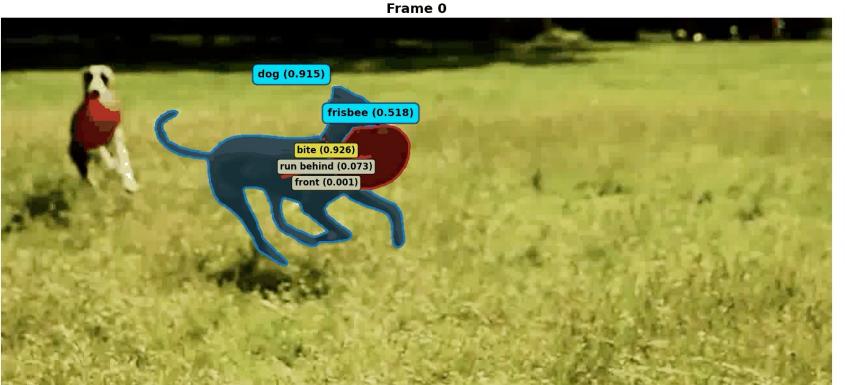
A dog runs around a turtle

A car quickly drives pass a person



A person riding on a horse

A dog playing with a frisbee



A dog jumps left on the sofa to play with a woman

Application to VLM-based Embodied Agents

Instruction:
🧠 ➡️ Stack the maroon triangular prism and the olive triangular prism in sequence.



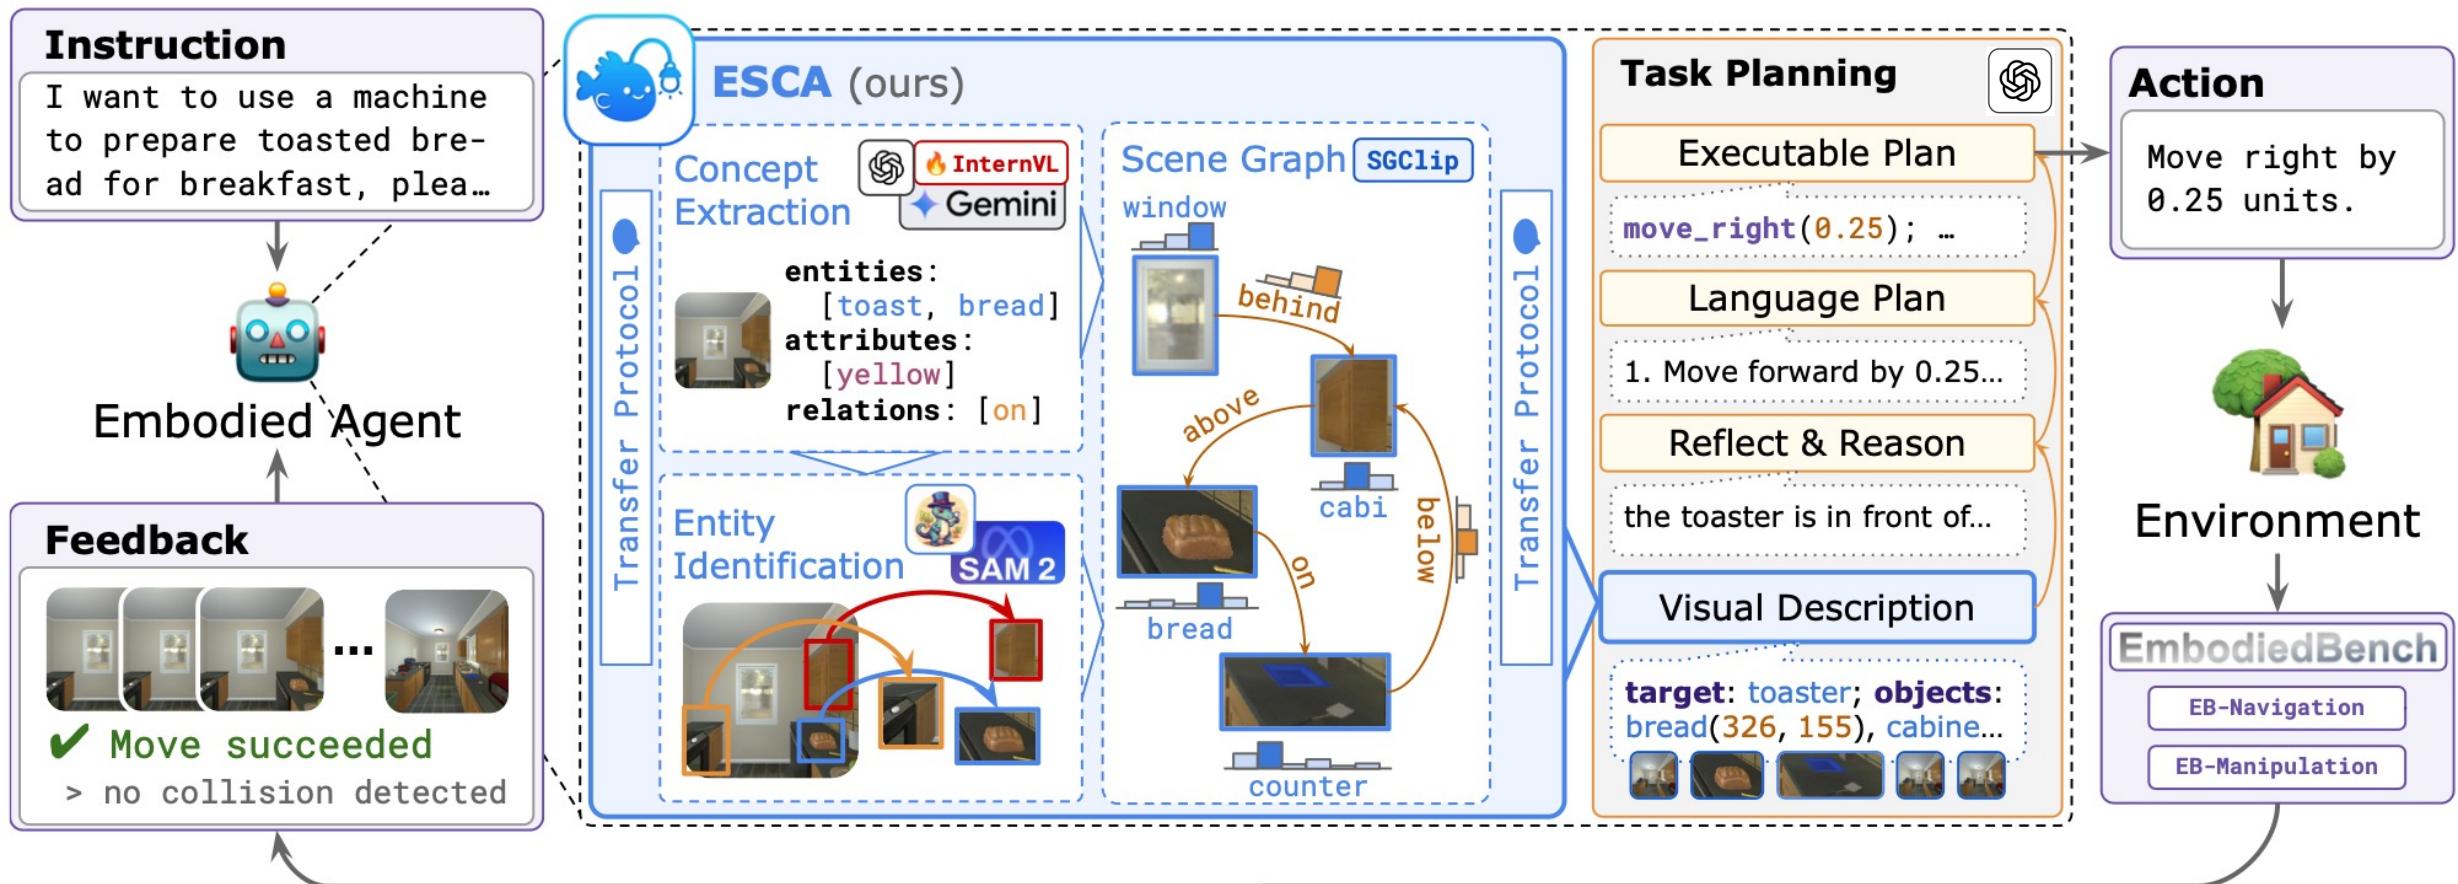
EB-Manipulation

Instruction:
🧠 ➡️ Navigate to the Toaster in the room and be as close as possible to it.



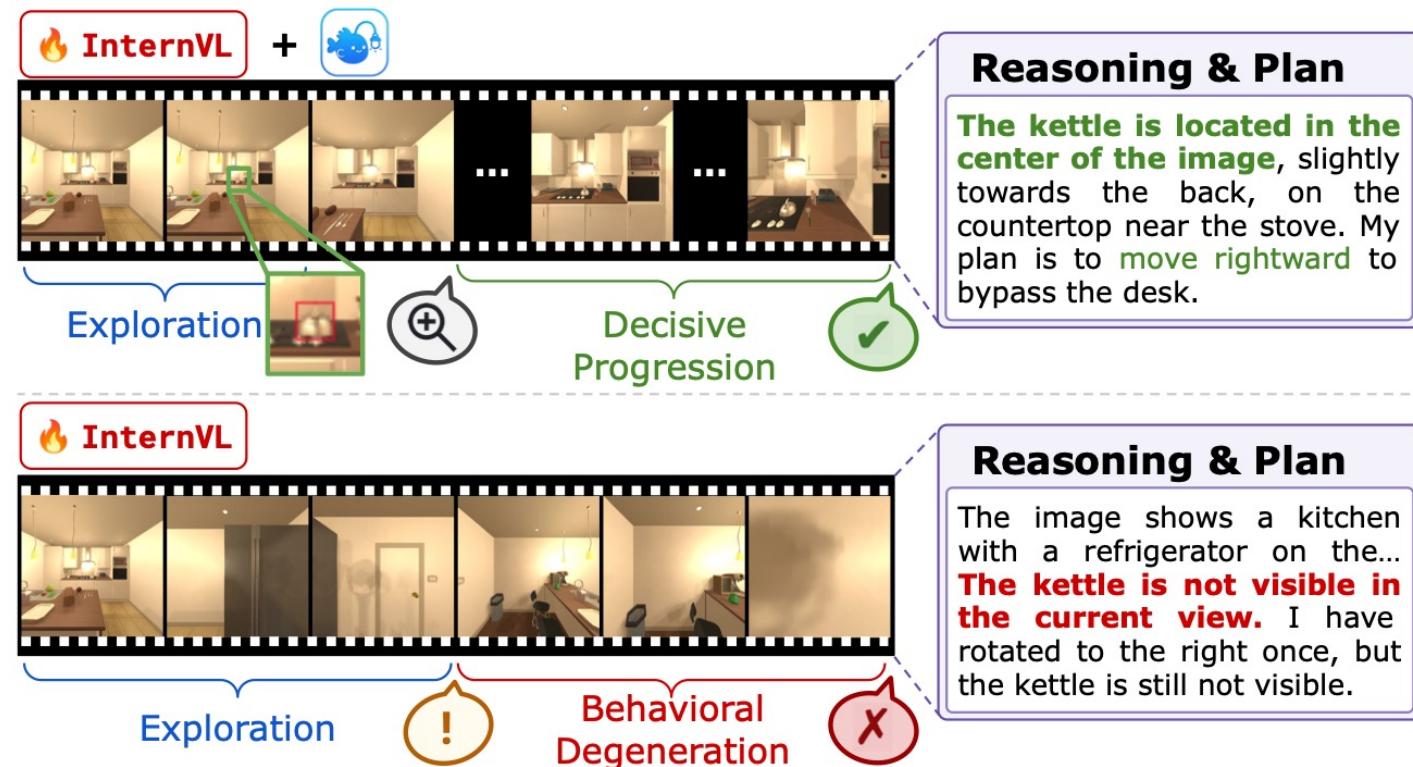
EB-Navigation

Application to Embodied Agents

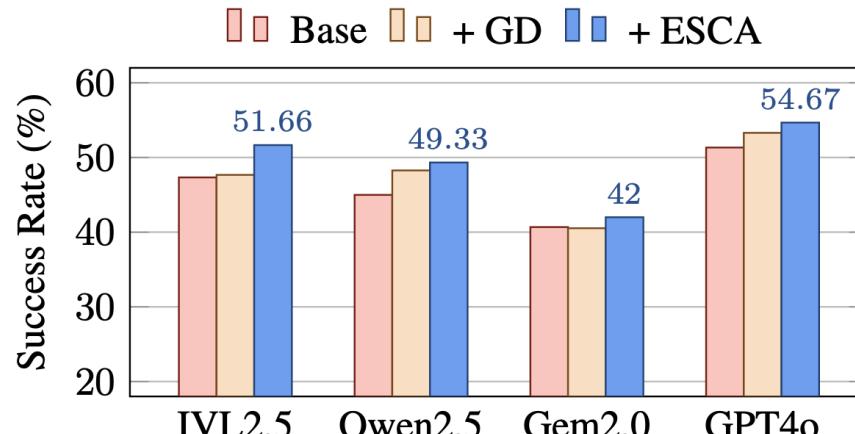


Application to Embodied Agents

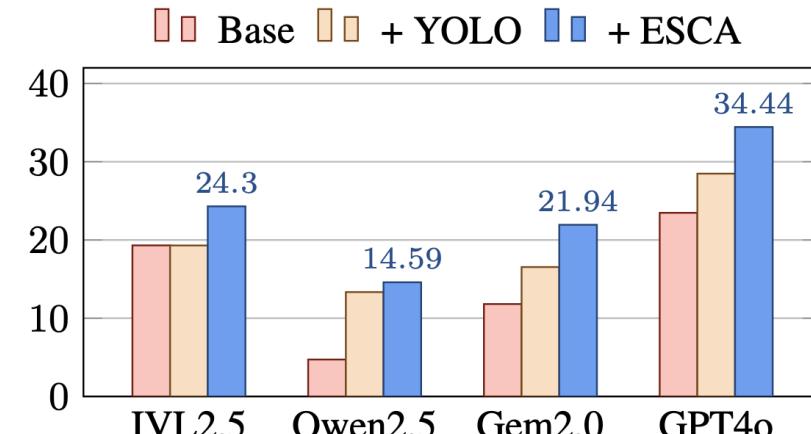
Instruction
“navigate to the kettle
on the stove”



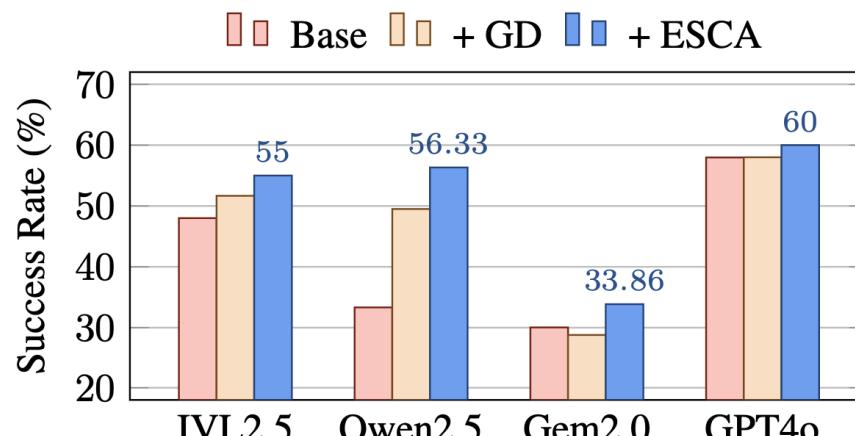
Results



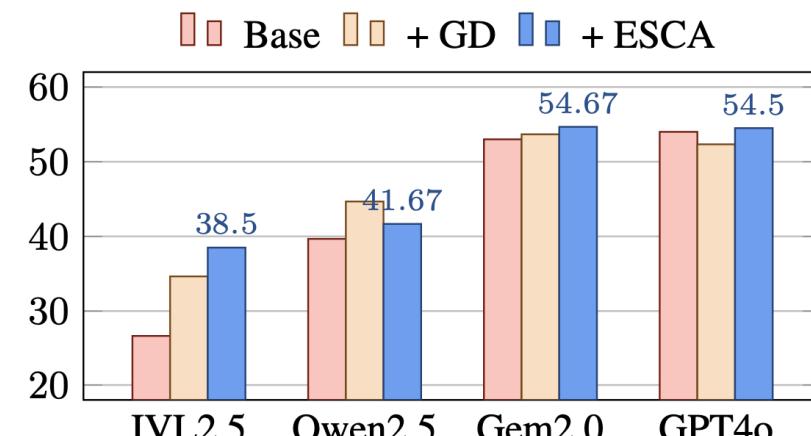
(a) EB-Navigation



(b) EB-Manipulation

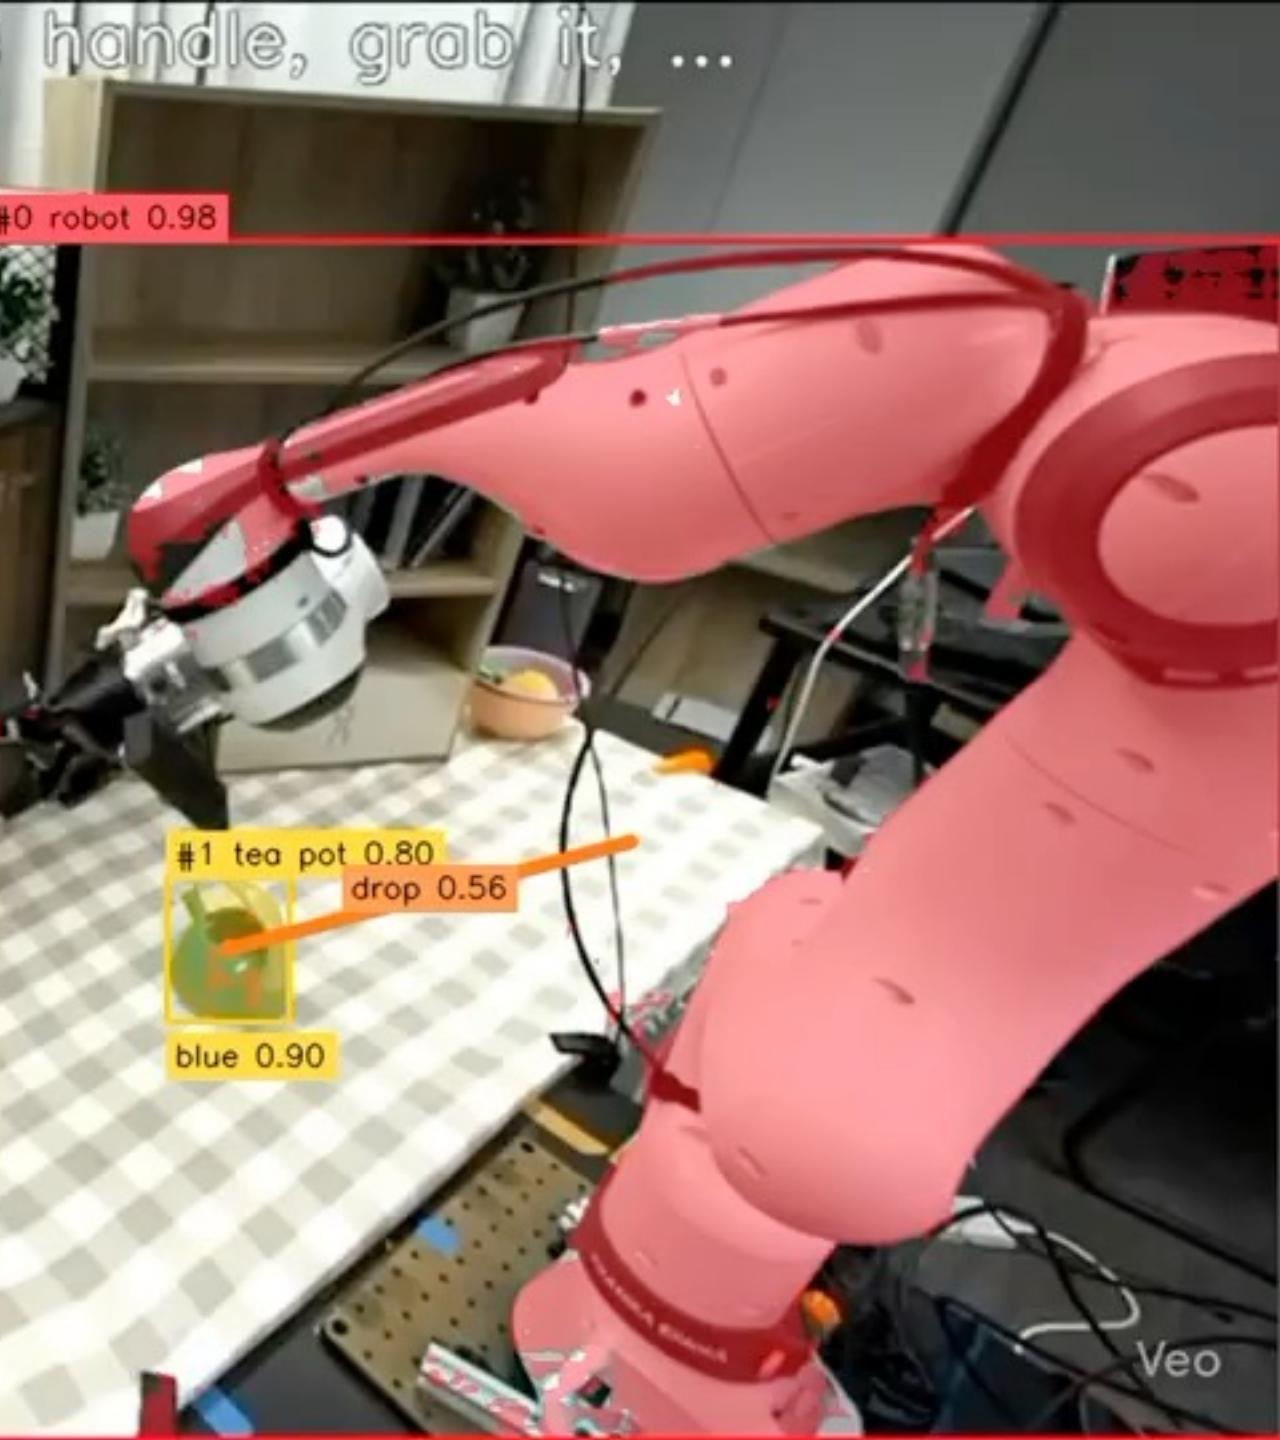


(c) EB-Habitat



(d) EB-Alfred

go to the teapot, go down to its handle, grab it, ...



Agenda

Scallop, A Language for Neurosymbolic Programming

Ziyang Li, Jiani Huang, Mayur Naik
[PLDI 2023]

LASER: A Neurosymbolic Framework For Learning Spatio-temporal Scene Graphs with Weak Supervision

Jiani Huang, Ziyang Li, Ser-Nam Lim, Mayur Naik
[ICLR 2025]

ESCA: Contextualizing Embodied Agents via Scene-Graph Generation

Jiani Huang, Ziyang Li, Matthew Kuo, Amish Sethi, Neelay Velingker, Mayank Keoliya, Ser-Nam Lim, Mayur Naik
[NeurIPS 2025]

Key Takeaway

- Grounded and accurate perception is key to embodied systems
- Structured representation for perception provides trustworthiness and transparency to downstream planning and control
- Neurosymbolic learning alleviates the need for full-supervision labels
- Neurosymbolic learning pipeline trains better mid-level and high-level perception systems
- Formal domain-specific languages (DSL) and program synthesis is the bridge connecting natural and symbolic domains

General Messages

- Programming Languages create **abstractions**
 - Scallop abstracts away probabilistic and differentiable reasoning
 - Spatial-Temporal Specifications abstract away naturalness
- Synthesized programs have usages **beyond execution**
 - Specifications help **training** the underlying vision-language model
 - (Similar to synthesized reward function or value functions during RL)
 - Synthesized Scallop program can be part of end-to-end neurosymbolic learning pipeline as a **reasoning module**

References

- [1] Scallop: From Probabilistic Deductive Databases to Scalable Differentiable Reasoning, Li et. al., NeurIPS 2021
- [2] Scallop, A Language for Neurosymbolic Programming, Li et. al., PLDI 2023
- [3] Improved Logical Reasoning of Language Models via Differentiable Symbolic Programming, Huang et. al., ACL 2023
- [4] Relational Programming with Foundation Models, Li et. al., AAAI 2024
- [5] Neurosymbolic Programming with Scallop: Principles and Practice, Li et. al., Foundations and Trends in PL, 2024
- [6] LASER: A Neuro-Symbolic Framework for Learning STSGs with Weak Supervision, Huang et. al., ICLR 2025
- [7] Lobster: A GPU-Accelerated Framework for Neurosymbolic Programming, Biberstein et. al., ASPLOS 2026
- [8] ESCA: Contextualizing Embodied Agents via Scene-Graph Generation, Huang et. al., In Submission